

Lecture 3

Instructor: Subrahmanyam Kalyanasundaram

19th August 2019

Last Class

- ▶ BST DELETE

Last Class

- ▶ BST DELETE
- ▶ Red-Black Trees
- ▶ Today we see, **INSERT** in RB Trees

Red-Black Trees

RBTs have the following properties:

1. All nodes are colored either Red or Black.
2. The root node and the leaf nodes (NIL) are black.
3. Both children of a red node are black.
No double red.
4. For any node x , all paths from x to the descendant leaves have the same number of black nodes. = Black height(x)

Red-Black Trees

RBTs have the following properties:

1. All nodes are colored either Red or Black.
2. The root node and the leaf nodes (NIL) are black.
3. Both children of a red node are black.
No double red.
4. For any node x , all paths from x to the descendant leaves have the same number of black nodes. = Black height(x)

Black height of a red black tree is the black height of its root.

A Red-Black Tree supports all procedures of a BST:

- ▶ $\text{INSERT}(val)$ – Inserts val into the RBT rooted at $node$.
- ▶ $\text{SEARCH}(val)$ – Returns True if val exists in the BST rooted at $node$. False otherwise.
- ▶ $\text{SUCC}(val)$ – Returns the smallest element greater than val in the RBT.
- ▶ $\text{PRED}(val)$ – Returns the largest element lesser than val in the RBT.
- ▶ $\text{DELETE}(val)$ – Deletes val from the RBT.

The procedures in green are implemented exactly like in a BST.

Claim

A red-black tree with black-height β has height at most 2β .

Proof sketch:

- ▶ Try to construct the longest possible path with at most β many black nodes.
- ▶ Property 4 will force you to color every alternate node black.

Observations

Theorem

If a red-black tree with n *internal* nodes and black height β , then

$$2^\beta \leq n + 1 \leq 4^\beta.$$

Observations

Theorem

If a red-black tree with n *internal* nodes and black height β , then

$$2^\beta \leq n + 1 \leq 4^\beta.$$

Proof sketch

- ▶ Merge each red node with its parent.
- ▶ Now each node has 1, 2 or 3 values with 2, 3 or 4 children.

This is a 2-3-4 tree!

Observations

Theorem

If a red-black tree with n *internal* nodes and black height β , then

$$2^\beta \leq n + 1 \leq 4^\beta.$$

Proof sketch

- ▶ Merge each red node with its parent.
- ▶ Now each node has 1, 2 or 3 values with 2, 3 or 4 children.
This is a 2-3-4 tree!
- ▶ The above 2-3-4 tree has height β .
- ▶ Thus $2^\beta - 1 \leq n \leq 4^\beta - 1$.

Observations

Theorem

A red-black tree with n *internal* nodes has height at most $2 \log(n + 1)$.

Observations

Theorem

A red-black tree with n *internal* nodes has height at most $2 \log(n + 1)$.

Proof

- ▶ We have seen that $2^\beta \leq n + 1 \leq 4^\beta$.
- ▶ That is, $1/2 \log(n + 1) \leq \beta \leq \log(n + 1)$.

Observations

Theorem

A red-black tree with n *internal* nodes has height at most $2 \log(n + 1)$.

Proof

- ▶ We have seen that $2^\beta \leq n + 1 \leq 4^\beta$.
- ▶ That is, $1/2 \log(n + 1) \leq \beta \leq \log(n + 1)$.
- ▶ Use previous claim that height is at most twice the black-height to conclude the Theorem.

INSERT procedure

$\text{INSERT}(x)$ – Insert value x into the red-back tree.

High level strategy:

- ▶ Create a node X with value x and color **red**.
- ▶ Insert node X just like inserting into a Binary Search Tree.
- ▶ Call procedure **FixINSERT** at node X .

FixINSERT procedure

Which properties might be broken when we insert a new red node?

1. All nodes are colored either Red or Black.
2. The root node is black.
3. The leaf nodes (NIL) are black.
4. Both children of a red node are black.
5. For any node, all paths from the node to the descendant leaves have the same number of black nodes.

FixINSERT procedure

Only properties 2 and 4 could be broken after inserting a red node:

1. All nodes are colored either Red or Black.
2. The root node is black.
3. The leaf nodes (NIL) are black.
4. Both children of a red node are black.
5. For any node, all paths from the node to the descendant leaves have the same number of black nodes.

FixINSERT procedure

Only properties 2 and 4 could be broken after inserting a red node:

1. All nodes are colored either Red or Black.
2. The root node is black.
3. The leaf nodes (NIL) are black.
4. Both children of a red node are black.
5. For any node, all paths from the node to the descendant leaves have the same number of black nodes.

FIXINSERT - Fixing property 2

Property 2

The root node is black.

Some invariants when FIXINSERT is called on a node Z :

- ▶ Z is colored Red.
- ▶ If Property 2 is violated, then node Z itself is the root.

FIXINSERT - Fixing property 2

Property 2

The root node is black.

Some invariants when FIXINSERT is called on a node Z :

- ▶ Z is colored Red.
- ▶ If Property 2 is violated, then node Z itself is the root.

Resolution: Simply color Z black.

FIXINSERT - Fixing Property 4

Property 4

A red node has black children.

FIXINSERT - Fixing Property 4

Property 4

A red node has black children.

Some invariants when FIXINSERT is called on a node Z :

- ▶ Z is colored Red.
- ▶ If Property 4 is violated, it is violated only by the node Z and its parent.

FIXINSERT - Fixing Property 4

Property 4

A red node has black children.

Some invariants when FIXINSERT is called on a node Z :

- ▶ Z is colored Red.
- ▶ If Property 4 is violated, it is violated only by the node Z and its parent.

There are three cases when FIXINSERT is called on a node Z .

Cases

- ▶ Case 1: Uncle of Z is Red.

Cases

- ▶ Case 1: Uncle of Z is Red.
- ▶ Case 2: Uncle of Z is black and Z is a right child of a left child.

Cases

- ▶ Case 1: Uncle of Z is Red.
- ▶ Case 2: Uncle of Z is black and Z is a right child of a left child.
- ▶ Case 3: Uncle of Z is black and Z is a left child of a left child.

Cases

- ▶ Case 1: Uncle of Z is Red.
- ▶ Case 2: Uncle of Z is black and Z is a right child of a left child.
- ▶ Case 3: Uncle of Z is black and Z is a left child of a left child.
- ▶ Other cases follow by symmetry.

Case 1

Case 1

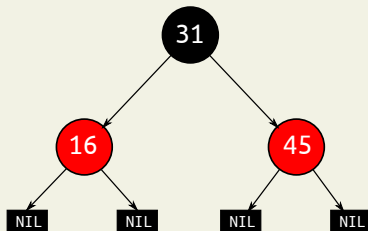
Uncle of Z is Red.

Resolution:

- ▶ Recolor parent, uncle and grandparent.
- ▶ Call `FIXINSERT(grandparent)`

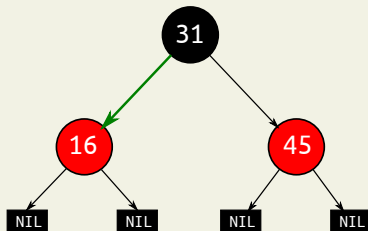
Example 1

Insert 9 to the following RBT:



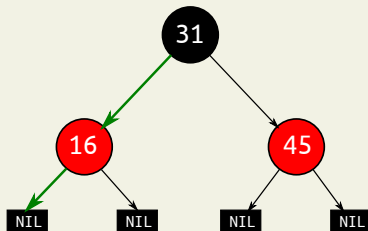
Example 1

Find the position where 9 should be inserted



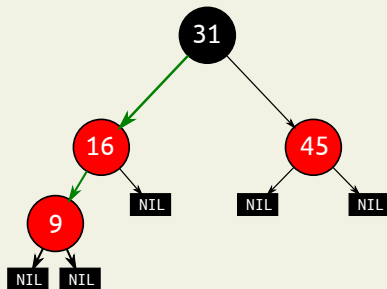
Example 1

Find the position where 9 should be inserted




Example 1

Insert 9 as a new node with color **red**



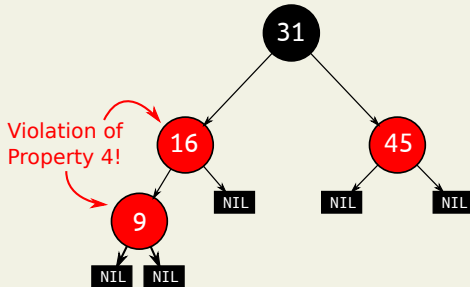
Example 1

Call FixInsert at the inserted location.



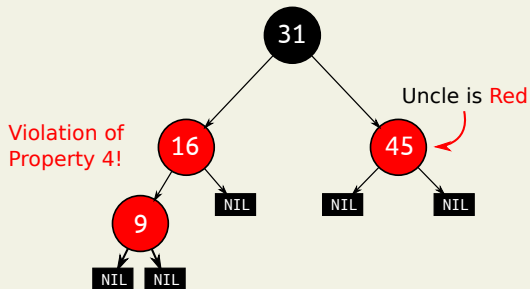
Example 1

Property 4 is violated. Check color of the uncle to determine case.



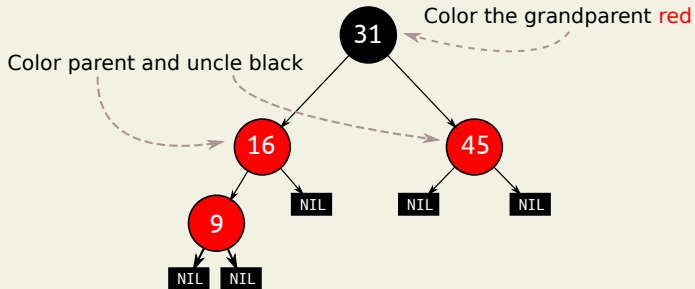
Example 1

We are in Case 1.



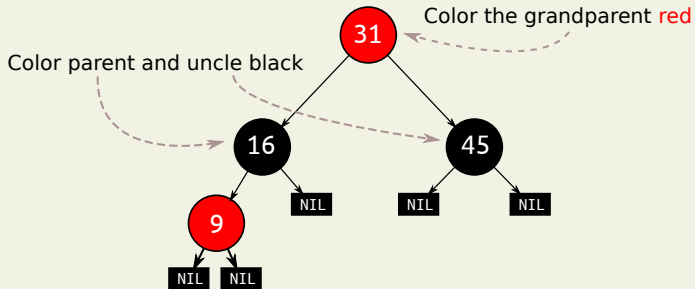
Example 1

Case 1 is resolved by recoloring.



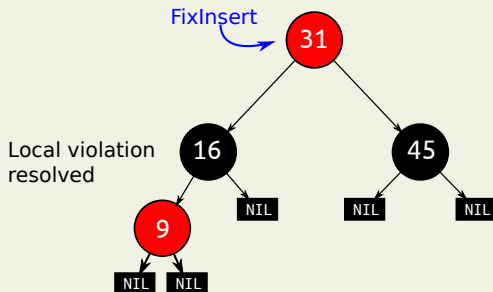
Example 1

Case 1 is resolved by recoloring.




Example 1

Now call FixInsert on the grandparent.



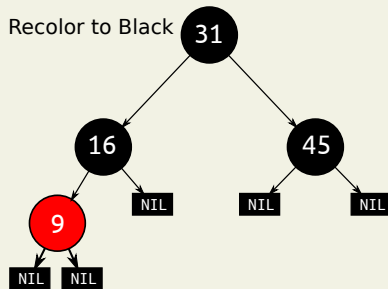
Example 1

Root node is not black.



Example 1

Simply recolor root to black.



Case 3

Case 3

Uncle of Z is Black and Z is left child of a left child.

Resolution:

- ▶ Recolor parent and grandparent.
- ▶ Rotate right at grandparent.

Case 3

Case 3

Uncle of Z is Black and Z is left child of a left child.

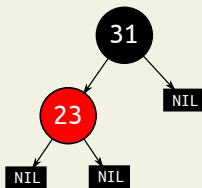
Resolution:

- ▶ Recolor parent and grandparent.
- ▶ Rotate right at grandparent.

Symmetric case: Z is right child of a right child.


Example 2

Want to insert 18 into this RBT.



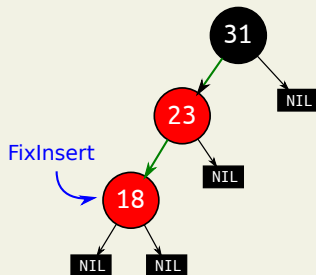
Example 2

Insert 18 as a red node according to BST property.



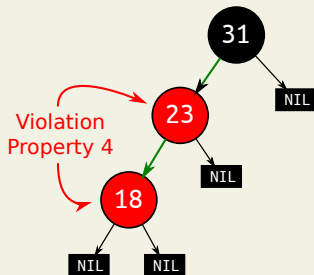
Example 2

Call FixInsert on the new node.



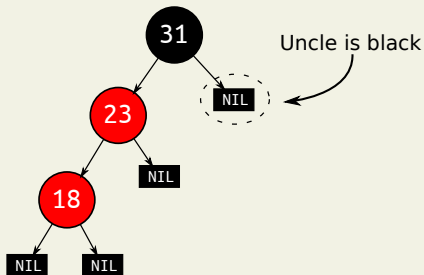
Example 2

FIXINSERT has to fix property 4.



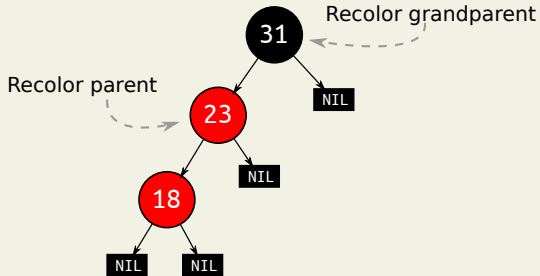
Example 2

Determine the case by looking at uncle.



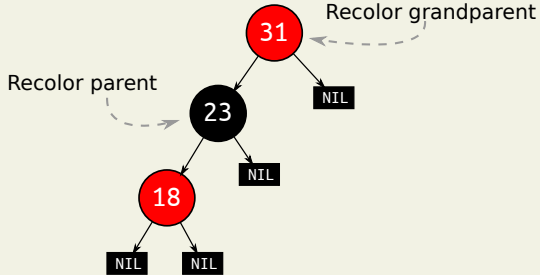
Example 2

Node 18 and its parent 23 are both left children. This is case 3.



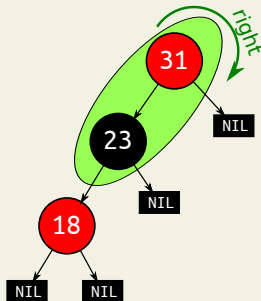
Example 2

Recolor parent to black and grandparent to red.



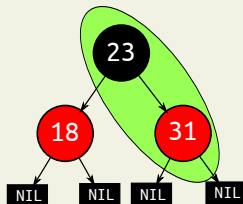
Example 2

Rotate right at grandparent.



Example 2

The resulting tree has no violations.



Case 2

Case 2

Uncle of Z is Black and Z is right child of a left child.

Resolution:

- ▶ Assign parent to Z .
- ▶ Rotate left at Z .
- ▶ Call `FIXINSERT` at Z .

The above procedure results in case 3.

Case 2

Case 2

Uncle of Z is Black and Z is right child of a left child.

Resolution:

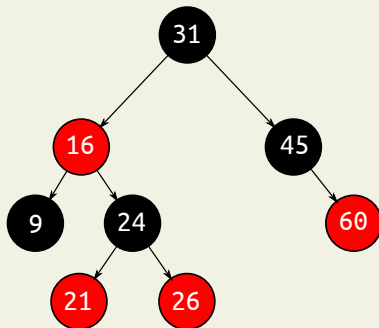
- ▶ Assign parent to Z .
- ▶ Rotate left at Z .
- ▶ Call `FIXINSERT` at Z .

The above procedure results in case 3.

Symmetric case: Z is left child of a right child.

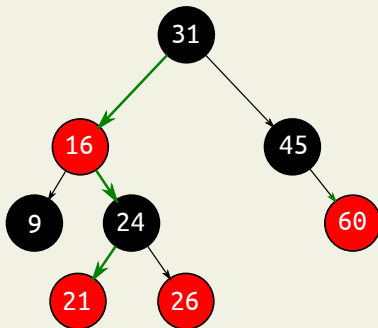
Example 3

Want to insert 18 into the following:



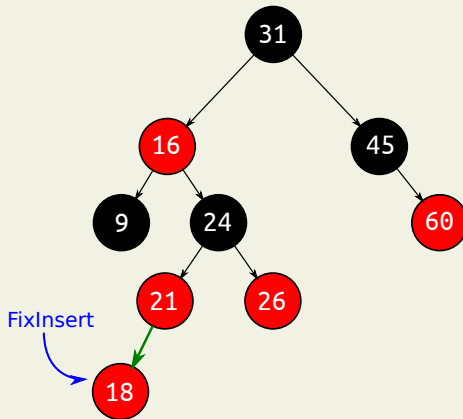
Example 3

Find the position for 18:



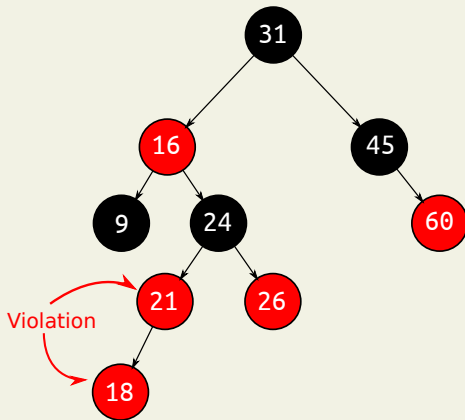
Example 3

Insert a new node with value 18 and color **red**. Call **FixInsert** at the inserted location.



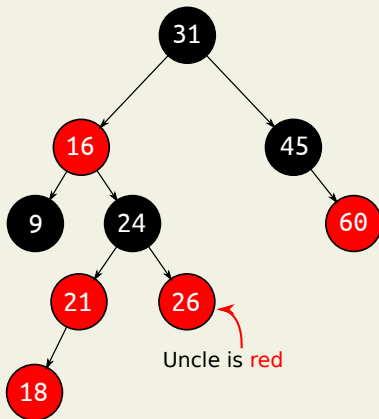
Example 3

Property 4 is violated. Check color of uncle to determine the case.




Example 3

Uncle is red. So we are in case 1.



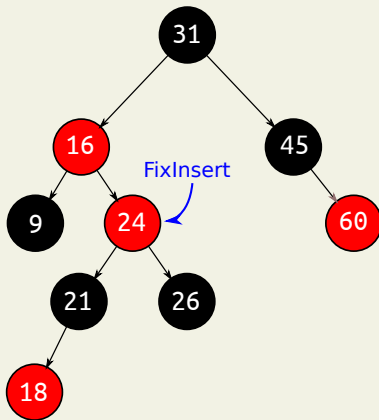
Example 3

Recolor as done earlier.



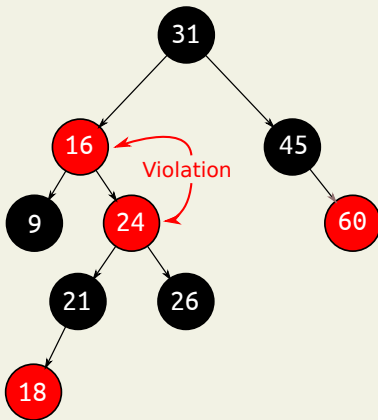
Example 3

Call FixINSERT on grandparent.




Example 3

Property 4 does not hold. Check color of Uncle to determine case.



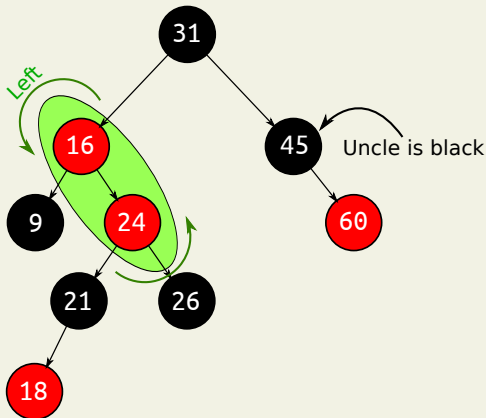
Example 3

Uncle is black and 24 is right child of a left child. So we are in case 2.



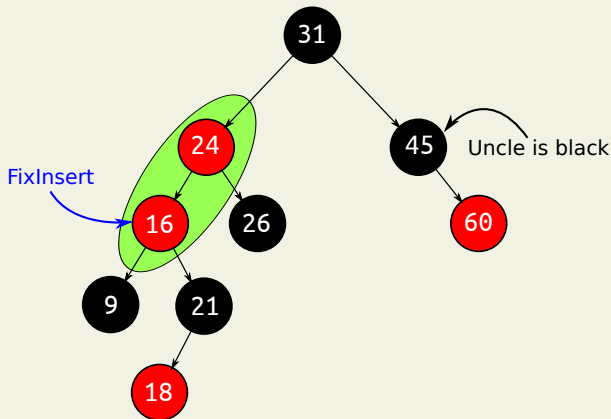
Example 3

Set Z as node with 16. Rotate left at Z and call `FixINSERT` on Z .




Example 3

Now, Uncle is black and 16 is left child of a left child. This is case 3.




Example 3

Recolor granparent red, parent black.




Example 3

Right rotate at grandparent.



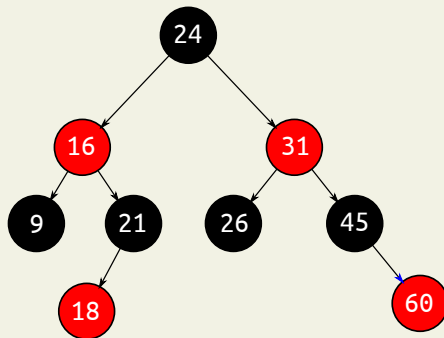
Example 3

Right rotate at grandparent.



Example 3

Done!



FixINSERT pseudocode

Algorithm 1 FixINSERT called on node Z

```
1: while color(parent( $Z$ )) = red do
2:    $U \leftarrow \text{Uncle}(Z)$ 
3:   if parent( $Z$ ) is the left child of the grandparent then
4:     if color( $U$ ) = red then
5:       Recolor parent, uncle and grandparent.
6:        $Z \leftarrow \text{grandparent}(Z)$ .
7:     else
8:       if  $Z$  is the right child then
9:          $Z \leftarrow \text{parent}(Z)$ ; Left rotate at ( $Z$ )
10:      end if
11:      Recolor parent and grandparent.
12:      Right rotate at grandparent( $Z$ ).
13:    end if
14:  end if
15: end while
16: color(root)  $\leftarrow$  black.
```
