

Lecture 8

Instructor: Subrahmanyam Kalyanasundaram

9th September 2019

Plan

- ▶ Last class, we saw 2-3-4 trees (or (2,4)-trees)
- ▶ We saw that the INSERT and DELETE procedures

Plan


- ▶ Last class, we saw 2-3-4 trees (or (2,4)-trees)
- ▶ We saw that the INSERT and DELETE procedures
- ▶ Today, we see B-trees
- ▶ After that, we see binary heaps

Exam on Thursday, 12 Sep

2-3-4 Trees

- ▶ Search trees, but not binary search trees
- ▶ Each node can store 1, 2, or 3 keys
- ▶ If a node stores d keys, then it has $d + 1$ children
- ▶ All leaf nodes are NIL nodes
- ▶ All leaf nodes are at the same level

Example



No NIL nodes are shown above

2-3-4 Trees

- ▶ What can we say about the height of a 2-3-4 tree?
- ▶ $1/2 \log(n+1) \leq h \leq \log(n+1)$
- ▶ All operations, query and modify, are $O(\log n)$

2-3-4 Trees: Implementation

Each node contains:

- ▶ d , the number of keys in the node
- ▶ x_1, x_2, \dots, x_d , the keys in increasing order
- ▶ The pointers to the $d + 1$ children
- ▶ A bit that indicates whether the node is an external node

(a, b) -tree

- ▶ 2-3-4 Trees are $(2, 4)$ -trees
- ▶ In (a, b) -trees, each node has at least $a - 1$ and at most $b - 1$ keys
- ▶ So each node has at least a and at most b children
- ▶ The lower bound of $a - 1$ is **not** applicable to the root

(a, b) -tree

- ▶ 2-3-4 Trees are $(2, 4)$ -trees
- ▶ In (a, b) -trees, each node has at least $a - 1$ and at most $b - 1$ keys
- ▶ So each node has at least a and at most b children
- ▶ The lower bound of $a - 1$ is **not** applicable to the root
- ▶ Exercise: Show that $2(a - 1) \leq b - 1$ has to be satisfied.

B-Tree

- ▶ B-Trees are (a, b) -trees with large values of a and b
- ▶ In a large database, the tree may be stored in the secondary memory
- ▶ Accessing a “page” takes time
- ▶ It helps if the entire page is a node

INSERT

- ▶ Search for the key, insert at leaf, may need to recurse up
- ▶ The procedure for INSERT in CLRS avoids recursing up
- ▶ This procedure splits every full node pre-emptively while searching
- ▶ If the leaf node needs to be split, then the parent is sure to have room to accommodate the median

INSERT

- ▶ Search for the key, insert at leaf, may need to recurse up
- ▶ The procedure for INSERT in CLRS avoids recursing up
- ▶ This procedure splits every full node pre-emptively while searching
- ▶ If the leaf node needs to be split, then the parent is sure to have room to accommodate the median
- ▶ Avoids the upwards recursion!

DELETE

- ▶ This also can be executed in one-pass
- ▶ During the search for the node, pre-emptively merge the nodes that have min no. of elements, and can't borrow

DELETE

- ▶ This also can be executed in one-pass
- ▶ During the search for the node, pre-emptively merge the nodes that have min no. of elements, and can't borrow
- ▶ Exercise: Read the procedure in CLRS

Heaps

Abstract Data Type - Heap

Heap

A max-heap supports the following functions:

- ▶ `INSERT(val)` – Inserts *val* into the heap.
- ▶ `EXTRACTMAX()` – Returns and removes the maximum element from the heap.


Binary max heap

A binary max-heap satisfies the following properties:

1. **Structural Property:** Is a complete binary tree except possibly for the lowest level, which is “left-filled”.
2. **Heap Property:** The value of a node is greater than that of both its children.


Data Structure

Heaps are usually implemented using arrays.



Data Structure


Read off from top to bottom, left to right.



Array Representation:

Data Structure

Read off from top to bottom, left to right.



Questions

About Heaps:

1. How many nodes does a height h heap have? (both bounds)
2. What is the maximum height of a heap with n nodes?

About the array implementation:

1. What is the array index of the children of the node at $A[i]$?
2. What is the array index of the right sibling of the node at $A[i]$?

Heaps using arrays

Typically, a heap is built starting with an arbitrary array:

- ▶ Procedure BUILDHEAP(Array A) – Takes an array and rearranges the elements to form a heap.

In Object Oriented languages, BUILDHEAP is essentially the *Constructor* of class Heap.

The procedure BUILDHEAP works by using a method called HEAPIFY(*node*).

Heapify

The $\text{HEAPIFY}(node)$ procedure:

- ▶ If $node$ violates the heap property:
 1. Swap value of $node$ with the largest of its two children.
 2. Call HEAPIFY on the child replaced.
- ▶ Else, do nothing and return.

Heapify

The HEAPIFY(*node*) procedure:


- ▶ If *node* violates the heap property:
 1. Swap value of *node* with the largest of its two children.
 2. Call HEAPIFY on the child replaced.
- ▶ Else, do nothing and return.

Note:

- ▶ The Heapify procedure assumes that both the subtrees under *node* are already heaps.
- ▶ It merely resolves the possible conflict between the value at *node* and its children and recurses.


Heapify

Example:




Heapify

Example:




Heapify

Example:




Heapify

Example:




Heapify

Example:



Heapify

Note that Heapify only resolves conflicts downwards.



Building a Heap

Two ways:

- ▶ William's method: Take each element and use INSERT procedure.
- ▶ Floyd's method: Take all elements in an arbitrary array. Heapify repeatedly.

Building a Heap

The procedure $\text{BUILDHEAP}(A)$ by Floyd is the following:

- ▶ For i from n to 1:
 - ▶ $\text{HEAPIFY}(i)$

Building a Heap

The procedure $\text{BUILDHEAP}(A)$ by Floyd is the following:

- ▶ For i from n to 1:
 - ▶ $\text{HEAPIFY}(i)$

Note: Indices $n/2$ to n form leaves of the heap.

The leaves are already heaps (trivially).

Hence it suffices to run the above loop from $n/2$ to 1.

Analysis of Floyd's Method

- ▶ We need to do $n/2$ HEAPIFY operations
- ▶ Each HEAPIFY can take $O(\log n)$ time
- ▶ So total time is $O(n \log n)$

Analysis of Floyd's Method

- ▶ We need to do $n/2$ HEAPIFY operations
 - ▶ Each HEAPIFY can take $O(\log n)$ time
 - ▶ So total time is $O(n \log n)$
-
- ▶ But most of the HEAPIFY operations are small
 - ▶ We have $n/2$ nodes at height 1, $n/4$ nodes at height 2 and so on
 - ▶ It can be shown that BUILDHEAP(A) takes only $O(n)$ time

Heap Sort

- ▶ Given an array A :
- ▶ Run $\text{BUILDHEAP}(A)$
- ▶ Repeatedly do $\text{EXTRACTMAX}()$
- ▶ What is the total time?

Exercises

Write the following procedures:

- ▶ **INSERT(*val*):**
 - ▶ Insert new value as the last element in the array.
 - ▶ Repeatedly Heapify *upwards* from the new element.
- ▶ **EXTRACTMAX():** Swap positions of root with last leaf. Heapfiy at new root.