

# Multiprocessing, multithreading and vectorization


Sparsh Mittal

IIT Hyderabad, India

Courtesy for some slides: S. R. Sarangi and others



A parallel program running on three computation units with three threads each with width 4 SIMD unit in each thread



# Background

# Strong and weak scaling

- **Strong scaling:** how solution time varies with the number of cores for a fixed **total** problem size.
  - Use 2X machines for a task => solve it in half the time
- **Weak scaling:** how solution time varies with the number of cores for a fixed problem size per cores.
  - if the dataset is twice as big, use 2X machines to solve the task in constant time.

# Shared Memory vs Message Passing

## \* Shared Memory

- \* All the threads share the virtual address space.
- \* They can communicate with each other by reading and writing values from/to shared memory.
  - Application ensures no data corruption (Lock/Unlock)
  - Example language: OpenMP, CUDA

## \* Message Passing

- \* Programs communicate between each other by sending and receiving messages (e.g., sending emails)
- \* They do not **share** memory addresses.
- \* Example language: MPI

# Types of Parallelism

- Instruction Level Parallelism
  - Different instructions within a stream can be executed in parallel
  - Pipelining, out-of-order execution, speculative execution, VLIW
  - Dataflow
- Data Parallelism
  - Different pieces of data can be operated on in parallel
  - SIMD: Vector processing, array processing
  - Systolic arrays, streaming processors
- Task Level Parallelism
  - Different “tasks/threads” can be executed in parallel
  - Multithreading
  - Multiprocessing (multi-core)

# Flynn's Taxonomy

# Flynn's Classification

- \* **Instruction stream** → Set of instructions that are executed
- \* **Data stream** → Data values that the instructions process
- \* Four **types** of multiprocessors : **SISD, SIMD, MISD, MIMD**



# SISD and SIMD

- \* SISD → Standard uniprocessor
- \* SIMD → One instruction, operates on multiple pieces of data. Vector processors have one instruction that operates on many pieces of data in parallel. For example, one instruction can compute the  $\sin^{-1}$  of 4 values in parallel.

# MISD

- \* MISD → Multiple Instruction Single Data
  - \* Very rare in practice
  - \* Consider an aircraft that has a MIPS, an ARM, and an X86 processor operating on the same data (multiple instruction streams)
  - \* We have different instructions operating on the same data
  - \* The final outcome is decided on the basis of a majority vote.

# MIMD

- \* MIMD → Multiple **instruction**, multiple **data** (two types, SPMD, MPMD)
- \* SPMD → Single **program**, multiple **data**.  
Examples: **OpenMP** or **MPI** programs. We typically have **multiple processes** or **threads** executing the same **program** with different inputs.
- \* MPMD → A **master program**, delegates work to multiple **slave programs**. The programs are different.

# Summary

- **SISD**: Single instruction operates on single data element
- **SIMD**: Single instruction operates on multiple data elements
  - Array processor
  - Vector processor
- **MISD**: Multiple instructions operate on single data element
  - Closest form: systolic array processor, streaming processor
- **MIMD**: Multiple instructions operate on multiple data elements (multiple instruction streams)
  - Multiprocessor
  - Multithreaded processor

# Multiprocessing

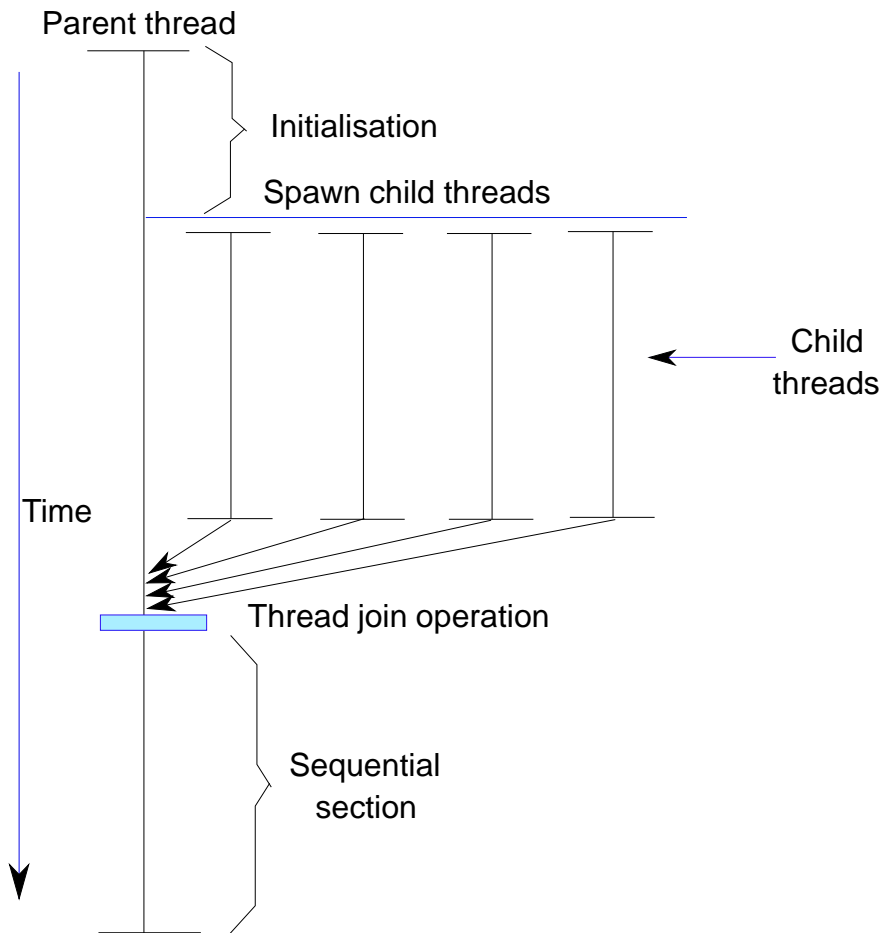
- \* The term **multiprocessing** refers to **multiple processors** working in **parallel**.
- \* This is a generic definition, and it can refer to multiple processors in the **same chip**, or processors across **different chips**.
- \* A **multicore processor** is a specific type of multiprocessor that contains all of its constituent processors in the **same chip**. Each such processor is known as a **core**.

# Multithreading

# The Notion of Threads

- \* We spawn a set of separate **threads**
- \* Properties of **threads**
  - \* A **thread** shares its **address space** with other threads
  - \* It has its own **program counter**, set of **registers**, and **stack**
  - \* A **process** contains multiple **threads**
  - \* Threads **communicate** with each other by writing values to memory or via synchronization operations

# Operation of the Program





# Multithreading


- \* **Multithreading** → A design paradigm that proposes to run multiple threads on the same pipeline.
- \* Three types
  - \* Coarse grained
  - \* Fine grained
  - \* Simultaneous

# Analogy

- Consider a car which can be shared by 4 people (A, B, C, D) in following way:
- Option1: Four people ride the car simultaneously almost all the time to go to their destinations.
- Option2: A uses the car for 15 days, B for 14 days, C for 18 days and D for 16 days (and repeat).
- Option3: A uses the car for 8 hours, B for 9 hours, C for 5 hours and D for 7 hours (and repeat).
- Match the above three options to three types of multithreading

# Coarse Grained Multithreading

- \* **Assume** that we want to run 4 threads on a pipeline
- \* Run thread 1 for  $n$  cycles, run thread 2 for  $n$  cycles, ....



# Implémentation

- \* Steps to minimize the context switch overhead
- \* For a 4-way coarse grained MT machine
  - \* 4 program counters
  - \* 4 register files
  - \* 4 flags registers
  - \* A context register that contains a thread id.
  - \* Zero overhead context switching → Change the thread id in the context register

# Advantages

- \* Assume that **thread 1** has an **L2 miss**
  - \* **Wait** for 200 cycles
  - \* Schedule thread 2
  - \* Now let us say that thread 2 has an L2 miss
  - \* Schedule thread 3
- \* We can have a **sophisticated algorithm** that switches every n cycles, or when there is a **long latency event** such as an L2 miss.
  - \* Minimises idle cycles for the entire system

# Fine Grained Multithreading

- \* The **switching granularity** is very small
  - \* 1-2 cycles
- \* **Advantage** :
  - \* Can take **advantage** of **low latency** events such as division, or L1 cache misses
  - \* **Minimise** idle cycles to an even greater extent
- \* **Correctness Issues**
  - \* We can have instructions of 2 threads simultaneously in the **pipeline**.
  - \* We never forward/interlock for **instructions** across **threads**

# Simultaneous Multithreading

- \* Most modern processors have **multiple issue slots**
  - \* Can issue **multiple instructions** to the **functional units**
  - \* For example, a 3 issue **processor** can **fetch, decode, and execute** 3 instructions per cycle
  - \* If a benchmark has low **ILP** (instruction level parallelism), then fine and coarse grained multithreading cannot **really help**.

# Simultaneous Multithreading

- \* Main Idea

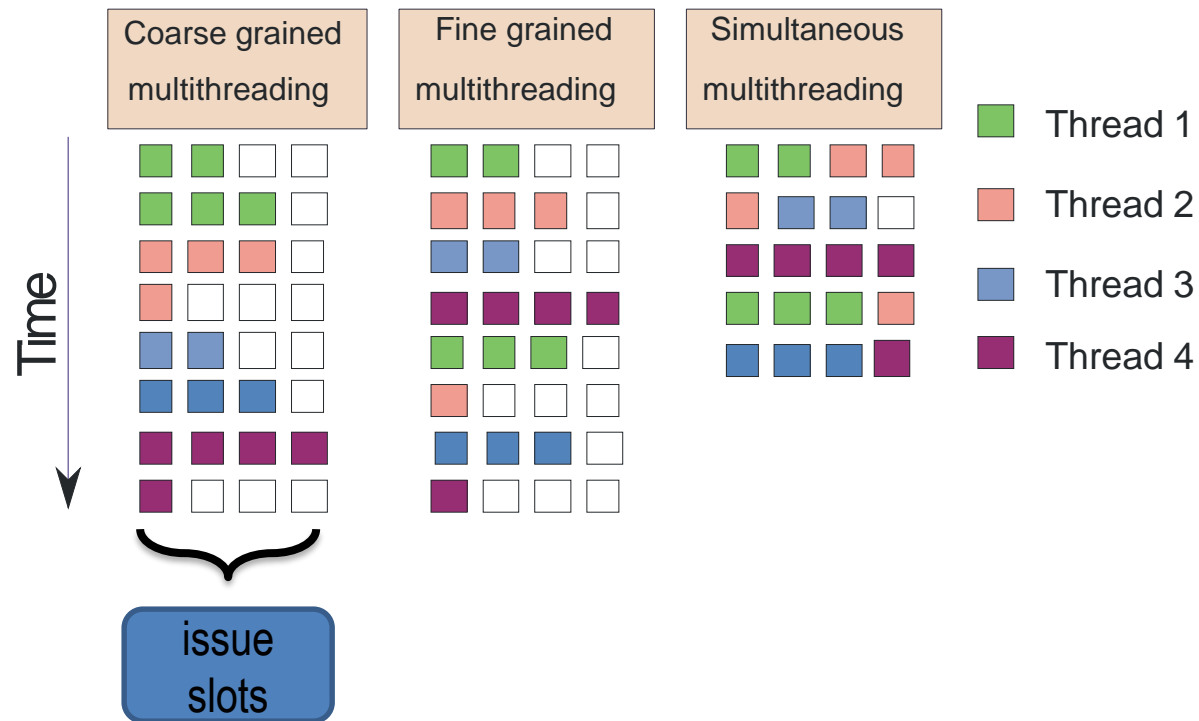
- \* Partition the issue slots across threads
- \* Scenario : In the same cycle
  - \* Issue 2 instructions for thread 1
  - \* and, issue 1 instruction for thread 2
  - \* and, issue 1 instruction for thread 3

- \* Support required

- \* Need smart instruction selection logic.
  - \* Balance fairness and throughput




# Summary



# Vectorization (and vector processor)

A parallel program running on three computation units with three threads each with width 4 SIMD unit in each thread



**BIG PICTURE**

# Vectorization

Scalar Instructions

$$\begin{array}{l} 4 + 1 = 5 \\ 0 + 3 = 3 \\ -2 + 8 = 6 \\ 9 + -7 = 2 \end{array}$$

© ColfaxResearch.com

Vector Instructions

$$\begin{array}{l} 4 \\ 0 \\ -2 \\ 9 \end{array} + \begin{array}{l} 1 \\ 3 \\ 8 \\ -7 \end{array} = \begin{array}{l} 5 \\ 3 \\ 6 \\ 2 \end{array}$$

Vector Length

# Some of the SIMD instruction sets used in industry

| Register size | Instruction set |
|---------------|-----------------|
| 80 bits       | MMX             |
| 128-bits      | SSE1/SSE2 etc.  |
| 256-bits      | AVX/AVX2        |
| 512-bits      | AVX-512         |


# Vector Processors

- \* A **vector instruction** operates on arrays of data
  - \* Example : There are **vector instructions** to add or **multiply** two **arrays** of data, and produce an array as output
  - \* **Advantage** : Can be used to perform all kinds of array, matrix, and linear algebra **operations**. These **operations** form the core of many scientific programs, high intensity graphics, and data analytics applications.

# Software Interface

- \* Let us define a **vector register**
  - \* Example : 128 bit registers in the **MMX instruction set**  
→ XMM0 ... XMM15
  - \* Can hold 4 floating point values, or 8 2-byte short integers
  - \* **Addition** of **vector registers** is equivalent to pairwise addition of each of the individual **elements**.
  - \* The **result** is saved in a **vector register** of the same size.

# Example of Vector Addition



Let us define 8 128 bit vector registers in SimpleRisc. vr0 ... vr7



# Loading Vector Registers

- \* There are two **options** :

- \* **Option 1** : We assume that the data elements are stored in **contiguous locations**
- \* Let us define the v.ld **instruction** that uses this **assumption**.

| Instruction      | Semantics   |
|------------------|---|
| v.ld vr1, 12[r1] | $vr1 \leftarrow ([r1+12], [r1+16], [r1+20], [r1+24])$ |

- \* **Option 2**: Assume that the elements are not saved in contiguous locations.
  - \* For this, there are scatter-gather instructions

# Scatter Gather Operation

- \* The data is **scattered** in memory
  - \* The **load operation** needs to **gather** the data and save it in a **vector register**.
- \* Let us define a **scatter gather** version of the **load instruction** → v.sg.ld
  - \* It uses another **vector register** that contains the addresses of each of the elements.

| Instruction      | Semantics   |
|------------------|---|
| v.sg.ld vr1, vr2 | $vr1 \leftarrow ([vr2[0]], [vr2[1]], [vr2[2]], [vr2[3]])$ |

# Vector Store Operation

- \* We can similarly define two vector store operations

| Instruction      | Semantics  |
|------------------|--|
| v.sg.st vr1, vr2 | $[vr2[0]] \leftarrow vr1[0]$<br>$[vr2[1]] \leftarrow vr1[1]$<br>$[vr2[2]] \leftarrow vr1[2]$<br>$[vr2[3]] \leftarrow vr1[3]$ |

| Instruction      | Semantics  |
|------------------|--|
| v.st vr1, 12[r1] | $[r1+12] \leftarrow vr1[0]$<br>$[r1+16] \leftarrow vr1[1]$<br>$[r1+20] \leftarrow vr1[2]$<br>$[r1+24] \leftarrow vr1[3]$ |

# Vector Operations

- \* We can now define **custom operations** on vector registers
  - \* `v.add` → Adds two vector registers
  - \* `v.mul` → Multiplies two vector registers
  - \* We can even have operations that have a **vector operand** and a **scalar operand** → Multiply a **vector** with a **scalar**.

# Design of a Vector Processor

## \* Salient Points

- \* We have a **vector register file** and a **scalar register file**
- \* There are scalar and vector functional units
- \* Unless we are converting a vector to a scalar or vice versa, we in **general** do not **forward** values between **vector** and **scalar** instructions
- \* The **memory** unit needs support for **regular** operations, **vector** operations, and possibly **scatter-gather** operations.

# References

- S. Mittal et al, “A Survey on Evaluating and Optimizing Performance of Intel Xeon Phi” 2019 ([pdf](#))