

CS3510

Operating Systems

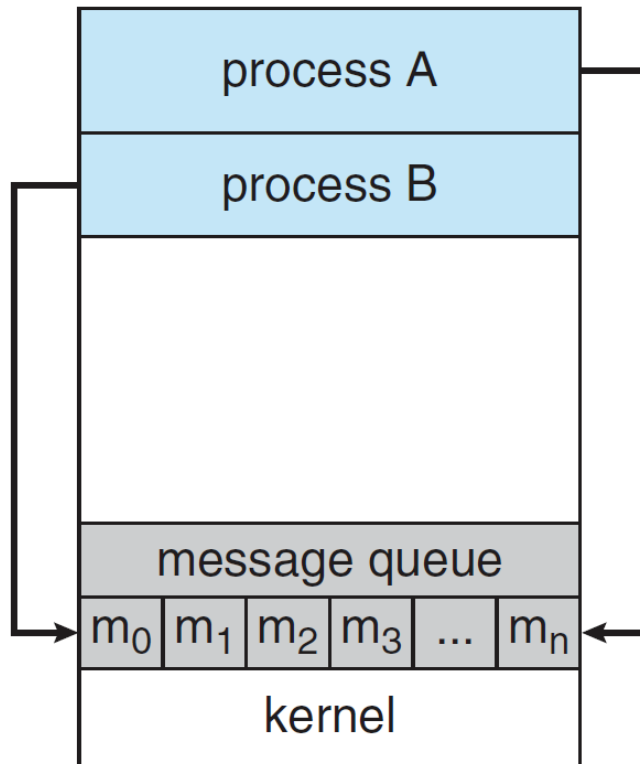
IPC

Bheemarjuna Reddy
IIT Hyderabad


Inter Process Communication (IPC)

- Independent vs Cooperating processes
- Why let processes cooperate?
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- Two fundamental models
 - Message Passing
 - Shared Memory

Communication Models



Message Passing




Shared Memory

Shared Memory

- Processes establish a segment of memory as shared
 - Typically part of the memory of the process creating the shared memory.
 - Other processes attach this to their memory space.
 - Good when we have to share a lots of data
- Requires processes to agree to remove memory protection for the shared section
 - Recall that OS normally protects processes from writing in each others memory.

Shared Memory Communication



- Communication occurs by “simply” reading/writing to shared address page
 - Really low overhead communication and faster way of communication
 - But introduces complex synchronization problems
 - Cache coherence problem in multi-core systems

```
#include <stdio.h>
#include <sys/shm.h>
#include <sys/stat.h>
// Example using System V shared memory objects
// shm_open and mmap in POSIX
main(int argc, char **argv) {
    char* shared_memory;
    const int size = 4096;
    int segment_id = shmget(IPC_PRIVATE, size, IPC_CREAT | 0666);
    int cpid = fork();
    if (cpid == 0)
    {
        shared_memory = (char*) shmat(segment_id, NULL, 0); //attach
        sprintf(shared_memory, "Hi from process %d", getpid());
    }
    else
    {
        wait(NULL);
        shared_memory = (char*) shmat(segment_id, NULL, 0); //attach
        printf("Process %d read: %s\n", getpid(), shared_memory);
        shmdt(shared_memory); //detach
        shmctl(segment_id, IPC_RMID, NULL); //remove segment
    }
}
```

E x a m p l e

```
#include <fcntl> #include <string.h>
#include <sys/shm.h> #include <sys/stat.h>
#define MAX_LEN 10000
struct region {      /* Defines "structure" of shared memory */
    int len;
    char buf[MAX_LEN]; };
struct region *rptr;
int fd; char * msg="Hello";
/* Create shared memory object and set its size */
fd = shm_open("/myregion", O_CREAT | O_RDWR, S_IRUSR |
S_IWUSR);
if (fd == -1) ... /* Handle error */;
if (ftruncate(fd, sizeof(struct region)) == -1) //set Size
... /* Handle error */;
/* Map shared memory object to process' address space */
rptr = mmap(NULL, sizeof(struct region), PROT_READ |
PROT_WRITE, MAP_SHARED, fd, 0);
if (rptr == MAP_FAILED)
    /* Handle error */;
/* Now we can refer to mapped region using fields of rptr */
sprintf(rptr,"%s",msg); //write to shared memory
rptr+=strlen(msg);
...
```

```
#include <fcntl> #include <string.h>
#include <sys/shm.h> #include <sys/stat.h>
```

```
int main()
{
    /* name of the shared memory object */
    const char *name = "/myregion";
    /* shared memory file descriptor */
    int shm fd;
    /* pointer to shared memory object */
    void *ptr;
    /* open the shared memory object */
    shm fd = shm open(name, O_RDONLY, 0666);
    /* memory map shared memory object to process' address space */
    ptr = mmap(0, sizeof(struct region), PROT_READ, MAP_SHARED,
    shm fd, 0);
    /* read from the shared memory object */
    printf("%s", (char *)ptr);
    /* remove the shared memory object */
    shm_unlink(name);
    return 0;
}
```


Message Passing

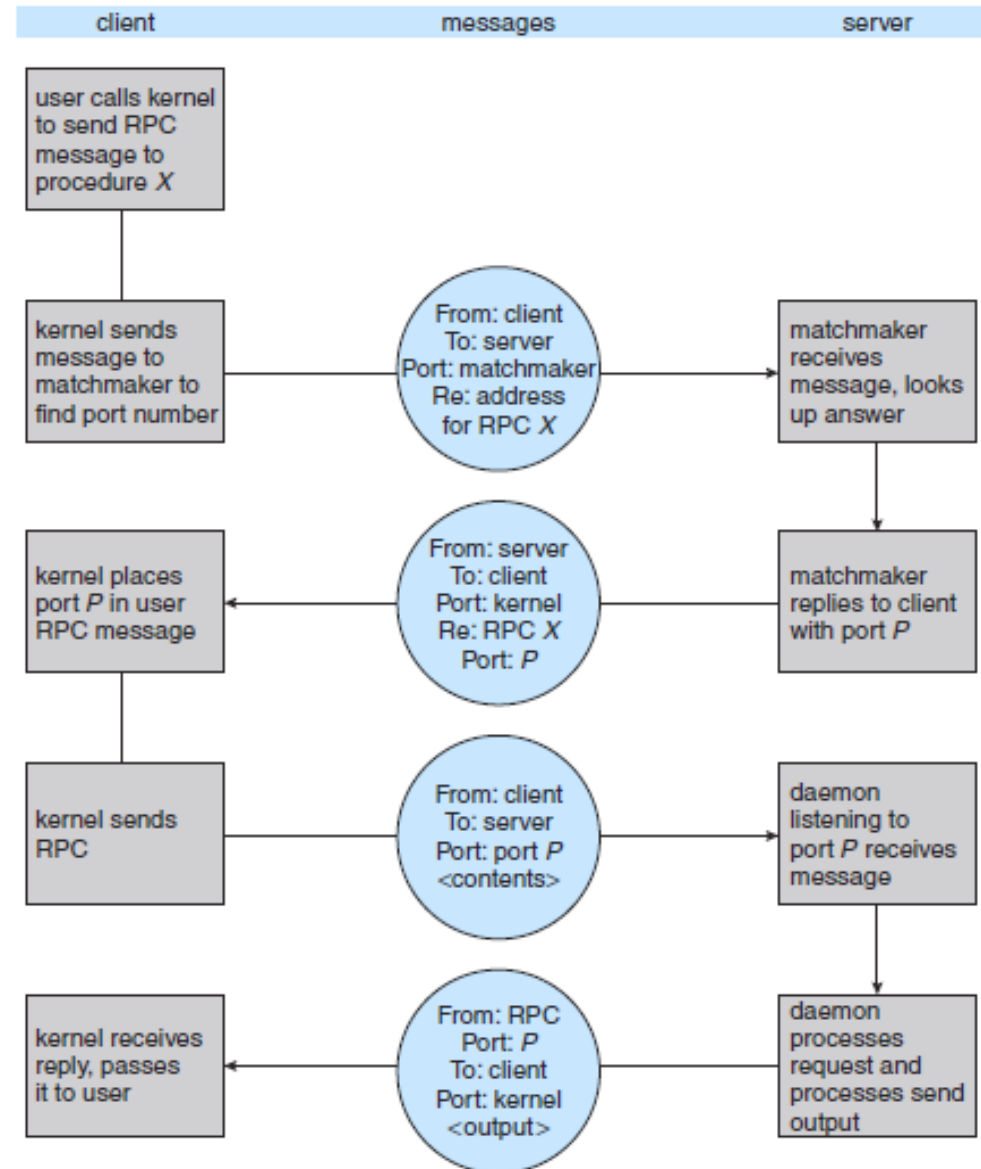
- Useful for exchanging small amount of data w/o any conflicts
- Send(P, msg): Send msg to process P
 - Fixed vs variable size msg
- Recv(Q, msg): Receive msg from process Q
- Typically requires kernel intervention
 - User mode to kernel mode for Sending
 - Kernel mode to User mode for Receiving
- Communication link is needed for msg passing
 - Physical link realization
 - » shared memory, hardware bus, network
 - Logical implementation of link and its basic operations
 - » Direct vs indirect communication
 - Direct communication
 - Hardcode sender/receiver IDs (Symmetry)
 - Hardcode sender only (Asymmetry)
 - Indirection using mailboxes/ports
 - Owned by Process (owner/User) vs Owned by OS
 - Send(boxA msg) and Receive(boxA msg)

Message Passing


- Logical implementation of link and its basic operations
 - » Synchronous vs asynchronous communication
 - » Automatic vs explicit buffering
- Possible impl. of send()/receive() primitives
 - Blocking send/receive: process is blocked till msg is tx/rx
 - Non-blocking send/receive: Send msg & resume; Receive valid msg or return NULL w/o blocking
- When both send and receive are blocking, no buffer is needed. Other combinations need *buffering*.
 - Zero capacity buffer
 - » Needs synchronous sender.
 - Bounded capacity buffer
 - » If the buffer is full, the sender blocks.
 - Unbounded capacity buffer
 - » The sender never blocks
- Easier to implement in distributed and multi-core sys with NUMA compared to shared memory, but slower than that (sys calls)

Additional Communication mechanisms for Client-Server Systems

- **Sockets:** end point of communication; pair of sockets for processes to talk over network
- **RPC:** Executing a procedure in remote m/c
 - Msg is addressed to a RPC daemon listening on a port in remote m/c & carries function details for exec
 - Output is sent back via a Msg to local m/c



Additional Communication mechanisms for Client-Server Systems



- **Pipes**

- A pipe is a stream of communication between two processes
- You can think of it as a virtual file stream shared between two processes
- A process can read and/or write to a pipe
- Two processes can communicate via a pipe without even knowing it.
- This forms the backbone of Unix-like environments.
- The pipe function gets two descriptors (integer labels)
- Read descriptor - read from the pipe
- Write descriptor - write to the pipe
- Both processes must know the descriptors
- read and write are used with the pipe

```

int main(void)
{
    int pid;
    char buffer[1024];
    int fd[2];

    pipe(fd); /* ordinary pipes; fd[0] is for read-end, fd[1] is for write-end of pipe */

    pid = fork();

    if (pid == 0) /* child */
    {
        int count;
        close(fd[0]); /* close unused READ end, child will write */

        /* prompt user for input */
        printf("input: ");
        fgets(buffer, sizeof(buffer), stdin);
        printf("child: message is %s", buffer);

        /* write to the pipe (include NUL terminator) */
        count = write(fd[1], buffer, strlen(buffer) + 1); //pipe is a special type of file
        printf("child: wrote %i bytes\n", count);
        close(fd[1]);
        exit(0);
    }
    else /* parent */
    {
        int count;
        close(fd[1]); /* close unused WRITE end */
        wait(NULL); /* reap the child */
        /* read from the pipe */
        count = read(fd[0], buffer, sizeof(buffer));
        printf("parent: message is %s", buffer);
        printf("parent: read %i bytes\n", count);
        close(fd[0]);
    }
}

```

Pipes

- An ordinary pipe cannot be accessed from outside the process that created it.
 - A parent will create a pipe, then fork so the child can access it.
 - Child processes inherit all open files (pipes are a special kind of file) from the parent.
 - Once the processes end, the pipes no longer exist.
 - Ordinary pipes can only be used with processes on the same machine
- Named pipes are more powerful than ordinary pipes.
 - They can be used by several processes at once.
 - They don't require a parent-child relationship.
 - They exist independently of the process that created them.
 - » Much like how files created on the disk by a process exist after the process ends.
 - See mkfifo on Unix-based systems and CreateNamedPipe on Windows
 - Unix named Pipes: FIFOs
 - » Bidirectional, but half-duplex; appear like files in fileSystem
 - » Communicating processes must reside in the same machine.

Summary

- **IPC**
 - Shared memory
 - Message passing
 - RPC
 - Pipes
- We discuss other IPC mechanisms, related to synchronization of processes, later

Reading Assignment

- Chapter 3 from OSC by Galvin et al
- Chapter 2 from MOS by Tanenbaum et al
- Chapter 3. Processes and Chapter 19. Process Communication from Understanding the Linux Kernel by Daniel P. Bovet and Marco Cesati (available on Intranet)
- The Linux Programming Interface by Michael Kerrisk
- http://man7.org/linux/man-pages/man3/shm_open.3.html
- http://pubs.opengroup.org/onlinepubs/009695399/functions/shm_open.html
- <http://man7.org/linux/man-pages/man2/open.2.html>
- <http://man7.org/linux/man-pages/man2/shmget.2.html>