

---

**CS3510**

**Threads**

**Bheemarjuna Reddy Tamma**  
**IIT HYD**

# Outline

---

- Case for Threads
- Threads vs Processes
- Thread details
  - Pthread Library

# Case for Parallelism

---

```
main()
  read_data();
  for(all data i ← 1 to N)
    compute(i);
    write_data(i);
  endfor
```

Blocking Write

```
main()
  read_data();
  for(all data i ← 1 to N)
    compute(i);
    CreateProcess(write_data(i));
  endfor
```

Does Writing in  
new process

# Case for Parallelism

---

Consider the following code fragment:

```
for(k = 0; k < n; k++)  
    a[k] = b[k] * c[k] + d[k] * e[k];
```

Instead:

```
fn(l, m) {  
    for(k = l; k < m; k++)  
        a[k] = b[k] * c[k] + d[k] * e[k];  
}
```

CreateProcess(fn, 0, n/2);

CreateProcess(fn, n/2, n);

# Case for Parallelism

---

Consider a Web server

create a number of processes, and for each process do

- get network message (HTTP REQ) from client
- get URL data from disk
- compose response
- send response (HTML Object)

☞ Server connections are fast, but client connections may not be

- Takes server a loooong time to feed the response to client
- While it's doing that it can't service any more requests

# Parallel Programs

---

- To build parallel programs, such as:
  - Parallel execution on a multiprocessor
  - Web server to handle multiple simultaneous web requests
- We will need to:
  - Create several processes that can execute in parallel
  - **Cause each to map to the same address space**
    - because they're part of the same computation
  - Give each its starting address and initial parameters
  - The OS will then schedule these processes in parallel

# Parallelism vs Concurrency

---

- Both are not same, but typically used interchangeably!
- **Parallelism:** Doing multiple tasks simultaneously
  - E.g., driving while talking!
  - E.g., Running Word App and Media player simultaneously, each one on separate Core/CPU
  - Not possible w/o multiple Cores
- **Concurrency:** Making progress for multiple tasks
  - Single CPU: time-sharing of CPU resource with time slices
  - Multiple CPU: same as parallelism
- So, Parallelism → Concurrency, but **not vice versa**.

# Processes Overheads

---

- A full process includes numerous things:
  - an address space (defining all the code and data segments)
  - OS resources and accounting information
  - a “thread of control”,
    - defines where the process is currently executing
    - That is the PC and registers

## Creating a new process is costly

- all of the structures (e.g., page tables, address space) that must be allocated
- Costly even after copy-on-write optimization

## Communicating between processes is costly

- most communication goes through the OS to avoid synchronization issues with shared resources



# Need “Lightweight” Processes

---

- What's similar in these processes?
  - They all share the same code, heap and data (most of the address space)
  - They all share the same privileges
  - They share almost everything in the process
- What don't they share?
  - Each has its own PC, register set, stack, and stack pointer
- Idea: why don't we separate the idea of process (address space, accounting, etc.) from that of the minimal “thread of control” (PC, SP, registers)?

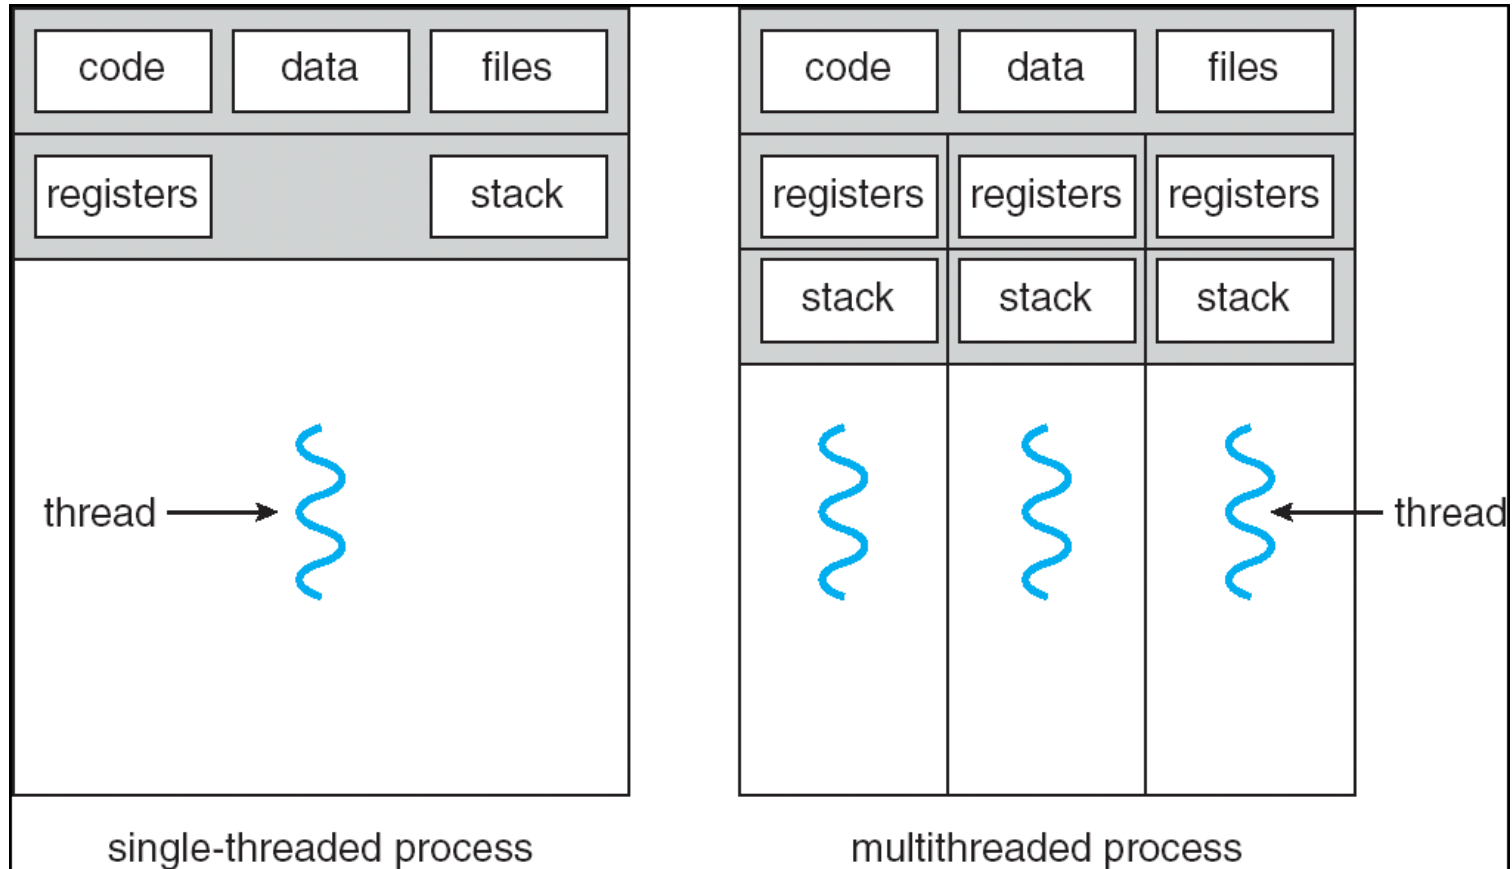
# Threads and Processes

---

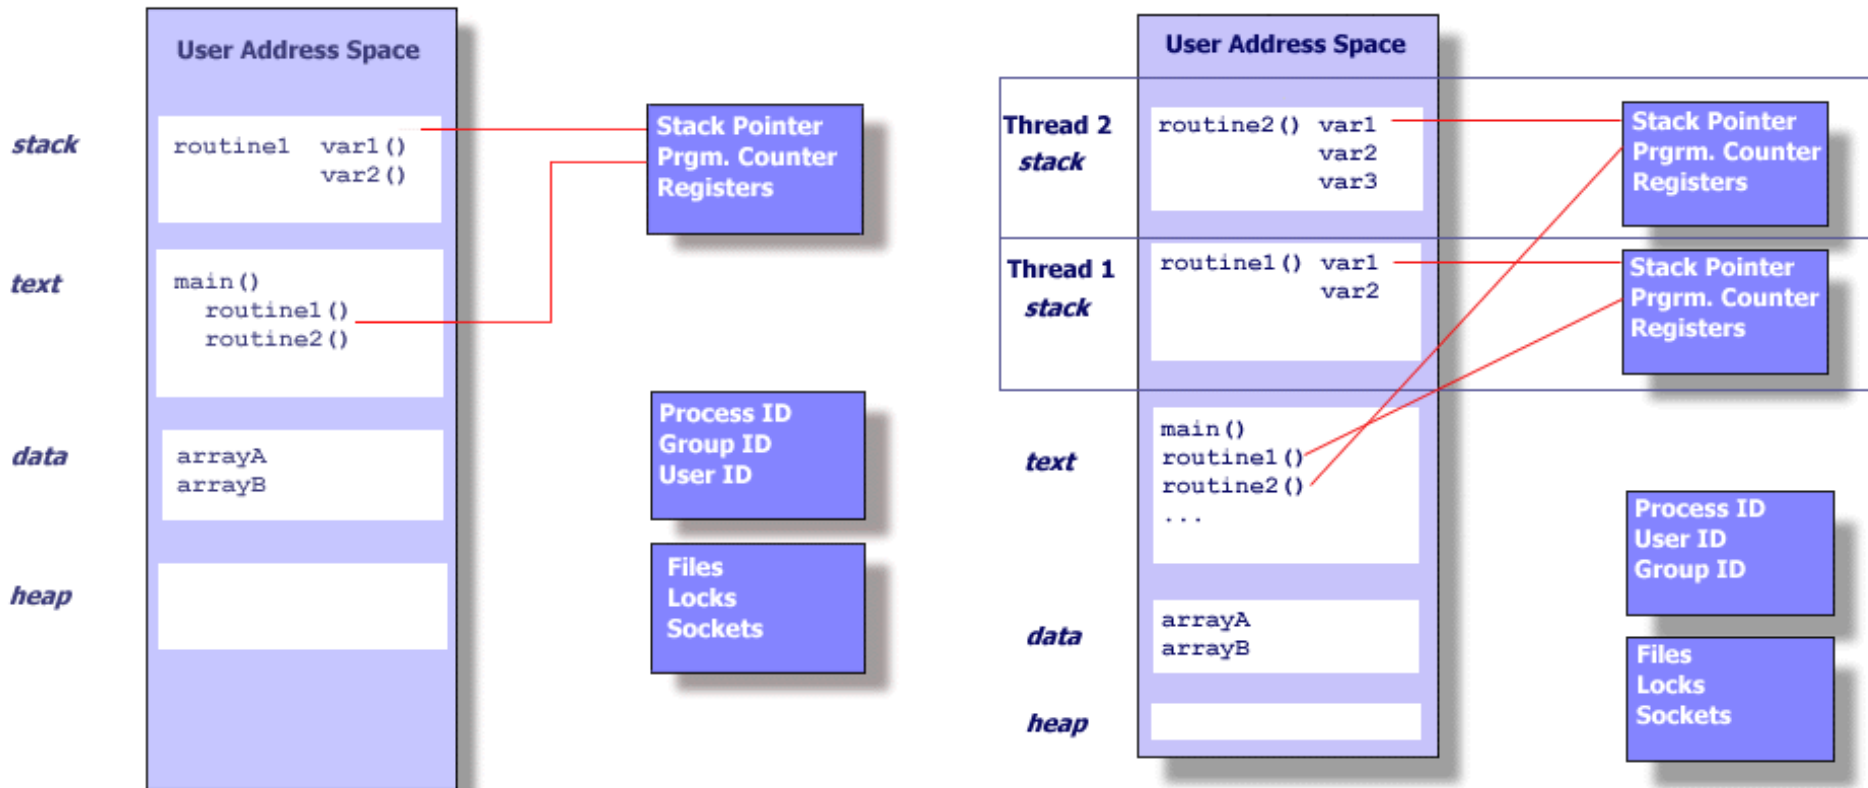
- Most operating systems therefore support two entities:
  - the process,
    - which defines the address space and general process attributes
  - the thread,
    - which defines a sequential execution stream within a process
- A thread is bound to a single process.
  - For each process, however, there may be many threads.
- Threads are the unit of scheduling
- Processes are *containers* in which threads execute

# Multithreaded Processes

---



# Single vs Multithreaded Processes



# Threads vs. Processes

---

- A thread has no separate code/data segment or heap
- A thread cannot live on its own, it must live within a process
- *There can be more than one thread in a process, the first thread calls main & has the process's stack*
- Inexpensive creation
- Inexpensive context switching
- If a thread dies, its stack is reclaimed
- A process has code/data/heap & other segments
- There must be at least one thread in a process
- *Threads within a process share code/data/heap, share I/O, but each has its own stack & registers*
- Expensive creation
- Expensive context switching
- If a process dies, its resources are reclaimed & all threads die

# Separating Threads and Processes

---

- Makes it easier to support multithreaded applications
  - Different from multiprocessing, multiprogramming
  - Multiprocessing  $\equiv$  Multiple CPUs, parallel job execution
  - Multiprogramming  $\equiv$  Multiple Jobs or Processes
  - Multithreading  $\equiv$  Multiple threads per Process
- Concurrency (multithreading) is useful for:
  - improving program structure
  - handling concurrent events (e.g., web requests)
  - building parallel programs
  - Resource sharing
  - Multiprocessor utilization
- Is multithreading useful even on a single processor?

# Benefits of multithreaded programs

---

- **Responsiveness**
  - Interactive apps like Web server
- **Resource sharing**
  - Implicit as threads share the same address space
- **Economy**
  - Creation of multithread process is cheaper
  - Solaris
    - Process creation is 30 times slower than that of thread creation
    - Context switch is about 5 times slower
- **Scalability**
  - Single-threaded process can only run on one CPU
  - Multithreading in multi-core m/cs increases parallelism

# Examples of multithreaded programs

---

- Embedded systems
  - Elevators, Planes, Medical systems
  - Single Program, concurrent operations
- Most modern OS kernels
  - Internally concurrent because have to deal with concurrent requests by multiple users
  - Threads for I/O devices, interrupt handling, managing amount of free memory, etc
  - But no protection needed within kernel
- Database Servers
  - Access to shared data by many concurrent users
  - Also background utility processing must be done



# Single-Threaded Example

---

- Consider the following *CHESS* program:

```
main() {  
    DisplayChessBoad();  
    ComputeNextMove();  
}
```

- What is the behavior here?
- Version of program with Threads:

```
main() {  
    CreateThread(DisplayChessBoad());  
    CreateThread(ComputeNextMove());  
}
```

# Single-Threaded Example

---

- Consider the following C program:

```
main() {  
    ComputePI("pi.txt");  
    PrintClassList("clist.text");  
}
```

- What is the behavior here?
  - Program would never print out class list
  - Why?
    - ComputePI would never finish
  - Solution?
    - (w/o using threads?)

# Use of Threads

---

- Version of program with Threads:

```
main() {  
    CreateThread(ComputePI("pi.txt"));  
    CreateThread(PrintClassList("clist.text"));  
}
```

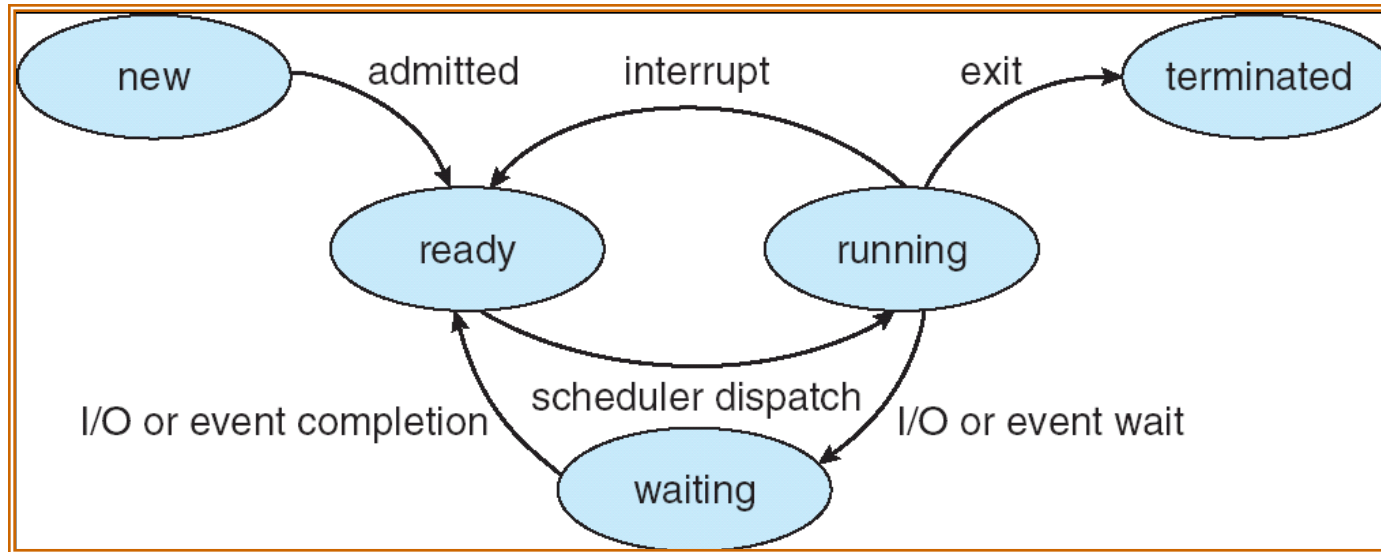
- What does "CreateThread" do?
  - Start independent thread running given procedure
- What is the behavior here?
  - Now, you would actually see the class list
  - This *should* behave as if there are two separate CPUs



Time →

# Lifecycle of a Thread (or Process)

---



- As a thread executes, it changes state:
  - **new**: The thread is being created
  - **ready**: The thread is waiting to run
  - **running**: Instructions are being executed
  - **waiting**: Thread waiting for some event to occur
  - **terminated**: The thread has finished execution
- “Active” threads are represented by their TCBs
  - TCBs organized into queues based on their state

# Cooperative Threads

---

*A cooperative thread runs until it decides to give up the CPU*

```
main()
```

```
{
```

```
    tid t1 = CreateThread(fn, arg);
```

```
    ...
```

```
    Yield(t1);
```

```
}
```

```
fn(int arg)
```

```
{
```

```
    ...
```

```
    Yield(any);
```

```
}
```

# Cooperative Threads

---

- Cooperative threads use **non pre-emptive** scheduling
- Advantages:
  - Simple
    - Scientific apps
- Disadvantages:
  - For badly written code
- Scheduler gets invoked only when Yield is called
- A thread could also yield the processor when it blocks for I/O

# Non-Cooperative Threads

---

- No explicit control passing among threads
- Rely on a scheduler to decide which thread to run
- A thread can be pre-empted at any point
- Often called **pre-emptive** threads
- Most modern thread packages use this approach
  - Pthreads API
  - Win32 threads API
  - Java API

# Classification of OS

# threads Per AS: # of addr spaces:	One	Many
One	MS/DOS, early Macintosh	Traditional UNIX
Many	Embedded systems (Geoworks, VxWorks, JavaOS, etc) JavaOS, Pilot(PC)	Mach, OS/2, Linux Windows 10, Solaris, HP-UX, OS X

- Most operating systems have either
  - One or many address spaces
  - One or many threads per address space



## Example User Thread Interface

---

`t = thread_fork(initial context)`

creates a new thread of control

`thread_stop()`

stops the calling thread, sometimes called `thread_block`

`thread_start(t)`

starts the named thread

`thread_yield()`

voluntarily gives up the processor


`thread_exit()`

terminates the calling thread, sometimes called  
`thread_destroy`

# Multithreading models

---

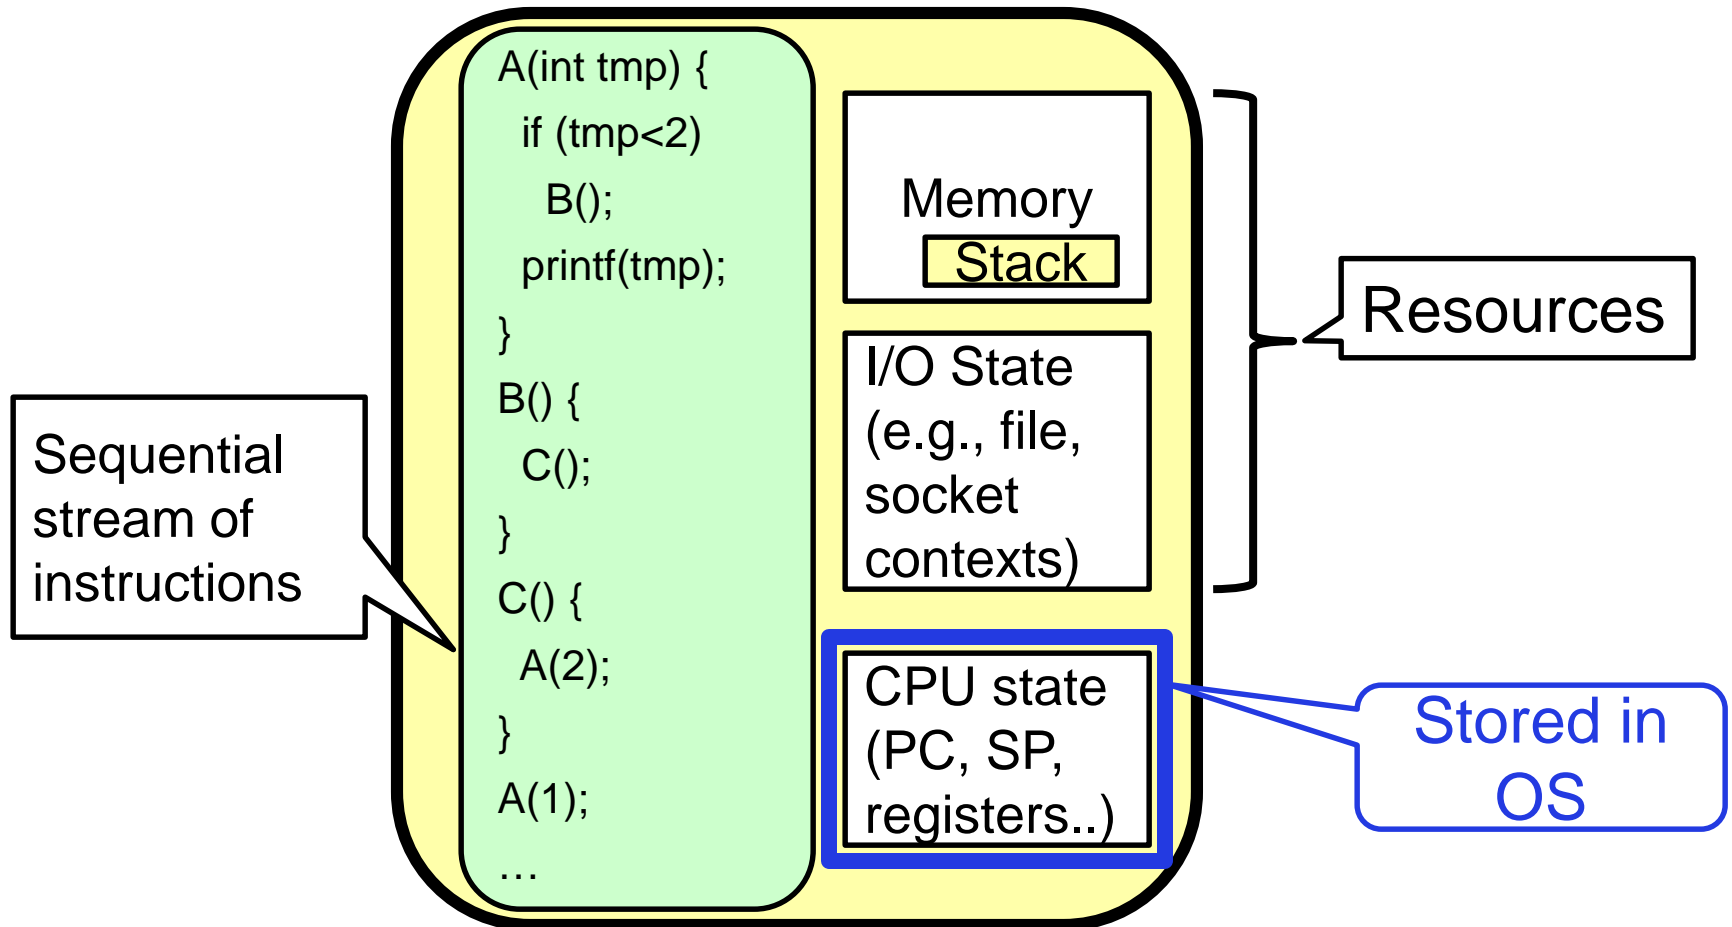
- There are actually 2 level of threads:
- Kernel threads:
  - Supported and managed directly by the kernel.
  - Windows, Mac, Linux
- User threads:
  - Supported above the kernel, and without kernel knowledge by user-level threads library.
  - E.g., **POSIX Pthreads API**



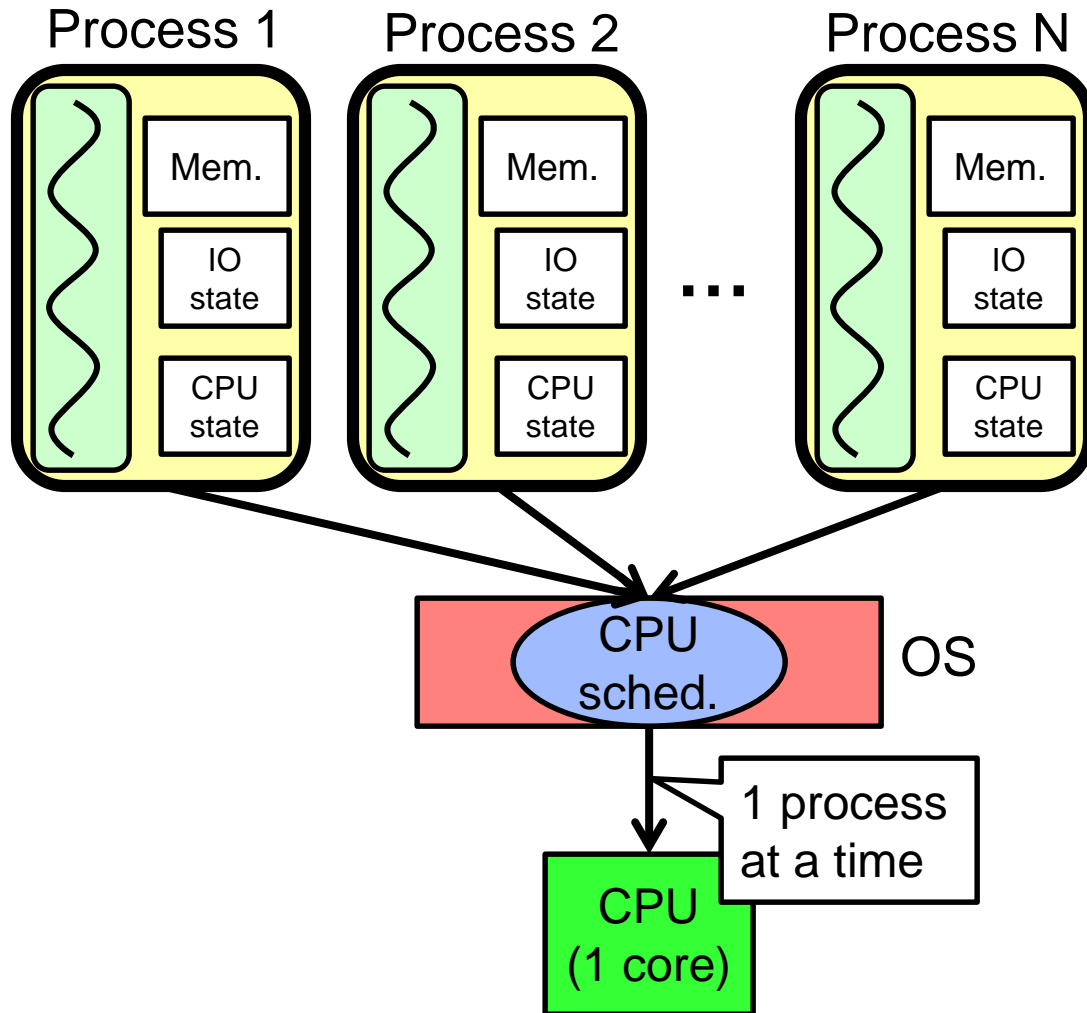
1:1 mapping b/w user and kernel threads in Windows & Linux

# Putting it Together: Process

## (Unix) Process

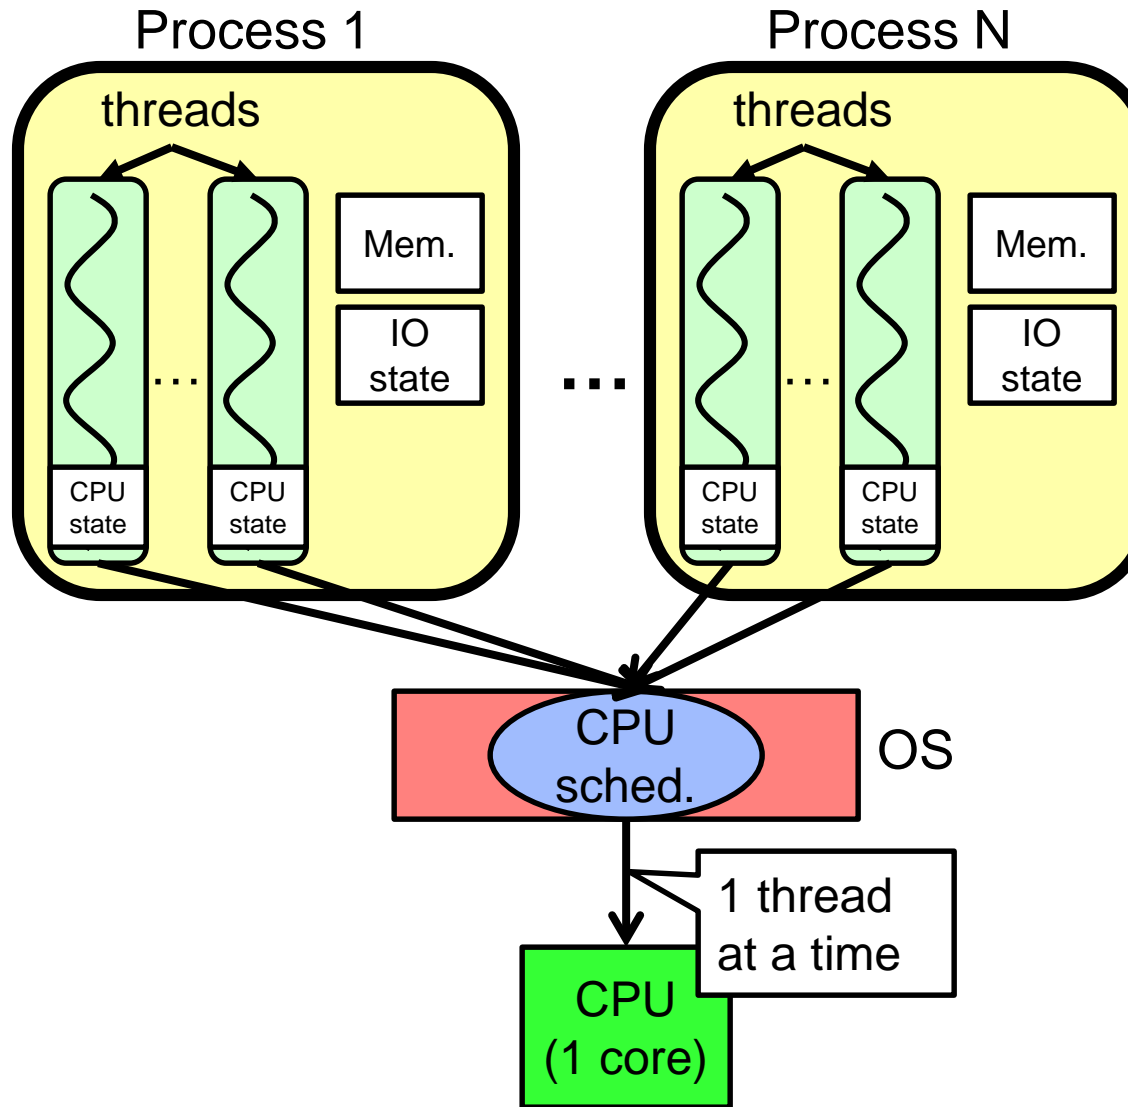


# Putting it Together: Processes



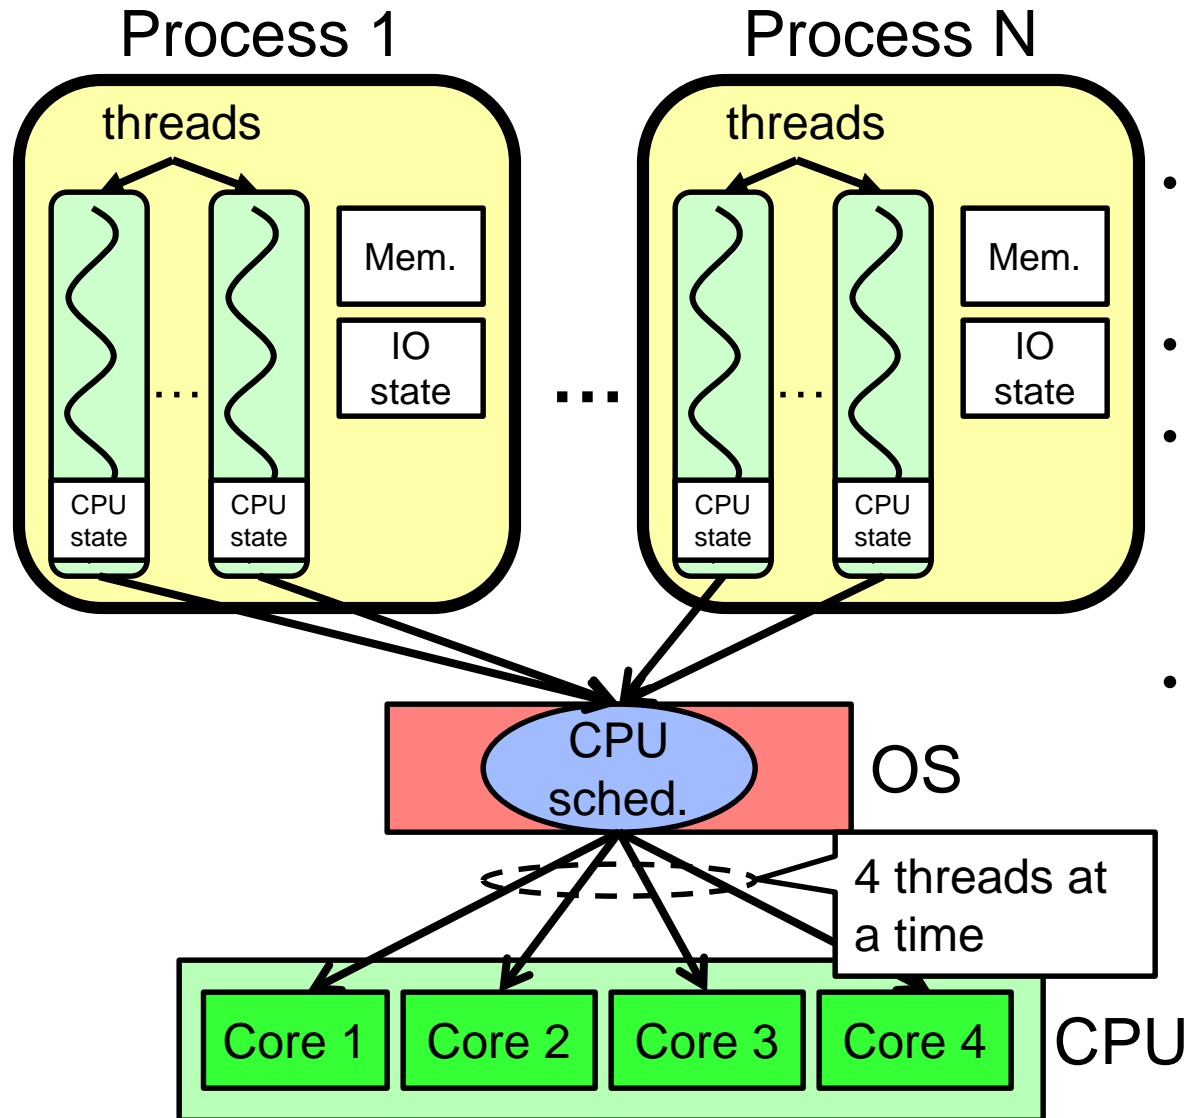
- Switch overhead:  
**high**
  - CPU state: **low**
  - Memory/IO state: **high**
- Process creation: **high**
- Protection
  - CPU: **yes**
  - Memory/IO: **yes**
- Sharing overhead: **high** (involves at least a context switch)

# Putting it Together: Threads



- **Switch overhead: medium**
  - CPU state: *low*
- **Thread creation: medium**
- **Protection**
  - CPU: *yes*
  - Memory/IO: *no*
- **Sharing overhead: *low(ish)* (thread switch overhead low)**

# Putting it Together: Multi-Cores



- Switch overhead: **low** (only CPU state)
- Thread creation: **low**
- Protection
  - CPU: **yes**
  - Memory/IO: **No**
- Sharing overhead: **low** (thread switch overhead low, may not need to switch at all!)

# Pthreads

---

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- *Specification, not implementation*
- API specifies behavior of the thread library, implementation is up to development of the library
- May be provided either as user-level or kernel-level threads
- Common in UNIX operating systems (Linux & Mac OS X)
  - Implemented by glibc

# Pthreads API

---

- `pthread_create` (thread id, attr, start\_routine, arguments\_start\_routine)
- `pthread_exit` (status)
- `pthread_cancel` (thread id)
- `pthread_attr_init` (attr)
- `pthread_attr_destroy` (attr)
- `pthread_join` (thread id, status)
- `pthread_detach` (thread id)
- `pthread_attr_setdetachstate` (attr, detachstate)
- `pthread_attr_getdetachstate` (attr, detachstate)
- `pthread_self` (), `pthread_equal` (TID1, TID2)



# Pthreads: Creating Threads

---

- **main()** program comprises a single, default thread
  - All other threads must be explicitly created using **pthread\_create**, anywhere in the program
- By default, thread is created with certain attributes (configurable)
  - `pthread_attr_init` and `pthread_attr_destroy` are used to initialize/destroy thread attribute object
  - Other routines are then used to query/set specific attributes in the thread attribute object
  - Detached or joinable state
  - Scheduling policy, Scheduling parameters
  - Stack address, Stack size, etc

# Pthread API: Creation and Termination

---

```
#include <pthread.h> //Implemented by glibc
#include <stdio.h>
#define NUM_THREADS    5
void *PrintHello(void *threadNo)
{
    long tid;
    tid = (long*)threadNo;
    printf("Hello World! It's me, thread #%ld!\n", tid);
    pthread_exit(void *threadNo); //can pass status to other threads
}
```

→ Use gcc & g++ with -pthread flag

→ getrlimit: get user limits of system resources like memory, no of processes, timeslice, etc

→ \$ulimit: get and set user limits of system resources

→ \$cat /proc/[PID]/limits

→ \$sudo cat /proc/[PID]/sched

# Pthread API: Creation and Termination

---

```
int main ( )
{
    pthread_t threads[NUM_THREADS];
    int rc;
    long t; void * status;
    for(t=0; t<NUM_THREADS; t++){
        printf("In main: creating thread %ld\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)&t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n",
rc);
            exit(-1);
        }
    }
    /* main() should wait for thread(s) to finish */
    for(t=0; t<NUM_THREADS; t++)
        pthread_join(threads[t],&status); //collect status
    pthread_exit(NULL); //makes
}
```

# Pthread API: Thread Argument Passing (safeway)

```
long *taskids[NUM_THREADS];

for(t=0; t<NUM_THREADS; t++)
{
    taskids[t] = (long *) malloc(sizeof(long));
    *taskids[t] = t;
    printf("Creating thread %ld\n", t);
    rc = pthread_create(&threads[t], NULL, PrintHello, (void *)
taskids[t]);
    ...
}
```

Full Source Code:

[https://computing.llnl.gov/tutorials/pthreads/samples/hello\\_arg1.c](https://computing.llnl.gov/tutorials/pthreads/samples/hello_arg1.c)

# Pthread API: Thread Argument Passing (Unsafe)

---

```
int rc;
```

```
long t;
```

```
for(t=0; t<NUM_THREADS; t++)
```

```
{
```

```
    printf("Creating thread %ld\n", t);
```

```
    rc = pthread_create(&threads[t], NULL, PrintHello, (void *) &t);
```

```
    ...
```

```
}
```

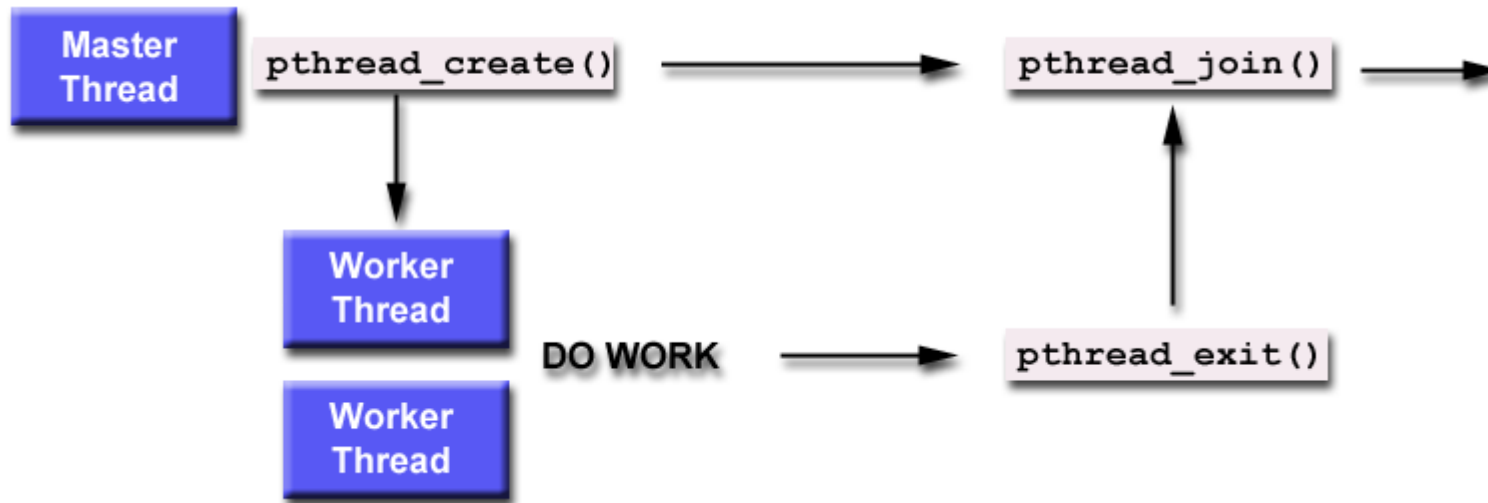
- Param **t** is changed by the main thread as it creates new threads

Full Src Code:

[https://computing.llnl.gov/tutorials/pthreads/samples/hello\\_arg3.c](https://computing.llnl.gov/tutorials/pthreads/samples/hello_arg3.c)

# Pthread API: Thread Join/Detach

---



- “joining” is one of the ways to accomplish synchronization between threads
- Calling `pthread_exit()` at last in `main( )` blocks the process till all its threads are done!
- Example:

<https://computing.llnl.gov/tutorials/pthreads/samples/join.c>

# Pthread: Issues

---

After a thread has been created, how do you know

a) when it will be scheduled to run by the OS

b) which processor/core it will run on?

Ans:

1. Depends on underlying thread scheduling algo (FIFO/RO/etc for pthreads) or
2. Implementation specific

Robust programs should not depend on threads running order or core on which a thread runs on

# Linux Threads

---

- Linux does not distinguish between processes and threads
  - Uses term **task** (**struct task\_struct**)
  - clone ( ) for creating threads
    - Flags passed as args determine level of sharing b/w parent and child tasks
    - CLONE\_FS, CLONE\_VM, CLONDE\_FILES
    - Sharing → threads
    - No sharing → processes
  - fork( ) for creating duplicate tasks (processes)



# Multithreading Issues

---

- Semantics of `fork()` and `exec()` system calls
  - Child process duplicates all threads of parent?
  - Two versions of `fork()`!!
  - `exec()` inside a thread will replace the entire process (inc all threads) with prg specified as arg for `exec()`
- Thread cancellation
  - Asynchronous vs. Deferred Cancellation
  - `pthread_cancel` (thread id) supports both, but deferred is recommended as it's safe
- Signal handling
  - Which thread to deliver it to?
  - `kill(pid, signal)`
  - `pthread_kill(tid, signal)`

# Thread Hazards

---

```
int a = 1, b = 2, w = 2;
```

```
main( ) {  
    CreateThread(fn, 4);  
    CreateThread(fn, 4);  
    while(w) ;  
}  
  
fn( ) {  
    int v = a + b;  
    w--;  
}
```

# Concurrency Problems

---

A statement like `w--` in C (or C++) is implemented by several machine instructions:

```
ld    r4, #w
add   r4, r4, -1
st    r4, #w
```

Now, imagine the following sequence, what is the value of `w`?

```
ld    r4, #w
_____
_____
_____
add   r4, r4, -1
st    r4, #w
```

```
_____
ld    r4, #w
add   r4, r4, -1
st    r4, #w
```

# Implicit Threading

---

- Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads
- Implicit Threading: Creation and management of threads done by compilers and run-time libraries rather than programmers!
- Three methods
  - Thread Pools
  - OpenMP
  - Grand Central Dispatch
- Other methods include Microsoft Threading Building Blocks (TBB), `java.util.concurrent` package

# Thread Pools

---

- Create a number of threads in a pool where they await work
- Advantages:
  - Usually slightly faster to service a request with an existing thread than create a new thread
  - Allows the number of threads in the application(s) to be bound to the size of the pool
  - Separating task to be performed from mechanics of creating task allows different strategies for running task
    - i.e., Tasks could be scheduled to run periodically
- Windows API supports thread pools:

```
DWORD WINAPI PoolFunction(AVOID Param) {  
    /*  
    * this function runs as a separate thread.  
    */  
}
```

# OpenMP

---

- Set of compiler directives and an API for C, C++, FORTRAN
- Provides support for parallel programming in shared-memory environments
- Identifies **parallel regions** - blocks of code that can run in parallel

**#pragma omp parallel**

Create as many threads as there are cores!

**#pragma omp parallel for**

```
for(i=0;i<N;i++) {  
    c[i] = a[i] + b[i];
```

```
}
```

Run **for loop** in parallel

```
#include <omp.h>  
#include <stdio.h>  
  
int main(int argc, char *argv[])  
{  
    /* sequential code */  
  
    #pragma omp parallel  
    {  
        printf("I am a parallel region.");  
    }  
  
    /* sequential code */  
  
    return 0;  
}
```

# Grand Central Dispatch

---

- Apple technology for Mac OS X and iOS operating systems
- Extensions to C, C++ languages, API, and run-time library
- Allows identification of parallel sections
- Manages most of the details of threading
- Block is in “`^{ }`” - `^{ printf("I am a block"); }`
- Blocks placed in dispatch queue
  - Assigned to available thread in thread pool when removed from queue
- Two types of dispatch queues:
  - serial - blocks removed in FIFO order, queue is per process, called **main queue**
  - concurrent - removed in FIFO order but several may be removed at a time
    - Three system wide queues with priorities low, default, high

```
dispatch_queue_t queue = dispatch_get_global_queue  
    (DISPATCH_QUEUE_PRIORITY_DEFAULT, 0);
```

```
dispatch_async(queue, ^{ printf("I am a block."); });
```

# Summary

---

- Threads increase concurrency/parallelism
- Threads may cause synchronization issues
  - Need to employ synchronization primitives to avoid thread hazards



# Reading Assignment

---

- Chapter 4 from OSC by Galvin et al
- Chapter 2 from MOS by Tanenbaum et al
- <http://www.yolinux.com/TUTORIALS/LinuxTutorialPosixThreads.html>
- <https://computing.llnl.gov/tutorials/pthreads/>