

Operating Systems–II: CS3523

Spring 2020

Lab Assignment 2: Linux Kernel Debugger and Process Management

Last date for submission: N/A

Setup:

- a. Install virt-manager on your laptop.
- b. Import or Install CentOS 7:

Import:

- i. Download pre-installed CentOS 7 with kernel debugger setup from :

<http://jenkins.cse.iith.ac.in:8080/userContent/cs3523-1.qcow2>

- ii. Import the image into a VM

```
$ virt-install --import --memory 4096 --vcpus 2 --name cs3523-1 --disk  
./cs3523-1.qcow2 --os-variant rhel7
```

Install

- i. Download CentOS 7 ISO image from:
http://intranet.iith.ac.in/files/os/CentOS-7-x86_64-DVD-1908.iso
 - ii. Create a virtual machine with 4 GiB RAM, 10GB HDD (dynamic vdi), 2 CPU, 1 network cards (NAT)
 - iii. Install Centos 7 64-bit on it (minimal centos / no GUI, enable kdump) . Check the kernel version.
 - iv. Yum install crash, gcc, perl, strace rpms.

```
# yum install -y strace crash vim gcc perl
```
 - v. Update all packages

```
# yum update
```
 - vi. Reboot
 - vii. Install kernel-debuginfo and kernel-debuginfo-common packages

```
# yum --enablerepo=base-debuginfo install -y kernel-debuginfo-$(uname -r)
```
- c. You should be able to login via ssh to the IP address of the VM.
User: root, Passwd: r123
 - d. Do your experiments in the VM.

NOTE: If you don't use VM setup, your laptop may hang, crash, and go bad. Use centos in VM, as Ubuntu is not suitable for kernel programmers.

Problem Statement:

Background:

- Linux maintains separate process descriptor (struct task) for each process and thread.

Goal: The goal of this assignment is to task structures of processes using kernel debugger.

Details: On the virtual machine running centos, run the following experiments.

Experiment #1:

- Collect kernel dump of the running system using “echo c > /proc/sysrq-trigger”
- The system reboots and dump is captured in a time-stamped directory under /var/crash/
- Launch kernel debugger “crash” on the kdump collected using

```
# cd /var/crash/<dumpdirectory>
# crash /usr/lib/debug/usr/lib/modules/`uname -r`/vmlinuz ./vmcore
```

```
crash> help

*          extend      log          rd          task
alias      files       mach       repeat     timer
ascii     foreach     mod       runq       tree
bpf        fuser       mount     search     union
bt         gdb         net       set        vm
btop       help        p         sig        vtop
dev        ipcs        ps        struct     waitq
dis        irq         ptb       swap       whatis
eval       kmem        ptov      sym        wr
exit       list        sys       sys        q

crash version: 7.2.3-10.el7  gdb version: 7.6
For help on any command above, enter "help <command>".
For help on input options, enter "help input".
For help on output options, enter "help output".
```

- Try out a few commands like “log”, “bt”, “dev”, “p”, “ipcs”
- Quit from kernel debugger

```
crash> quit
```

Experiment #2:

- Run some processes in background
- Collect kernel dump of the running system using “echo c > /proc/sysrq-trigger”
- The system reboots and dump is captured in a time-stamped directory under /var/crash/
- Launch kernel debugger “crash” on the kdump collected using

```
# cd /var/crash/<dumpdirectory>
# crash /usr/lib/debug/usr/lib/modules/`uname -r`/vmlinuz ./vmcore
```

- Try out commands “ps”, “runq”, “task”, “set”

```
crash> help ps
```

```
crash> ps
```

	PID	PPID	CPU	TASK	ST	%MEM	VSZ	RSS	COMM
	0	0	0	ffffffffffb0e18480	RU	0.0	0	0	[swapper/0]
>	0	0	1	ffff9b0cb94941c0	RU	0.0	0	0	[swapper/1]

```

1      0  0  ffff9b0cb9450000  IN  0.1  127940  6568  systemd
2      0  1  ffff9b0cb9451070  IN  0.0  0      0  [kthreadd]
4      2  0  ffff9b0cb9453150  IN  0.0  0      0  [kworker/0:0H]
...
> 1642  811  0  ffff9b0bb60141c0  RU  0.0  115448  2100  bash
1828   1  1  ffff9b0c3778b150  IN  0.0  39084  1540  systemd-journal
1833   2  1  ffff9b0cb5c641c0  IN  0.0  0      0  [kworker/u4:0]
1844   1  1  ffff9b0bb66362a0  IN  0.0  26380  1772  systemd-logind
1946  1469  1  ffff9b0c376fc1c0  IN  0.1  89804  4068  pickup
1968   2  1  ffff9b0c376f9070  IN  0.0  0      0  [kworker/1:0]
1989   2  0  ffff9b0cb720a0e0  RU  0.0  0      0  [kworker/0:1]
1992   2  1  ffff9b0cb5d041c0  IN  0.0  0      0  [kworker/1:2]
1993   2  1  ffff9b0cb5d062a0  IN  0.0  0      0  [kworker/1:1]
2015  1642  1  ffff9b0c376fe2a0  ST  0.0  107976  600  cat
2016  1642  1  ffff9b0c376fa0e0  ST  0.0  107976  600  cat
2017  1642  0  ffff9b0c376fb150  ST  0.0  159364  1784  top
2018  1642  0  ffff9b0c376f8000  ST  0.1  149004  4956  vim
2021   2  1  ffff9b0cb94541c0  IN  0.0  0      0  [kworker/1:3]

```

crash> help runq

crash> runq

CPU 0 RUNQUEUE: ffff9b0cbfc1ac80

CURRENT: PID: 1642 TASK: ffff9b0bb60141c0 COMMAND: "bash"

RT_PRIO_ARRAY: ffff9b0cbfc1ae20

[no tasks queued]

CFS_RB_ROOT: ffff9b0cbfc1ad28

[120] PID: 1989 TASK: ffff9b0cb720a0e0 COMMAND: "kworker/0:1"

CPU 1 RUNQUEUE: ffff9b0cbfd1ac80

CURRENT: PID: 0 TASK: ffff9b0cb94941c0 COMMAND: "swapper/1"

RT_PRIO_ARRAY: ffff9b0cbfd1ae20

[no tasks queued]

CFS_RB_ROOT: ffff9b0cbfd1ad28

[no tasks queued]

crash> runq -m

CPU 0: [0 00:00:00.001] PID: 1642 TASK: ffff9b0bb60141c0 COMMAND: "bash"

CPU 1: [0 05:46:06.069] PID: 0 TASK: ffff9b0cb94941c0 COMMAND: "swapper/1"

crash> runq -t

CPU 0: 20766064292679

20766063199929 PID: 1642 TASK: ffff9b0bb60141c0 COMMAND: "bash"

CPU 1: 20766069904525

0000000000000000 PID: 0 TASK: ffff9b0cb94941c0 COMMAND: "swapper/1"

crash> runq -g

CPU 0

CURRENT: PID: 1642 TASK: ffff9b0bb60141c0 COMMAND: "bash"

ROOT_TASK_GROUP: ffffffff1307c00 RT_RQ: ffff9b0cbfc1ae20

[no tasks queued]

```
ROOT_TASK_GROUP: ffffffff1307c00 CFS_RQ: ffff9b0cbfc1ad00
[120] PID: 1642 TASK: ffff9b0bb60141c0 COMMAND: "bash" [CURRENT]
[120] PID: 1989 TASK: ffff9b0cb720a0e0 COMMAND: "kworker/0:1"
```

CPU 1

```
CURRENT: PID: 0 TASK: ffff9b0cb94941c0 COMMAND: "swapper/1"
ROOT_TASK_GROUP: ffffffff1307c00 RT_RQ: ffff9b0cbfd1ae20
[no tasks queued]
ROOT_TASK_GROUP: ffffffff1307c00 CFS_RQ: ffff9b0cbfd1ad00
[no tasks queued]
```

crash> help task

crash> whatis task_struct

```
crash> whatis wait_queue_head_t
typedef struct __wait_queue_head {
    spinlock_t lock;
    struct list_head task_list;
} wait_queue_head_t;
SIZE: 24
```

crash> waitq <pointer>

f. Quit from kernel debugger

crash> quit