

Concurrency and Process Synchronization: Synchronizing Access to Shared Objects

Outline

- Background
- Classical Producer-Consumer Problem
 - Basic solutions
- Race Conditions
- The Critical-Section Problem
- Peterson's Solution, Bakery Solution
- Hardware Support
- Mutex locks
- Semaphores
- Solution to P-C problem using Semaphores

Thread vs Process State

- Process-wide state:
 - Memory contents (global variables, heap)
 - I/O bookkeeping
 - Kept in **Process Control Block (PCB)**
- Thread-"local" state:
 - CPU registers including program counter (PC)
 - Execution stack
 - Kept in **Thread Control Block (TCB)**

Shared vs Per-Thread State


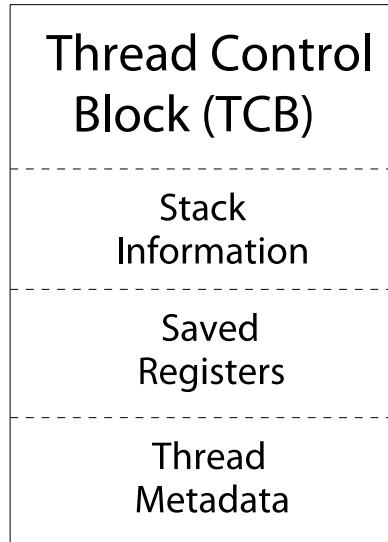
Shared State

Heap


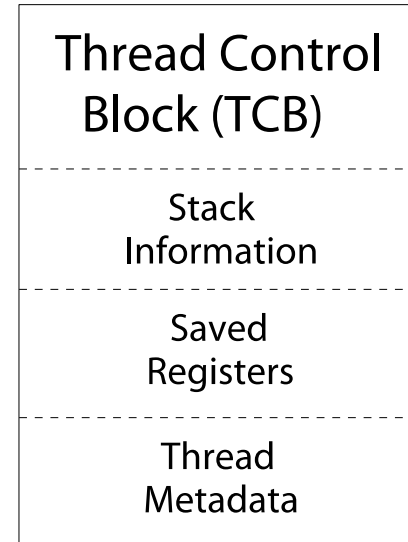
Global Variables

Code

Per-Thread State



Per-Thread State




Programmer vs. Processor View


Programmer's View	Possible Execution #1
----------------------	-----------------------------

.	.
.	.
.	.
$x = x + 1;$	$x = x + 1;$
$y = y + x;$	$y = y + x;$
$z = x + 5y;$	$z = x + 5y;$
.	.
.	.
.	.


Possible Executions



a) One execution



b) Another execution



c) Another execution

Correctness with Concurrent Threads

- Non-determinism:
 - Scheduler can run threads in **any order**
 - Scheduler can switch threads **at any time**
 - This can make testing very difficult
- *Independent Threads*
 - No state shared with other threads
 - Deterministic, reproducible conditions
- *Cooperating Threads*
 - Shared state between multiple threads
- **Goal: Correctness by Design**

Shared Objects

- Concurrent access to shared (data) objects may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes/threads
- First look at some toy examples to illustrate thread hazards
- Then look at the **classical consumer-producer problem**
- *Producer* process produces information consumed by *Consumer* process
 - Very common paradigm for cooperating processes/threads
 - Shared memory (IPC) is used in case of processes
 - Heap/Global variables are used in case of threads

Race Conditions

- What are the possible final values of x below?

Thread A
 $x = 1;$

Thread B
 $x = 2;$

- It can be $x = 1$ or 2 depending on which thread wins or loses the race to set x → **Race condition**
- Definition: *a timing dependent error involving shared state*
 - Whether it happens depends on how threads scheduled or interleaved while accessing/manipulating shared objects
 - In effect, once thread A starts doing something, it needs to “race” to finish it because if thread B looks at the shared object before A is done, it may see something inconsistent

Race Conditions

- What are the possible values of x below?
- Initially $y = 12$;

Thread A

$x = y + 1;$

Thread B

$y = y * 2;$

- 13 or 25 (non-deterministic)
- Race Condition

Two threads, one counter

Popular web server


- Uses multiple threads to speed things up.
- Simple shared state error:
 - each thread increments a shared counter to track number of hits

```
...  
hits = hits + 1;  
...
```

- What happens when two threads execute concurrently?

Shared counters


- Possible result: lost update!



- One other possible result: everything works.
⇒ Difficult to debug
- Race condition

Producer-Consumer Problem


- Start by imagining an unbounded (**infinite**) buffer
- Producer process writes data to buffer
 - Writes to **In** and moves rightwards
- Consumer process reads data from buffer
 - Reads from **Out** and moves rightwards
 - Should not try to consume if there is no data



Need an infinite buffer

Producer-Consumer Problem

- Bounded buffer: size 'N'
 - Access entry 0... N-1, then “wrap around” to 0 again
- Producer process writes data to buffer
 - Must not write more than 'N' items more than consumer “ate”
- Consumer process reads data from buffer
 - Should not try to consume if there is no data



Solution #1: Producer/Consumer Problem

Bounded Buffer Case:

```
#define BUFFER_SIZE 10

typedef struct{
    ..some stuff..
}item;

item buffer[BUFFER_SIZE];
int in = 0
int out = 0;
```

Threads: Keep above variables in Data segment;
Processes: Use shared memory;

Solution #1: Producer/Consumer (1/2)

- Producer process/thread:

```
item nextProduced;

while(true)
{
    /*Produce an item in nextProduced*/
    while(((in + 1) % BUFFER_SIZE) == out)
        continue; //do nothing...

    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
}
```


Solution #1: Producer/Consumer (2/2)

- Consumer process/thread:

```
item nextConsumed;

while(true)
{
    while(in == out)
        continue; //do nothing..

    /* Consume item in nextConsumed */
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
}
```

Solution #1: Producer/Consumer

Producer:

```
item nextProduced;

while(true)
{
    while(((in+1)%BUFFER_SIZE)==out)
        continue; //do nothing...

    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Consumer:

```
item nextConsumed;

while(true)
{
    while(in==out)
        continue; //do nothing..
    nextConsumed=buffer[out];
    out=(out+1)%BUFFER_SIZE;
}
```

Solution #1: Analysis

- Solution #1 does not suffer from any data inconsistency issues for 1 producer and 1 consumer scenario
 - Reason: **In** is manipulated (**write operation**) only by Producer
 - Reason: **Out** is manipulated only by Consumer
- But, it fails in the case of multiple producers and/or multiple consumers scenarios
 - Reason: **In** is now manipulated by more than one Producer, similarly **Out** is by more than one Consumer process/thread
- In any case, solution #1 only allows BUFFER_SIZE-1 items filled-up at the same time
- To remedy this, the processes would need to **synchronize** their access to the shared buffer

Solution #2: Producer/Consumer Problem

- Suppose that we wanted to provide a solution to the consumer-producer problem that fills in **all** the buffer items.
 - Assume an integer **count** keeps track of the number of full buffers.
 - Initially, **count** is set to 0.
 - It is incremented by the producer after it produces a new buffer item
 - It is decremented by the consumer after it consumes a buffer item
- Any issues with above solution?
 - Scenario 1: 1 Producer and 1 Consumer?
 - Scenario 2: More than 1 Producer and/or More than 1 Consumer?
 - Same variable is being manipulated by more than one process/thread which could be either Producer/Consumer.
 - Recall thread hazards discussed earlier in OS-1 course!

Sol #2: Producer-Consumer

- **Producer**

```
while (true) {  
  
    /* produce an item and */  
    /* put in nextProduced */  
  
    while (count == BUFFER_SIZE)  
        ; // do nothing b/c full  
  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

- **Consumer**

```
while (true) {  
  
    while (count == 0)  
        ; // do nothing b/c empty  
  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
  
    /* consume the item */  
    /* in nextConsumed */  
}
```

Sol #2: Analysis

- **count++** *not* atomic operation. Could be implemented as

```
register1 = count
register1 = register1 + 1
count = register1
```

- **count--** *not* atomic operation. Could be implemented as

```
register2 = count
register2 = register2 - 1
count = register2
```

- Consider this execution interleaving with “count = 5” initially:
 - S0: producer executes **register1 = count** {register1 = 5}
 - S1: producer executes **register1 = register1 + 1** {register1 = 6}
 - S2: consumer executes **register2 = count** {register2 = 5}
 - S3: consumer executes **register2 = register2 - 1** {register2 = 4}
 - S4: producer executes **count = register1** {count = 6}
 - S5: consumer executes **count = register2** {count = 4}

What just happened?

- Incorrect state of “count” as both processes manipulating this shared variable concurrently
- This is an example of **race condition**: When multiple jobs share & manipulate the same data concurrently, the outcome depends on the particular order in which the access takes place for the **shared data (critical section)**
- Threads share global memory (data segment) and many other except stack segment
 - This can result in *race conditions*
- Race conditions can arise even for processes that use IPC to coordinate
- To prevent race conditions, **concurrent tasks must be synchronized** to ensure that **only one task at a time** can manipulate **shared structures/variables in critical section**

Race conditions

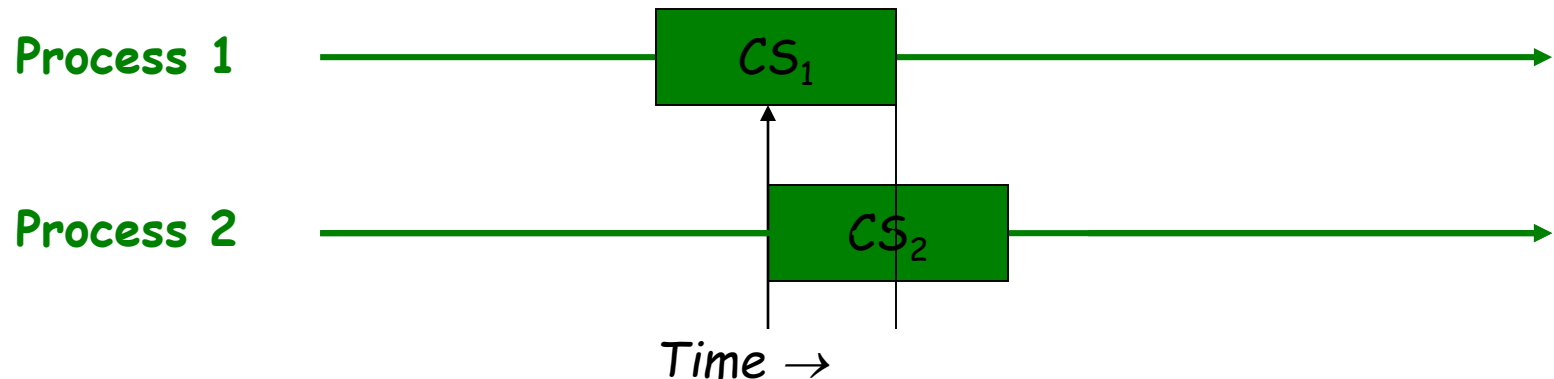
- Hard to detect:
 - All possible schedules/interleavings have to be safe
 - Number of possible schedule permutations is huge
 - Some bad schedules? Some that will work sometimes?
 - they are intermittent
 - Timing dependent = small changes can hide bug
 - Debugging is very hard as errors due to race conditions might surface very rarely
 - Celebrate if bug is deterministic and repeatable!

Race Conditions in Kernel

- OS is subject to several possible race conditions as it is implemented as multi-threaded process
- E.g., Kernel Data Structures like
 - List of Open files
 - Process table
 - DS for maintaining memory allocation
 - DS for interrupt handling
- Kernel developers have to ensure that OS is free from race conditions!
- Two general approaches to handle CSs in OS
 - Non-preemptive kernels → only one job is active in kernel at a time
 - Preemptive kernels → high priority job can preempt low priority one during its execution in kernel
 - More responsive
 - Needed for RT scheduling of jobs

Critical Section Problem (CSP)

- Problem: Design a protocol for processes to cooperate, such that only one process is in its critical section at a time to avoid race conditions
 - How to make multiple instructions of CS seem like one atomic instruction irrespective of underlying scheduling decisions?



Processes progress with non-zero speed, no assumption on clock speed

Used extensively in operating systems:
Queues, shared variables, interrupt handlers, etc.

Scheduler assumptions

Thread a:

```
while(i < 10)
    i = i + 1;
print "A won!";
```

Thread b:

```
while(i > -10)
    i = i - 1;
print "B won!";
```

If var **i** is shared, and initialized to 0

- Who wins on a single-CPU system?
- Is it guaranteed that someone wins?

Scheduler Assumptions

- Normally we assume that
 - A scheduler always gives every executable thread opportunities to run
 - In effect, each thread makes *finite progress*
 - But schedulers aren't always fair
 - Some threads may get more chances than others
 - To reason about worst case behavior we sometimes think of the scheduler as an adversary trying to “mess up” the algorithm

CSP Solution Structure

Shared vars:

Initialization:

Process/Thread:

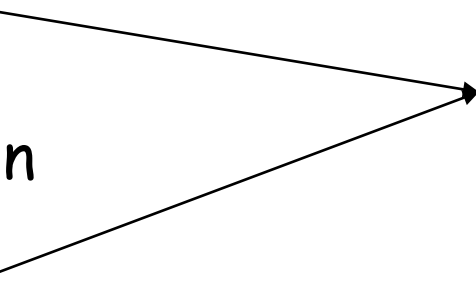
...

...

Entry Section

Critical Section

Exit Section




Added to solve the CS problem

Example Problem Setting

- Only 2 processes, P_0 and P_1
- General structure of process P_i (other process P_j)
do {
 CS entry
 critical section
 CS exit
 remainder section
} while (1);
- Processes may share some common variables to synchronize their actions.

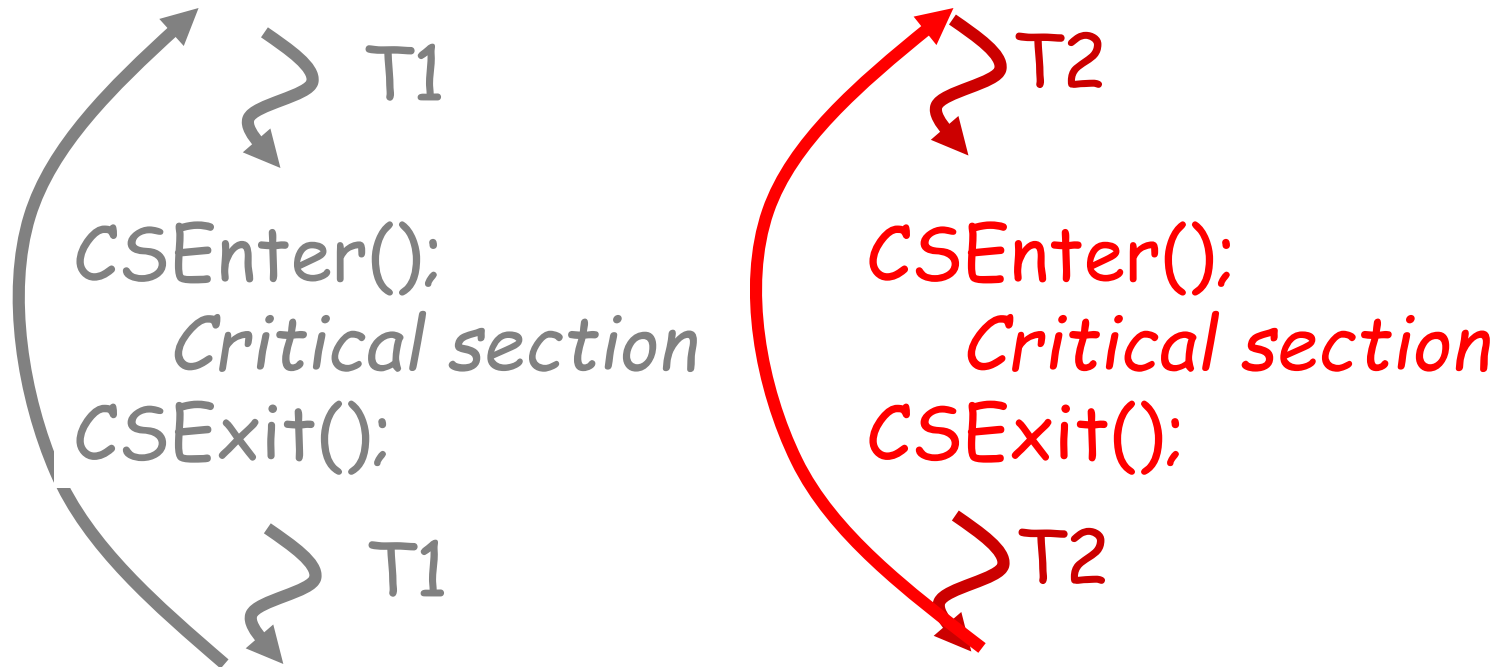
Critical Section Assumptions

- Tasks do some stuff but eventually might try to access shared data



Critical Section Assumptions

- Perhaps they loop (perhaps not!)



Critical Section Goals

- We would like
 1. **Safety (aka mutual exclusion)**
 - No more than one thread can be in a critical section (CS) at any time.
 2. **Liveness (aka progress)**
 - Only threads that are not executing in Remainder Sections can participate in deciding which will enter its CS next
 - A thread that is seeking to enter the CS will eventually succeed (i.e., selection can't be postponed indefinitely)
 3. **Bounded waiting**
 - A bound must exist on the number of times that other threads are allowed to enter their CSs after a thread has made a request to enter its CS and before that request is granted
- Assume that each process executes at a **nonzero** speed
- No assumption concerning relative speed of the **N** processes
- Ideally we would like **fairness** as well
 - If two threads are both trying to enter a critical section, they have equal chances of success
 - ... in practice, fairness is rarely guaranteed!

Solving the problem

- A first idea:
 - Have a boolean flag, *inside*. Initially false.

CSEnter()

{

while(inside) continue;

inside = true;

}

Code is unsafe: thread 0 could finish the while test when inside is false, but then 1 might call CSEnter() before 0 can set inside to true!

inside = false;

Race condition inside CSEnter code!!

- Now ask:
 - Is this Safe? Live? Bounded waiting?

Solving the problem: Take 2

- A different idea (assumes just two threads):
 - Have a boolean flag array. *inside[i]*. Initially false.

CSEnter(int i)

{

}

inside[i] = true;
while(inside[j]) continue;

{

}

Inside[i] = false;

Code isn't live: with bad luck, both threads could be looping, with 0 looking at 1, and 1 looking at 0

- Now ask:
 - Is this Safe? Live? Bounded waiting?

Solving the problem: Take 3

- Another broken solution, for two threads
 - Have a turn variable

```
CSEnter(int i)
```

```
{
```

```
    while(turn != i) continue;
```

```
}
```

```
    turn = j;
```

```
}
```

Code isn't live: thread 1 can't enter unless thread 0 did first, and vice-versa. But perhaps one thread needs to enter many times and the other fewer times, or not at all

- Now ask:
 - Is this Safe? Live? Bounded waiting?

Real-Life Analogy: Too Much Milk

Time	Person A (Alice)	Person B (Bob)
7:00	Look in Fridge. Out of milk	
7:05	Leave for store A	
7:10	Arrive at store A	Look in Fridge. Out of milk
7:15	Buy milk	Leave for store B
7:20	Arrive home, put milk away	Arrive at store B
7:25		Buy milk
7:30		Arrive home, put milk away

Too Much Milk: Correctness

1. At most one person buys milk
2. At least one person buys milk if needed

Solution Attempt #1

- Leave a note
 - Place on fridge before buying
 - Remove after buying
 - Don't go to store if there's already a note
- Leaving/checking a note is atomic (~ load/store)

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove Note;  
    }  
}
```

Attempt #1 in Action

Alice

```
if (noMilk) {  
    if (noNote) {  
  
        leave Note;  
        buy milk;  
        remove Note;  
    }  
}
```

Bob

```
if (noMilk) {  
    if (noNote) {  
  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```


Solution Attempt #2

Leave Note;

```
if (noMilk) {
```

```
  if (noNote) {
```

```
    leave Note;
```

```
    buy milk;
```

```
  }
```

```
}
```

```
remove Note;
```

But there's always a
note – you just left one!

At least you don't
buy milk twice...

Solution Attempt #3

- Leave a named note – each person ignores their own

Alice

```
leave note Alice
```

```
if (noMilk) {
```

```
    if (noNote Bob) {
```

```
        buy milk
```

```
    }
```

```
}
```

```
remove note Alice;
```

Bob

```
leave note Bob
```

```
if (noMilk) {
```

```
    if (noNote Alice) {
```

```
        buy milk
```

```
    }
```

```
}
```

```
remove note Bob;
```

Attempt #3 in Action

Alice

leave note Alice

if (noMilk) {

 if (noNote Bob) {

~~buy milk~~

 }

}

remove note Alice

Bob

leave note Bob

if (noMilk) {

 if (noNote Alice) {

~~buy milk~~

 }

remove note Bob

Solution Attempt #4

Alice

```
leave note Alice
while (note Bob) {
    do nothing
}
if (noMilk) {
    buy milk
}
remove note Alice;
```

Bob

```
leave note Bob
if (noNote Alice) {
    if (noMilk) {
        buy milk
    }
}
remove note Bob;
```

- This is a correct solution, but ...

Issues with Solution #4

- Complexity
 - Proving that it works is hard
 - How do you add another thread?
- Busy-waiting
 - Alice **consumes CPU time to wait**
- Fairness
 - Who is more likely to buy milk?

Peterson's Solution

- Process P_i
do {
 inside [i] := true;
 turn = j;
 while (inside [j] && turn == j) continue;
 critical section
 inside [i] = false;
 remainder section
} while (true);
- Meets all three requirements; solves the critical-section problem for **two processes**.

Analysis of Peterson's algorithm:

- Safety (by contradiction):
 - Assume that both processes (P1 and P2) are in their critical section (and thus have their *inside* flags set). Since only one, say P1, can have the *turn*, the other (P2) **must** have reached the `while()` test before P1 set his *inside* flag.
 - However, after setting his *inside* flag, P1 gave away the *turn* to P2. P2 has already changed the turn and **cannot** change it again, contradicting our assumption.

Liveness & Bounded waiting => the turn variable.

Issues with Peterson's algorithm

- Complexity: 3 variables: turn and a two-element array
- Limited for two processes, though it can be extended to any N processes with more complexity
- Relies on busy waiting (aka spin locks)
- It may not work on modern systems
 - load and store may not be atomic instructions
 - Ex: double-precision floating point store instruction
 - **store** mem, R1 && **store** mem, R2
 - `int l = 0; //Atomic`
 - `long long a = 0; // Not atomic, 2 cycles as memory bus is 32-bit`
 - memory reference ordering may not be preserved at the memory controller in multi-core/CPU systems

Dekker's Solution

CSEnter(int i)

```
{
    inside[i] = true;
    while(inside[J])
    {
        if (turn == J)
        {
            inside[i] = false;
            while(turn == J) continue;
            inside[i] = true;
        }
    }
}
```

CSExit(int i)

```
{
    turn = J;
    inside[i] = false;
}
```

Napkin analysis of Dekker's algorithm:

- Safety: No process will enter its CS without setting its *inside* flag. Every process checks the other process *inside* flag after setting its own. If both are set, the *turn* variable is used to allow only one process to proceed.
- Liveness: The *turn* variable is only considered when both processes are using, or trying to use, the resource
- Bounded waiting: The *turn* variable ensures alternate access to the resource when both are competing for access

Why does it work?

- Safety: Suppose thread 0 is in the CS.
 - Then `inside[0]` is true.
 - If thread 1 was simultaneously trying to enter, then `turn` must equal 0 and thread 1 waits
 - If thread 1 tries to enter “now”, it sets `turn` to 0 and waits
- Liveness: Suppose thread 1 wants to enter and can't (stuck in while loop)
 - Thread 0 will eventually exit the CS
 - When `inside[0]` becomes false, thread 1 can enter
 - If thread 0 tries to reenter immediately, it sets `turn=1` and hence will wait politely for thread 1 to go first!

Summary

- Dekker's algorithm does not provide *strict* alternation
 - Initially, a thread can enter critical section without accessing *turn*
- Dekker's algorithm will not work with many modern CPUs
 - CPUs execute their instructions in an out-of-order (OOO) fashion
 - This algorithm won't work on Symmetric MultiProcessors (SMP) CPUs equipped with OOO without the use of *memory barriers*
- Additionally, Dekker's algorithm can fail regardless of platform due to many optimizing compilers
 - Compiler may remove writes to *flag* since never accessed in loop
 - Further, compiler may remove *turn* since never accessed in loop
 - Creating an infinite loop!

Can we generalize Peterson's solution to more than 2 processes/threads?

- Obvious approach won't work:

```
CSEnter(int i)
{
    inside[i] = true;
    for(J = 0; J < N && J!=i; J++)
        while(inside[J] && turn == J)
            continue;
}
```

```
CSExit(int i)
{
    inside[i] = false;
}
```

- Issue: Who's turn next?

Bakery “concept”

- Described by Leslie Lamport
- Think of a popular store with a crowded counter
 - People take a ticket from a machine
 - If nobody is waiting, tickets don't matter
 - When several people are waiting, ticket order determines order in which they can make purchases

Bakery Algorithm: “Take 1”

- `int ticket[n]; // initialized to 0s`
- `int next_ticket; //initialized to 0`

`CSEnter(int i)`

```
{  
    ticket[i] = ++next_ticket;  
    for(J = 0; J < n && J!=i; J++)  
        while(ticket[J] && ticket[J] < ticket[i])  
            continue;  
}
```

`CSExit(int i)`

```
{  
    ticket[i] = 0;  
}
```

- Oops... access to `next_ticket` is a problem!
- `++next_ticket` is not atomic

Bakery Algorithm: “Take 2”

- `int ticket[n];`

Just add 1 to the max!

`CSEnter(int i)`

`{`

`ticket[i] = max(ticket[0], ... ticket[N-1])+1;`

`for(J = 0; J < n && J!=i; J++)`

`while(ticket[J] && ticket[j] < ticket[i])`

`continue;`

`}`

`CSExit(int i)`

`{`

`ticket[i] = 0;`

`}`

- Clever idea: just add one to the max.
- Oops... two could pick the same value!

Bakery Algorithm: “Take 3”

If i, j pick same ticket value, id's break tie:

$(\text{ticket}[J] < \text{ticket}[i]) \parallel (\text{ticket}[J] == \text{ticket}[i] \ \&\& \ J < i)$

Notation: $(B, J) < (A, i)$ to simplify the code:

$(B < A \parallel (B == A \ \&\& \ J < i)), \text{ e.g.:}$

$(\text{ticket}[J], J) < (\text{ticket}[i], i)$

Bakery Algorithm: “Take 4”

- `int ticket[N];`
- `boolean picking[N] = false;`

`CSEnter(int i)`

```
{  
    ticket[i] = max(ticket[0], ... ticket[N-1])+1;  
    for(J = 0; J < N && J!=i; J++)  
        while(ticket[J] && (ticket[J],J) < (ticket[i],i))  
            continue;  
}
```

`CSExit(int i)`

```
{  
    ticket[i] = 0;  
}
```

- Both `i` & `J` enter into `max()` at the same time (i.e, get same ticket number), but `i` goes fast and enters while
- Oops... `i` could look at `J` when `J` is still storing its ticket, and yet `J` could have a lower id than me (`i`)!
 - So, both enter CS!!

Bakery Algorithm: Almost final

- `int ticket[N];`
- `boolean choosing[N] = false;`

`CSEnter(int i)`

```
{
    choosing[i] = true;
    ticket[i] = max(ticket[0], ... ticket[N-1])+1;
    choosing[i] = false;
    for(J = 0; J < N && J!=i; J++) {
        while(choosing[J]) continue;
        while(ticket[J] && (ticket[J],J) < (ticket[i],i))
            continue;
    }
}
```

`CSExit(int i)`

```
{
    ticket[i] = 0;
}
```

Bakery Algorithm: Final

- `int ticket[N];`
- `boolean choosing[N] = false;`

`CSEnter(int i)`

```
{
    do {
        ticket[i] = 0;
        choosing[i] = true;
        ticket[i] = max(ticket[0], ... ticket[N-1])+1;
        choosing[i] = false;
    } while(ticket[i] >= MAXIMUM);
    for(J = 0; J < N && J!=i; J++) {
        while(choosing[J]) continue;
        while(ticket[J] && (ticket[J],J) < (ticket[i],i))
            continue;
    }
}
```

`CSExit(int i)`

```
{
    ticket[i] = 0;
}
```

Bakery Algorithm: Issues?

- What if we don't know how many threads might be running?
 - The algorithm depends on having an agreed upon value for N
 - Somehow would need a way to adjust N when a thread is created or one goes away
- Also, technically speaking, ticket can overflow!
 - Solution: Change code so that if ticket is “too big”, set it back to zero and try again.

How do real systems do it?

- Some real systems actually use algorithms such as the bakery algorithm
 - A good choice where busy-waiting isn't going to be super-inefficient
 - For example, if you have enough CPUs so each thread has a CPU of its own
- Some systems disable interrupts briefly when entering Critical Sections
 - Time consuming in multi-CPU systems for msg passing others about interrupt disabling
- Some systems take hardware “help”: Special *atomic instructions*
 - *Test-and-Set, Swap, Read-Modify-Write, etc*
 - *They guarantee exclusive access to memory locations (aka lock variables)*

Concurrent Applications

Shared Objects

Bounded Buffer Barrier

Synchronization Variables

Semaphores Locks Condition Variables

Atomic Instructions

Interrupt Disable Test-and-Set

Hardware

Multiple Processors Hardware Interrupts

Lock Variables

- Lock variable is a shared object that is guaranteed to be accessed by a single process at any given time
- Lock variables are used to achieve mutual exclusion

Process/Thread:

...

Acquire Lock

Critical Section

Release Lock

...

Test_and_Set Instruction

- Definition:

```
boolean test_and_set (boolean *target)
{
    boolean retValue = *target;
    *target = TRUE;
    return retValue;
}
```

Critical Sections with Atomic Hardware Primitives

Process i

Share: int lock;
Initialize: lock = false;

```
While(test_and_set(&lock))  
CONTINUE;
```

Assumes that test_and_set is compiled to a special hardware instruction that sets the lock and returns the OLD value (true: locked; false: unlocked)

Critical Section

```
lock = false;
```

Problem: Does not satisfy bounded waiting)
(see Galvin 9th Ed for correct solution, Figure 5.7)

Solution using TestAndSet

- Shared boolean variable *lock* initialized to FALSE.
- Solution:

```
while (true) {  
    while ( test_and_set (&lock ))  
        Continue; /* do nothing  
  
        //    critical section  
  
    lock = FALSE;  
  
        //    remainder section  
  
}
```

Compare_and_Swap Instruction

- Definition:

```
int compare_and_swap (int *value, int expected, int new_value)
{
    int temp = *value;
    if (*value == expected)
        *value = new_value;

    return temp;
}
```

Solution using Compare_and_Swap

- Shared Boolean variable **lock** initialized to FALSE (0)
- Solution:

```
    while (true) {  
        while (compare_and_swap(&lock, 0, 1) !=0)  
            continue;
```

```
        //    critical section
```

```
        lock = 0;
```

```
        //    remainder section
```

```
    }
```

H/W Implementation of Test_and_Set and Compare_and_Swap Instructions

- Test_and_Set is implemented using TSL h/w instruction
 - TSL register, lock
 - It loads the content of the variable named lock to register and writes a not NULL value into the lock variable atomically
- Compare_and_Swap is implemented using CMPXCHG h/w instruction
 - CMPXCHG register, lock
 - Swaps the contents of the register and the lock variable atomically
- H/W guarantees atomicity by disabling the access to the memory bus to the other processors
- Nowadays the [x86 instruction set](#) allows all instructions to be atomic
 - The lock prefix makes atomic any instruction following it
 - **lock store register, memLoc**
 - **This lock prefix can be used with all the memory access instructions** 70
 - **INC, DEC, ADD, SUB, XOR, etc**

Mutex Locks

- Previous solutions are complicated and generally inaccessible to application programmers
- OS designers build software tools to solve critical section problem
- Simplest is the mutex lock
- Protect critical regions with it by first **acquire()** a lock then **release()** it
 - Boolean variable indicating if lock is available or not
- Calls to **acquire()** and **release()** must be atomic
 - CSEntry and CSExit facing race conditions
 - To solve this CS problem in CSEntry and CSExit method, they are usually implemented via hardware atomic instructions!!
 - Test_and_Set and Compare_and_Swap
- But this solution requires **busy waiting**
 - This lock therefore called a **spinlock**
- **Even hardware instructions** Test_and_Set and Compare_and_Swap suffer from busy waiting!

acquire() and release()

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}  
release() {  
    available = true;  
}  
do {  
    acquire()  
    critical section  
  
    release()  
    remainder section  
} while (true);
```



```
while(TestAndSet(&lock));  
while (!available);  
    available = false;  
lock=false;
```

- Code for release()?
- Is preemption possible in critical section?

Presenting critical sections to users

- CSEnter and CSExit are possibilities
- But more commonly, operating systems have offered a kind of locking primitive
- We call these semaphores

Semaphores

- Non-negative integer with atomic increment and decrement
- Integer 'S' that (besides initialization) can only be modified by:
 - P(S) or S.wait(): decrement or block if already 0
 - V(S) or S.signal(): increment and wake up process if any
- These operations are *atomic* (indivisible)

semaphore

These systems use the operation **signal()** instead of **V()**

```
P(S) {  
    while(S ≤ 0)  
        ;  
    S--;  
}  
  
V(S) {  
    S++;  
}
```

Semaphore Types

- Counting Semaphores:
 - Any integer
 - Used for synchronization
- Binary Semaphores
 - Value is limited to 0 or 1
 - Used for mutual exclusion (mutex)

Process i

Shared: semaphore S

P(S);

Init: S = 1;

Critical Section

V(S);

Semaphore Implementation

- Must guarantee that no two processes can execute $P()$ and $V()$ on the same semaphore at the same time
 - No process may be interrupted in the middle of these operations
- Thus, implementation becomes the critical section problem where the P and V codes are placed in the critical section.
 - Could now **have busy waiting in critical section of critical section** implementation!
 - But implementation code is short
 - Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution.
- Busy waiting (spinlocks)
 - ☞ Consumes CPU resources
 - ☞ No context switch overhead, so OK in multi-CPU systems

Semaphore Implementation with no Busy waiting

- Alternative approach: Blocking
- With each semaphore there is an associated **waiting queue**. Each entry/node in a waiting queue (linked list) has two data items:
 - pointer to PCB of a blocked process
 - pointer to next record in the list
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue.
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue.
- Should spin or block?
 - Less time \Rightarrow spin
 - More time \Rightarrow block
 - A theory result:
 - Spin for as long as block cost
 - If lock is still not available, then block
 - Shown factor of 2-optimal!

Implementing Semaphores

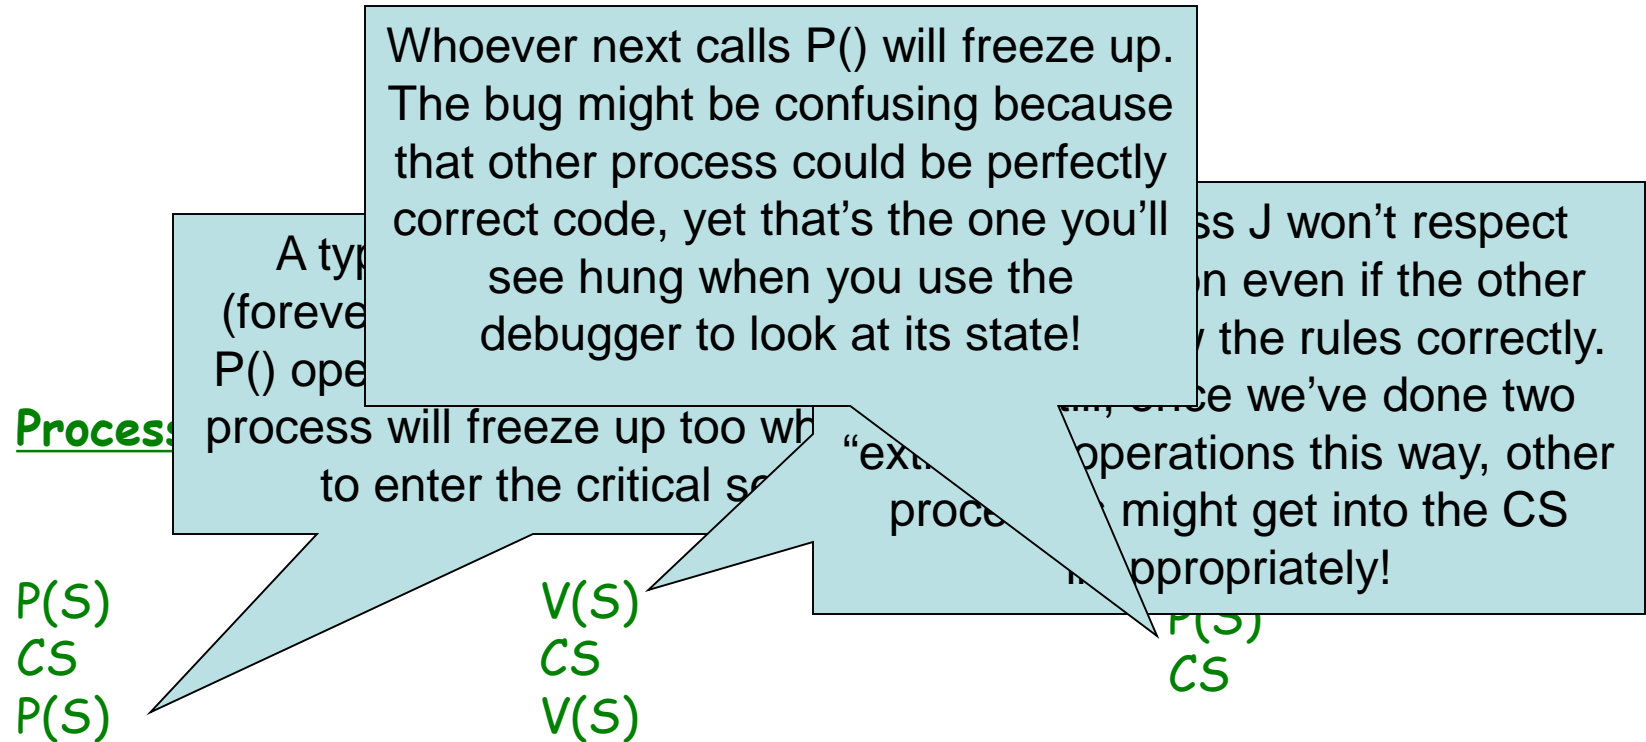
```
typedef struct semaphore {  
    int value;  
    ProcessList L;  
} Semaphore;  
  
void P(Semaphore *S) {  
    S->value = S->value - 1;  
    if (S->value < 0) {  
        add this process P to S.L;  
        block(P); //moved to  
        waiting Queue of Semaphore  
    }  
}
```

```
void V(Semaphore *S) {  
    S->value = S->value + 1;  
    if (S->value <= 0) {  
        remove a process Q from  
        S.L;  
        wakeup(Q); //moved from  
        waiting to Ready Queue of CPU  
    }  
}
```

Implementing Semaphores

- Per-semaphore list of processes
 - Implemented using PCB link field
 - Queuing Strategy: FIFO works fine
 - Will LIFO work?

Common programming errors



More common mistakes

- Conditional code that can break the normal top-to-bottom flow of code in the critical section
- Often a result of someone trying to maintain a program, e.g. to fix a bug or add functionality in code written by someone else

```
P(S)
if(something or other)
    return;
CS
V(S)
```

Producer-Consumer Problem

- Solving with semaphores
 - We'll use two kinds of semaphores
 - We'll use *counters* to track how much data is in the buffer
 - One **counter** counts as we add data and stops the producer if there are N objects in the buffer
 - A second **counter** counts as we remove data and stops a consumer if there are 0 in the buffer
 - Idea: since general semaphores can count for us, we don't need a separate counter variable
- Why do we need a second kind of semaphore?
 - We'll also need a mutex semaphore

Producer-consumer with a bounded buffer

- Problem Definition
 - Producer puts things into a shared buffer (wait if full)
 - Consumer takes them out (wait if empty)
 - Use a fixed-size buffer between them to avoid lockstep
 - Need to synchronize access to this buffer
- Correctness Constraints:
 - Consumer must wait for producer to fill buffers, if none full
 - scheduling constraint
 - Producer must wait for consumer to empty buffers, if all full
 - scheduling constraint
 - Only one thread can manipulate buffer queue at a time
 - mutual exclusion
- Remember why we need mutual exclusion
- General rule of thumb:
Use a separate semaphore for each constraint
 - Semaphore full; // consumer's constraint
 - Semaphore empty; // producer's constraint
 - Semaphore mutex; // mutual exclusion

Producer-Consumer Problem

```
Init: Semaphore mutex = 1; /* for mutual exclusion*/  
      Semaphore empty = N; /* number empty buf entries */  
      Semaphore full = 0;   /* number full buf entries */  
      any_t buf[N];  
      int tail = 0, head = 0;
```

Producer

```
void put(char ch) {  
  
    P(empty);  
    P(mutex);  
  
    // add ch to buffer  
    buf[head%N] = ch;  
    head++;  
  
    V(mutex);  
    V(full); //not V(empty)  
}
```

Consumer

```
char get() {  
  
    P(full);  
    P(mutex);  
  
    // remove ch from buffer  
    ch = buf[tail%N];  
    tail++;  
  
    V(mutex);  
    V(empty); //not V(full)  
  
    return ch;  
}
```

What's wrong?

```
Init: Semaphore mutex = 1; /* for mutual exclusion*/  
      Semaphore empty = N; /* number empty buf entries */  
      Semaphore full = 0;   /* number full buf entries */  
      char *buf[N];
```

Producer

```
void put(char ch
```

```
P(mutex);  
P(empty);
```

```
// add ch to buffer  
buf[head%N] = ch;  
head++;
```

```
V(mutex);  
V(full);
```

```
}
```

Oops! Even if you do the correct operations, the order in which you do semaphore operations can have an incredible impact on correctness

What if buffer is full?

Consumer

```
get() {
```

```
full);
```

```
P(mutex);
```

```
// remove ch from buffer  
ch = buf[tail%N];  
tail++;
```

```
V(mutex);  
V(empty);
```

```
return ch;
```

```
}
```

What's wrong?

```
Init: Semaphore mutex = 1; /* for mutual exclusion*/  
      Semaphore empty = N; /* number empty buf entries */  
      Semaphore full = 0;   /* number full buf entries */  
      any_t buf[N];  
      int tail = 0, head = 0;
```

Producer

```
void put(char ch) {  
  
    P(empty);  
    P(mutex);  
  
    // add ch to buffer  
    buf[head%N] = ch;  
    head++;  
  
    V(full);  
    V(mutex);  
}
```

Consumer

```
char get() {  
  
    P(full);  
    P(mutex);  
  
    // remove ch from buffer  
    ch = buf[tail%N];  
    tail++;  
  
    V(mutex);  
    V(empty);  
  
    return ch;  
}
```

What's wrong?

```
Init: Semaphore mutex = 1; /* for mutual exclusion*/  
      Semaphore empty = N; /* number empty buf entries */  
      Semaphore full = 0;   /* number full buf entries */  
      any_t buf[N];  
      int tail = 0, head = 0;
```

Producer

```
void put(char ch) {  
  
    P(empty);  
    P(mutex);  
  
    // add ch to buffer  
    buf[head%N] = ch;  
    head++;  
  
    V(mutex);  
    V(full);  
}
```

Consumer

```
char get() {  
  
    P(full);  
    P(mutex);  
  
    // remove ch from buffer  
    ch = buf[tail%N];  
    tail++;  
  
    return ch;  
  
    V(mutex);  
    V(empty);  
}
```

Discussion about Bounded Buffer Solution

- Why asymmetry?
 - Producer does: $P(\text{empty}), V(\text{full})$
 - Consumer does: $P(\text{full}), V(\text{empty})$
- Is order of P's important?
 - Yes! Can cause deadlock
- Is order of V's important?
 - No, except that it might affect scheduling efficiency
- What if we have 2 producers or 2 consumers?
 - Do we need to change anything?

In a Nut Shell...

- Fundamental Issue
 - Programmers *atomic* operation is not done atomically
 - Atomic Unit: instruction sequence guaranteed to execute indivisibly
 - Also called “critical section” (CS)
- Critical Section Implementation
 - Software: Dekker’s, Peterson’s, Baker’s algorithm
 - Hardware: test_and_set, swap
 - Hard for programmers to use
 - Operating System: semaphores
- Implementing Semaphores
 - Multithread synchronization algorithms shown earlier
 - Could have a thread disable interrupts, put itself on a “wait queue”, then context switch to some other thread (an “idle thread” if needed)
 - The O/S designer makes these decisions and the end user shouldn’t need to know

Reading Assignment

- Chapter 5 from OSC by Galvin et al 9th Edition
- Chapter 2 from MOS by Tanenbaum et al

References

- "Numbers Everyone Should Know" from Jeff Dean: http://brenocon.com/dean_perf.html
- Peterson's algorithm:
- https://en.wikipedia.org/wiki/Peterson%27s_algorithm
- Bakery algorithm: <http://www.cs.umd.edu/~shankar/412-S99/note-7.html>
- [Intel 64 and IA-32 Archs: Chapter 8, Multi-Processor Mgmt](#)