

Other Classical Synchronization Problems


Outline

- Classic Synchronization problems
 - Producer-Consumer with bounded buffer (already covered)
 - Reader-Writer
 - Dining Philosophers
- Solutions to classic problems via semaphores
- Synchronization in Windows/Linux/Pthreads

Readers and Writers

- In this model, threads share data that
 - some threads “read” and other threads “write”.
- Instead of CSEnter and CSExit we want
 - StartRead...EndRead; StartWrite...EndWrite
- Goal: allow multiple concurrent readers but only a single writer at a time, and if a writer is active, readers and other writers wait for it to finish

Readers/Writers Problem



- Motivation: Consider a shared database
 - Two classes of users:
 - Readers – never modify database
 - Writers – read and modify database
 - Is using a single lock on the whole database sufficient?
 - Like to have many readers at the same time
 - Only one writer at a time

Readers-Writers Problem

- Courtois et al 1971
- Models access to a database
 - A reader is a thread that needs to look at the database but won't change it.
 - A writer is a thread that modifies the database
- Example: booking movie tickets on BookMyShow
 - When you browse to look at movie schedules and seat availability, the website is acting as a reader on your behalf
 - When you reserve a ticket, the website has to write into the database to make the reservation

Readers-Writers Problem

- Many threads share an object in memory
 - Some write to it, some only read it
 - Only one writer can be active (enters CS) at a time
 - Any number of readers can be active simultaneously
- Key insight: generalizes the critical section concept
- One issue we need to settle, to clarify problem statement.
 - Suppose that a writer is active and a mixture of readers and writers now shows up. Who should get in next?
 - Or suppose that a writer is waiting and an endless stream of readers keeps showing up. Is it fair for them to become active?
- Ideally we would like a kind of back-and-forth form of fairness:
 - Once a reader is waiting, readers will get in next.
 - If a writer is waiting, one writer will get in next.

Readers-Writers

Shared variables: Semaphore mutex, rwl;
integer rcount;

Init: mutex = 1, rwl = 1, rcount = 0;

Writer

```
do {  
  
    Wait(rwl);  
    ...  
    /*writing is performed*/  
    ...  
    Signal(rwl);  
  
}while(TRUE);
```

Reader

```
do {  
    Wait(mutex);  
    rcount++;  
    if (rcount == 1)  
        Wait(rwl);  
  
    Signal(mutex);  
    ...  
    /*reading is performed*/  
    ...  
    Wait(mutex);  
    rcount--;  
    if (rcount == 0)  
        Signal(rwl);  
  
    Signal(mutex);  
}while(TRUE);
```

Readers-Writers Notes


- If there is a writer active
 - First reader blocks on **rwl**
 - Other readers block on **mutex**
- Once a reader is active, all readers get to go through
 - Any problem?
- Once a writer exits, which one gets in first?
- The last reader to exit signals a writer
 - If no writer, then readers can continue
- If readers and writers waiting on **rwl**, and writer exits
 - Who gets to go in first?
 - Reader (which Reader gets in first?) or Writer?
- Why doesn't a writer need to use **mutex**?

Does this work as we hoped?

- If readers are active, no writer can enter
 - The writers wait doing a `Wait(rwl)`
- While writer is active, nobody can enter
 - Any other reader or writer will wait
- But back-and-forth switching is buggy:
 - Any number of readers can enter in a row
 - Readers can “starve” writers
- With semaphores, building a solution that has the desired back-and-forth behavior is really, really tricky!
 - We recommend that you try, but not too hard...

Dining Philosophers (1)

- Philosophers eat/think
- Eating needs 2 forks
- Pick one fork at a time
- Fork → Binary Semaphore
 - take_fork(i): Wait(i)
 - put_fork(i): Signal(i)
- Any problem?
- How to prevent deadlock?



Dining Philosophers (2)

```
#define N 5                                /* number of philosophers */

void philosopher(int i)                    /* i: philosopher number, from 0 to 4 */
{
    while (TRUE) {
        think( );                          /* philosopher is thinking */
        take_fork(i);                      /* take left fork */
        take_fork((i+1) % N);              /* take right fork; % is modulo operator */
        eat();                             /* yum-yum, spaghetti */
        put_fork(i);                       /* put left fork back on the table */
        put_fork((i+1) % N);               /* put right fork back on the table */
    }
}
```

A nonsolution to the dining philosophers problem

Dining Philosophers (3)

```
#define N          5          /* number of philosophers */
#define LEFT      (i+N-1)%N   /* number of i's left neighbor */
#define RIGHT     (i+1)%N     /* number of i's right neighbor */
#define THINKING  0          /* philosopher is thinking */
#define HUNGRY    1          /* philosopher is trying to get forks */
#define EATING    2          /* philosopher is eating */
typedef int semaphore;       /* semaphores are a special kind of int */
int state[N];               /* array to keep track of everyone's state */
semaphore mutex = 1;        /* mutual exclusion for critical regions */
semaphore s[N]; = {0}      /* one semaphore per philosopher */

void philosopher(int i)     /* i: philosopher number, from 0 to N-1 */
{
    while (TRUE) {          /* repeat forever */
        think( );           /* philosopher is thinking */
        take_forks(i);      /* acquire two forks or block */
        eat( );             /* yum-yum, spaghetti */
        put_forks(i);       /* put both forks back on table */
    }
}
```

Philosopher → Binary Semaphore

Solution to dining philosophers problem (part 1)

Dining Philosophers (4)

```
void take_forks(int i)                /* i: philosopher number, from 0 to N-1 */
{
    down(&mutex);                      /* enter critical region */
    state[i] = HUNGRY;                 /* record fact that philosopher i is hungry */
    test(i);                           /* try to acquire 2 forks */
    up(&mutex);                         /* exit critical region */
    down(&s[i]);                        /* block if forks were not acquired */
}

void put_forks(i)                     /* i: philosopher number, from 0 to N-1 */
{
    down(&mutex);                      /* enter critical region */
    state[i] = THINKING;               /* philosopher has finished eating */
    test(LEFT);                        /* see if left neighbor can now eat */
    test(RIGHT);                       /* see if right neighbor can now eat */
    up(&mutex);                         /* exit critical region */
}

void test(i)                          /* i: philosopher number, from 0 to N-1 */
{
    if (state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
        state[i] = EATING;
        up(&s[i]);
    }
}
```

Summary: semaphores


- Semaphores are very “low-level” primitives
 - Users could easily make small errors
 - Similar to programming in assembly language
 - Small error brings system to grinding halt
 - Very difficult to debug
- Also, we seem to be using them in two ways
 - For mutual exclusion, the “real” abstraction is a critical section
 - But the bounded buffer example illustrates something different, where threads “communicate” using semaphores
- Simplification: Provide concurrency support in compiler
 - Monitors (next class)

Synchronization Examples

- Windows
- Linux
- Pthreads

Windows Synchronization

- Multi-threaded kernel with so many global, shared resources
- Uses interrupt masks to protect access to global resources on single-CPU systems
- Uses **spinlocks** on multi-CPU systems
 - Spinlocked-thread will never be preempted by the kernel
- Outside kernel, it provides **dispatcher objects** which may act as mutexes, semaphores, events, and timers
 - **Events**
 - An event acts much like a condition variable
 - Timers notify one or more threads when time expired
 - Dispatcher objects in either **signaled-state** (object available) or **non-signaled state** (thread will block)



Linux Synchronization

- Linux:
 - Prior to kernel Version 2.6, disables interrupts to implement short critical sections
 - Version 2.6 and later, fully preemptive and reentrant
- Linux provides following locks for sync in kernel:
 - Atomic integer (***atomic_t*** counter)
 - semaphores
 - spinlocks
 - Mutex locks (no spin, but blocking)
 - reader-writer versions of both
- SMP systems: fundamental locking mechanism is spinlocks
 - But kernel is coded such that these are held only for a shot while
- On single-CPU system, spinlocks are replaced by enabling and disabling kernel preemption

Pthreads Synchronization

- Pthreads API is OS-independent and available to user/application programs
- It provides:
 - mutex locks (no spin, but blocking)
 - condition variables (no state associated unlike counting semaphores)
 - Refer Tanenbaum MOS: Fig 2.32 for sol. to P-C prob.
 - Read-write locks
 - Semaphores as part of POSIX SEM extension
 - Named → actual name in file system and shared among processes
 - Unnamed → among threads of same process
 - Wait: do not decrease its value below zero

Pthreads Synchronization Example

```
#include <pthread.h>
pthread_mutex_t mutex;
...
/* create the mutex lock */
pthread_mutex_init(&mutex,NULL); //dynamic
....
/* acquire the mutex lock */
pthread_mutex_lock(&mutex);
/* critical section */
/* release the mutex lock */
...
pthread_mutex_unlock(&mutex);
```

```
mutex=PTHREAD_MUTEX_INITIALIZER; //static
pthread_mutex_trylock(&mutex)
pthread_mutex_timedlock(&mutex,time)
pthread_mutex_destroy(&mutex); //for dynamic
```

Pthreads Synchronization Example

```
#include <semaphore.h>

sem_t sem;

/* Create the semaphore and initialize it to 1 */
sem_init(&sem, 0, 1); // 0: no sharing with other
processes, but only among other threads

....

/* acquire the semaphore */
sem_wait(&sem);
/* critical section */
/* release the semaphore */
sem_post(&sem);

...
```

Other Approaches

- Multi-core systems
 - Make designing deadlock-free and starvation-free multithreading apps very difficult
- Transaction memory
 - From Databases
 - If all operations in a transaction (set of RW operations) are completed, it is committed
 - If not, roll back
 - It guarantees atomicity, so chances of deadlock and starvation bcz of developers mistakes!
- Compiler's helping hand: Example: OpenMP
 - `#pragma omp critical`
Count++

Reading Assignment

- Chapter 5 from OSC by Galvin et al 9th Edition
- Chapter 2 from MOS by Tanenbaum et al
- http://man7.org/linux/man-pages/man7/sem_overview.7.html
- <https://computing.llnl.gov/tutorials/openMP/>
- <https://computing.llnl.gov/tutorials/pthreads/>

Quiz 1

- January 21st (TUE) at 2:30PM
- Syllabus: Lessons 1-4 on Synchronization