

Deadlocks

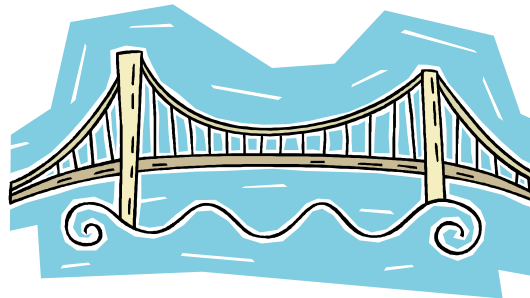
Outline

Part I: Prevention and Avoidance

- Discussion of Deadlocks
- Conditions for its occurrence
- Solutions for preventing/avoiding deadlocks

Deadlocks

- Complex multi-threaded programs contain several shared objects where we require multi-object synchronization with many synchronization primitives
- Definition: Deadlock exists among a set of processes if
 - Every process is waiting for an event
 - This event can be caused only by another process in the set
 - Event is the acquire or release of another resource



One-lane bridge

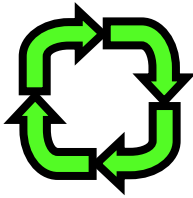


System Model

- There are non-sharable computer resources
 - Printers, Locks, Semaphores, DVD Drives, Tape drives, CPUs
 - Some of these can be made sharable with clever techniques
 - Maybe more than one instance of the resource
- Processes/threads need access to these resources
 - Acquire resource
 - If resource is available, access is granted
 - If not available, the process is blocked
 - Use resource
 - Release resource
- Undesirable scenario:
 - Process A acquires resource 1, and is waiting for resource 2
 - Process B acquires resource 2, and is waiting for resource 1

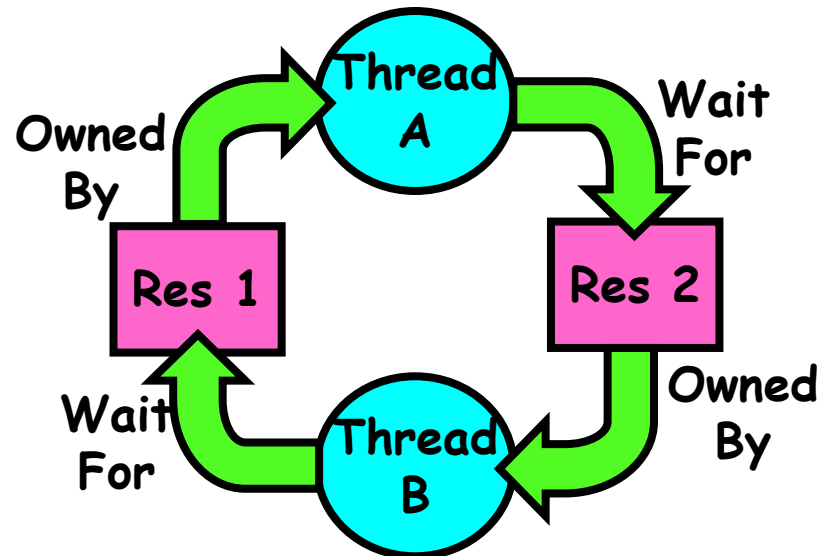
⇒ Deadlock!

Starvation vs Deadlock



- Starvation vs. Deadlock

- Starvation: thread waits indefinitely
 - Example, low-priority thread waiting for resources constantly in use by high-priority threads
- Deadlock: circular waiting for resources
 - Thread A owns Res 1 and is waiting for Res 2
 - Thread B owns Res 2 and is waiting for Res 1



- Deadlock \Rightarrow Starvation but not vice versa
 - Starvation can end (but doesn't have to)
 - Deadlock can't end without external intervention

For example: Semaphores

```
semaphore:      mutex1 = 1    /* protects resource 1 */  
                mutex2 = 1    /* protects resource 2 */
```

Process A code:

```
{  
    /* initial compute */  
    P(mutex1)  
    P(mutex2)  
  
    /* use both resources */  
  
    V(mutex2)  
    V(mutex1)  
}
```

Process B code:

```
{  
    /* initial compute */  
    P(mutex2)  
    P(mutex1)  
  
    /* use both resources */  
  
    V(mutex1)  
    V(mutex2)  
}
```



Swapping Ps?

Deadlock or starvation ? Always??

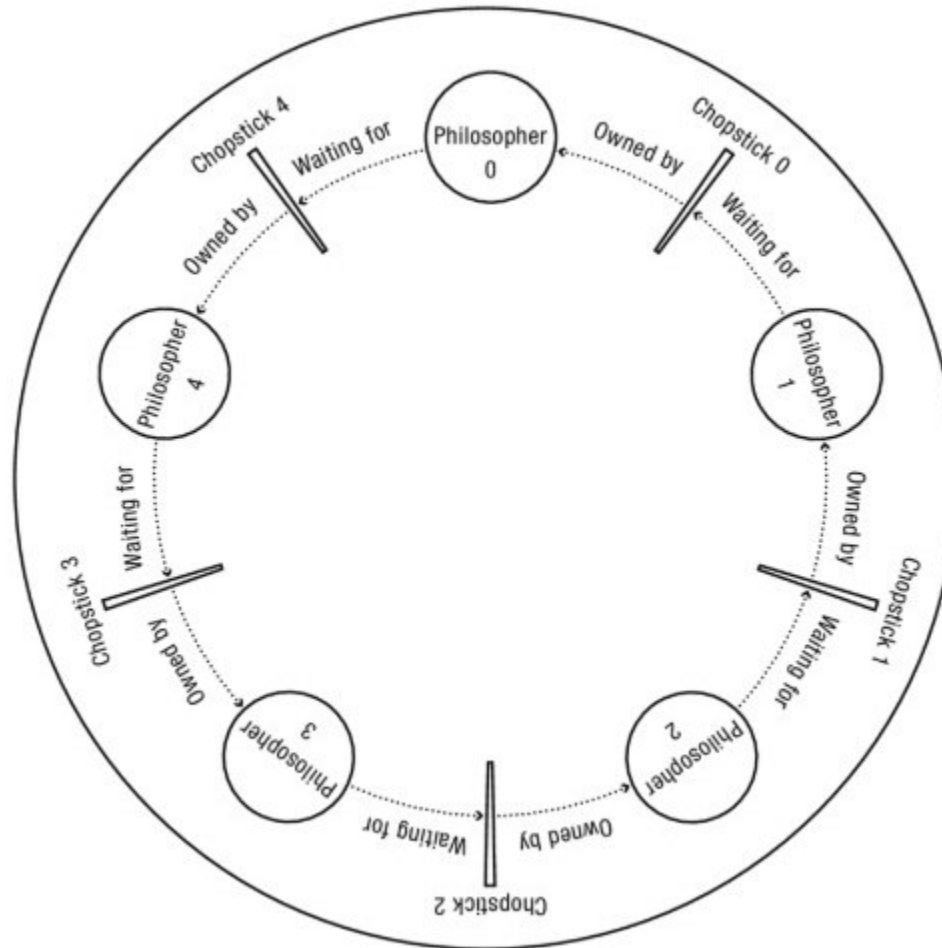
Since testing may not discover deadlocks, it's important to construct systems that are deadlock-free by design.

Four Conditions for Deadlock

- Coffman et. al. 1971
- Necessary conditions for deadlock to exist:
 - **Mutual Exclusion**
 - At least one resource must be held in non-sharable mode by a process
 - **Hold and wait**
 - There exists a process holding at least one resource, and waiting for additional resources
 - **No preemption**
 - Resources cannot be preempted
 - **Circular wait**
 - There exists a set of processes $\{P_1, P_2, \dots, P_N\}$, such that
 - P_1 is waiting for P_2 , P_2 for P_3 , and P_N for P_1

All four conditions must hold for deadlock to occur

Dining Philosophers meet ALL Four Conditions for Deadlock



Real World Deadlocks?

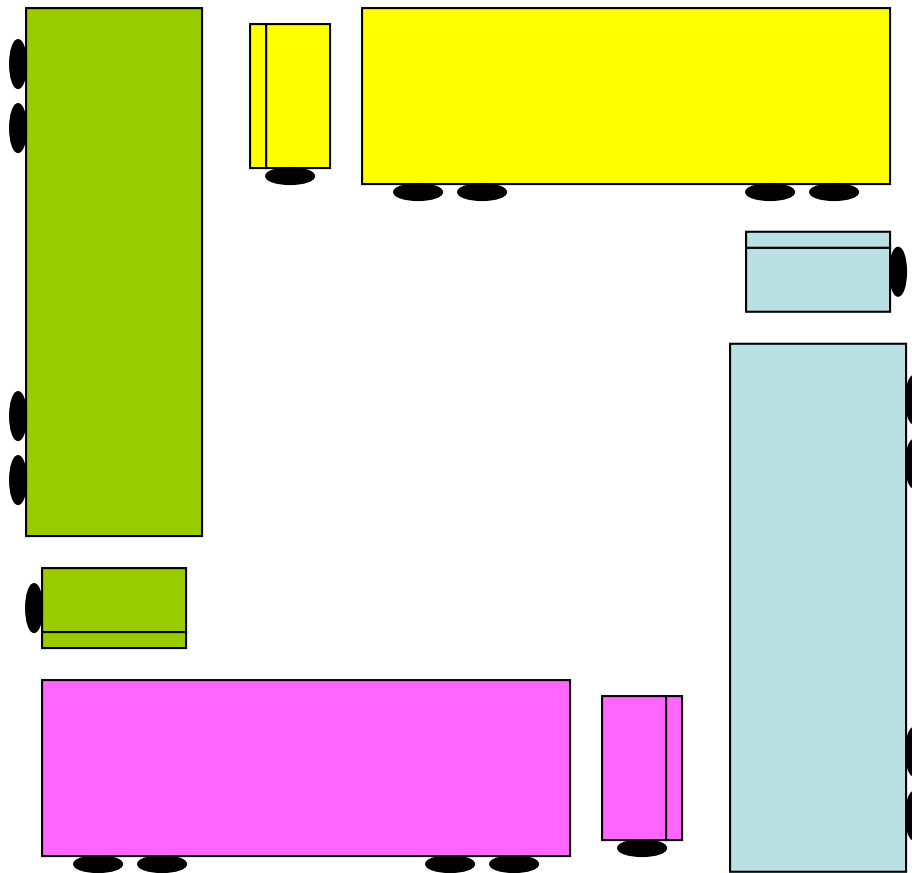
- Truck A has to wait for truck B to move



- Not deadlocked

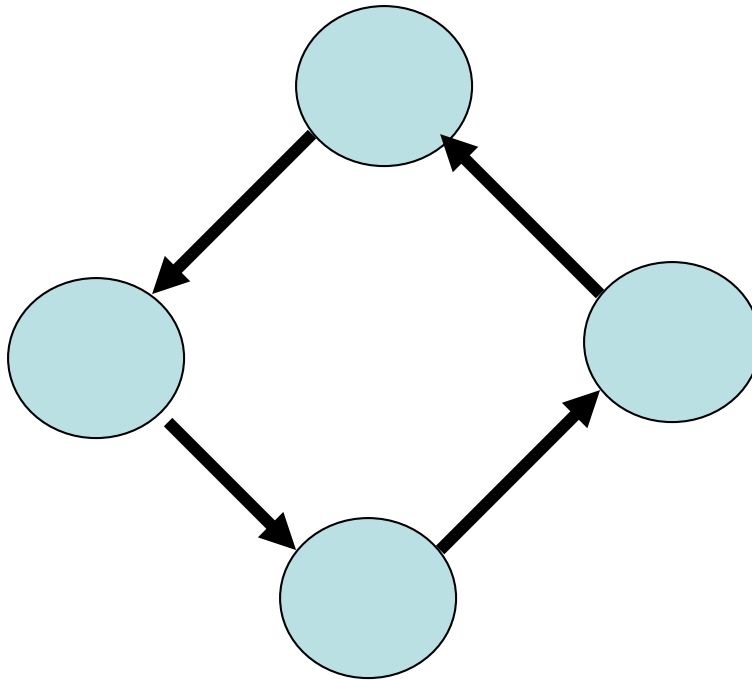
Real World Deadlocks?

- Gridlock



Real World Deadlocks?

- Gridlock



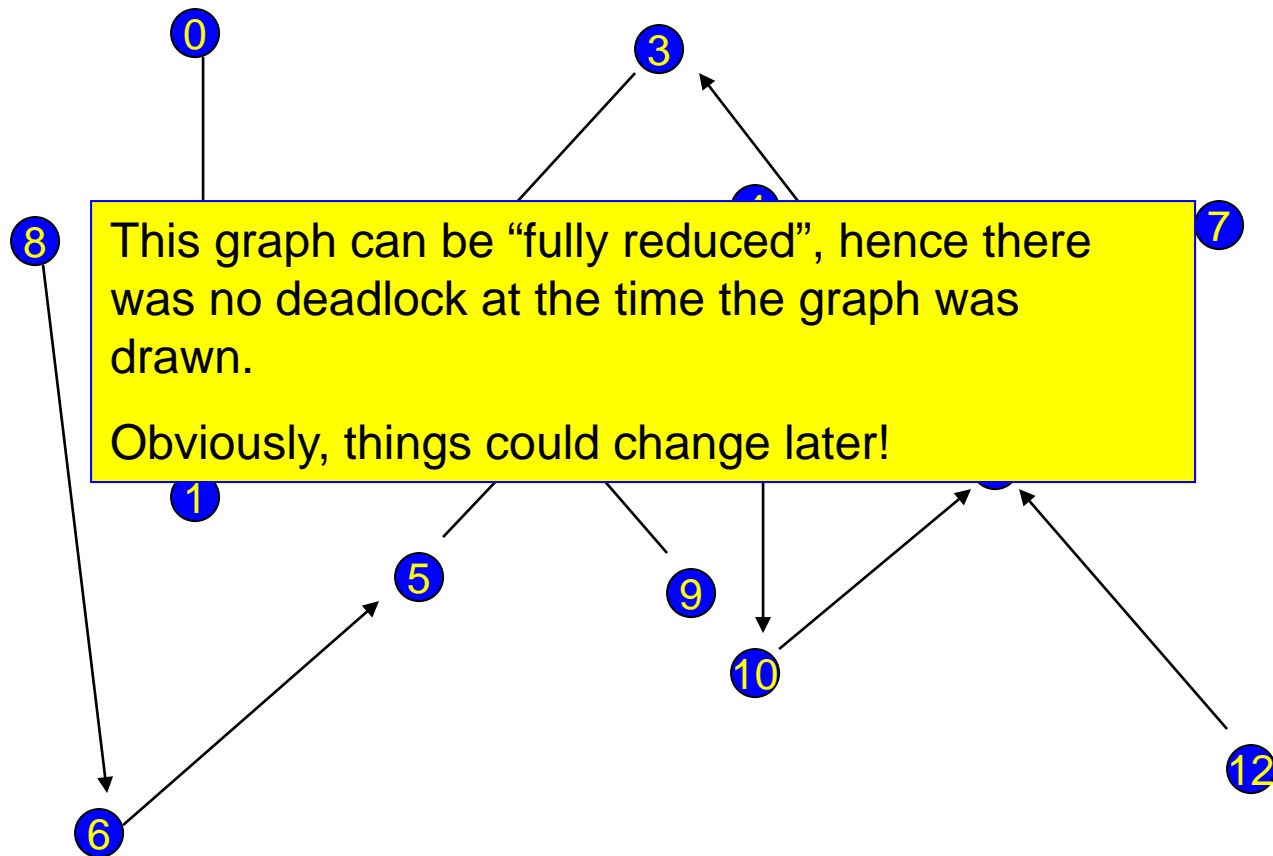
Testing for deadlock

- Steps
 - Collect “process state” and use it to build a graph
 - Ask each process “are you waiting for anything” held by others?
 - Put an edge in the graph between the processes if so
 - We need to do this in a single instant of time, not while things might be changing
- Now need a way to test for cycles in our graph

Testing for deadlock

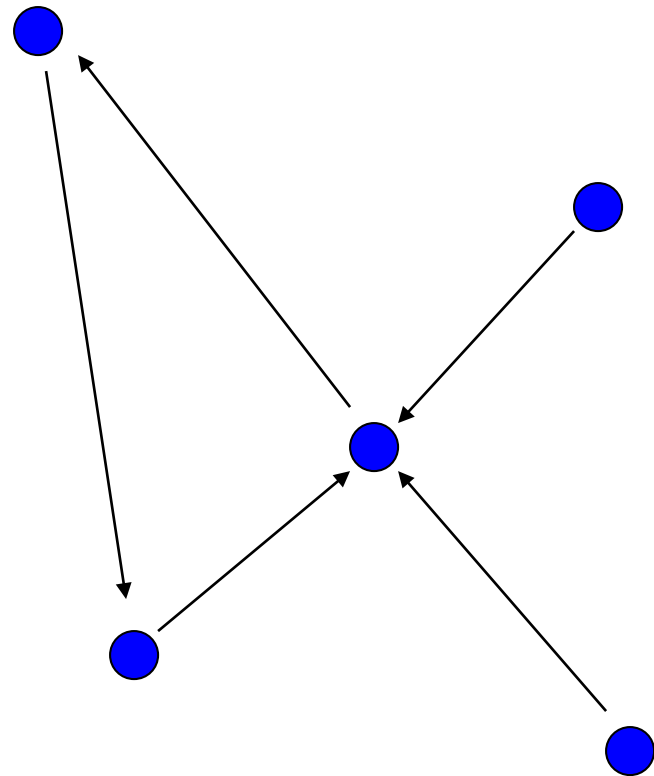
- One way to find cycles
 - Look for a node with no outgoing edges i.e., a process which is not waiting for resources held by others
 - Erase this node, and also erase any edges coming into it
 - Idea: This was a process others might have been waiting for, but it wasn't waiting for anything else to finish off
 - If (and only if) the graph has no cycles, we'll eventually be able to erase the whole graph!
- This is called a graph reduction algorithm

Graph reduction example



Graph reduction example

- This is an example of an “irreducible” graph
- It contains a cycle and represents a deadlock, although only some processes are in the cycle
- Once deadlock occurs, no of processes in the deadlocked state grows over the time...

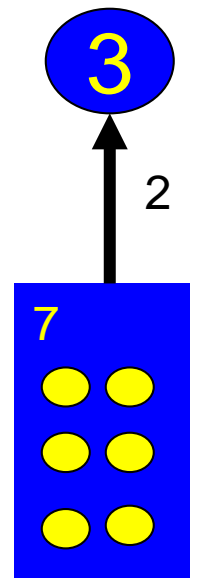
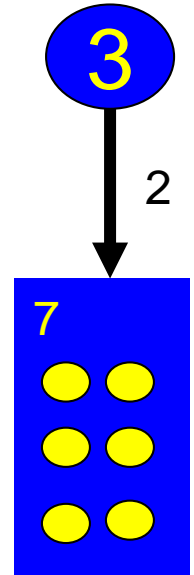


What about “resource” waits?

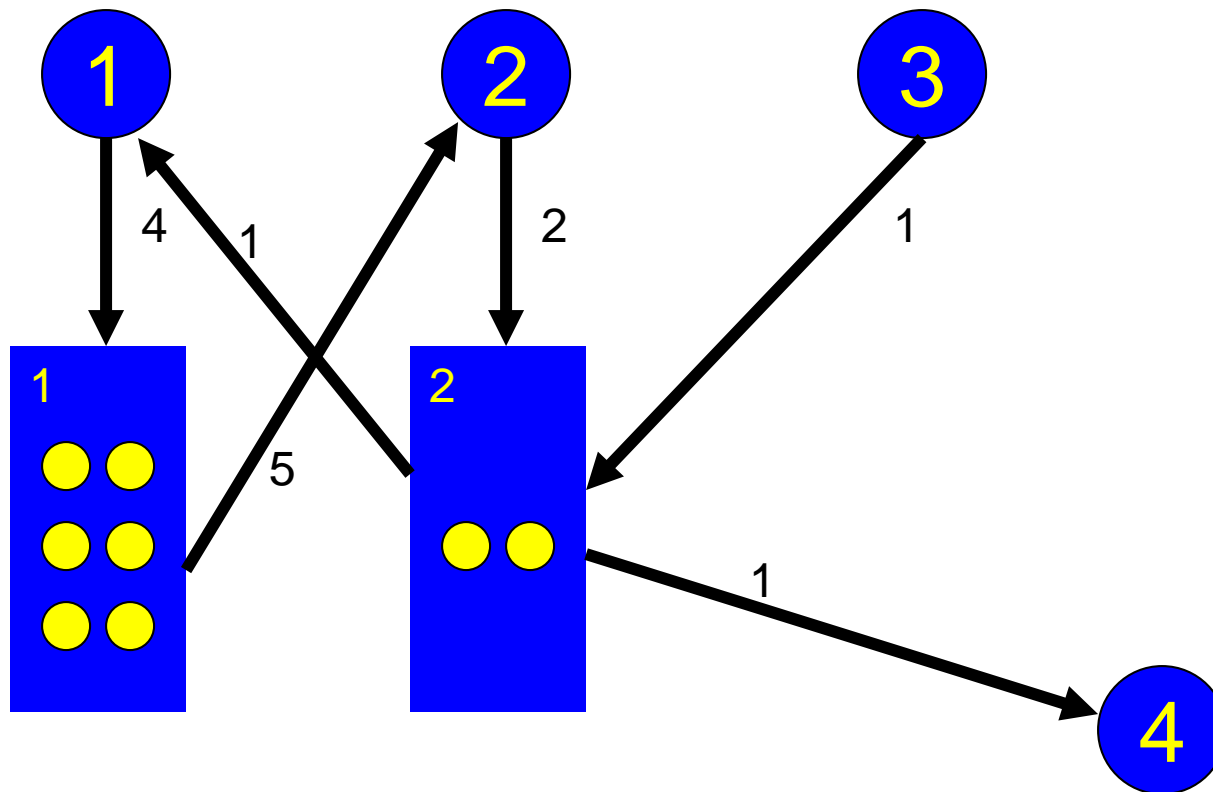
- When dining philosophers wait for one-another, they don't do so directly
- Instead, they wait for resources
- Can we extend our graphs to represent resource wait?

Resource-wait graphs

- We'll use two kinds of nodes
- A process: P_3 will be represented as circle:
- A resource: R_7 will be represented as rectangle:
 - A resource often has multiple identical units, such as “blocks of memory”
 - Represent these as circles in the box
- Arrow from a process to a resource: “I want k units of this resource.”
 - P_3 wants 2 units of R_7
- Arrow to a process: “I hold k units of this resource”
 - P_3 holds 2 units of R_7



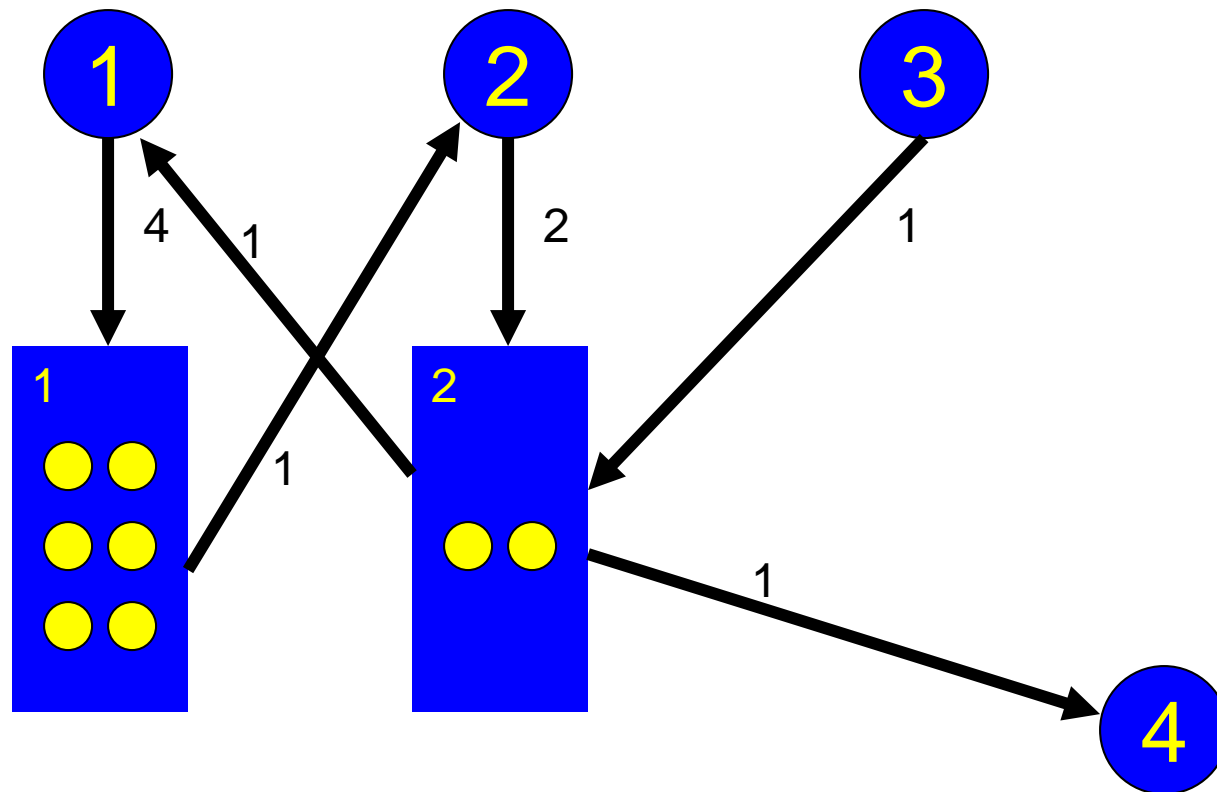
Resource-wait graphs



Reduction rules?

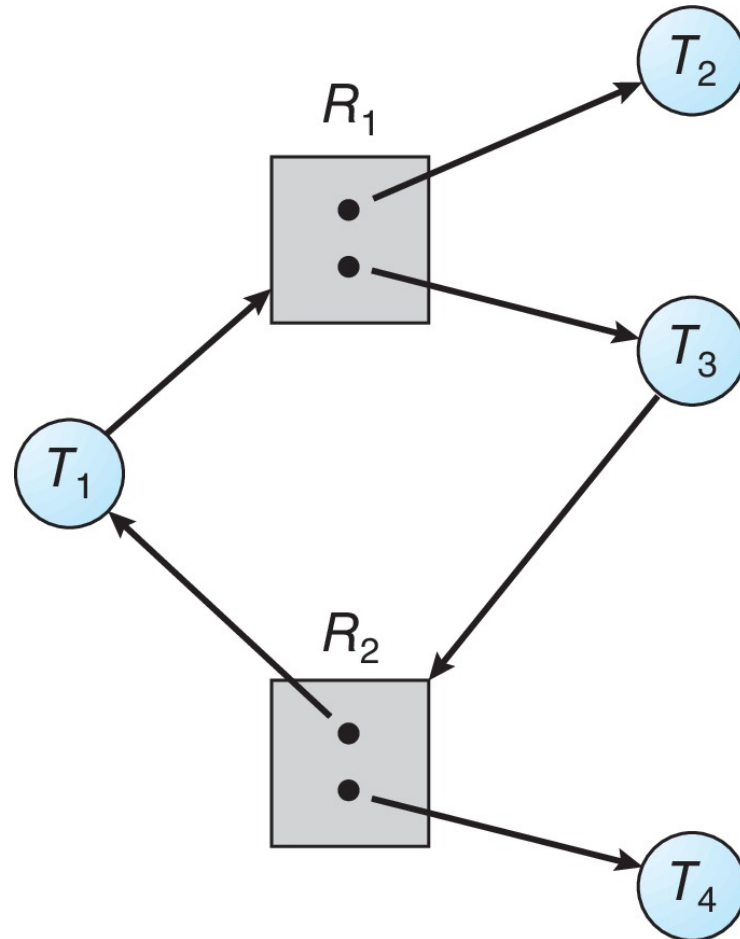
- Find a process that can have all its current requests satisfied (e.g. the “available amount” of any resource it wants is at least enough to satisfy the request)
- Erase that process (in effect: grant the request, let it run, and eventually it will release the resource)
- Continue until we either erase the graph or have an irreducible component. In the latter case we might have identified a deadlock.

This graph is reducible: The system is not deadlocked

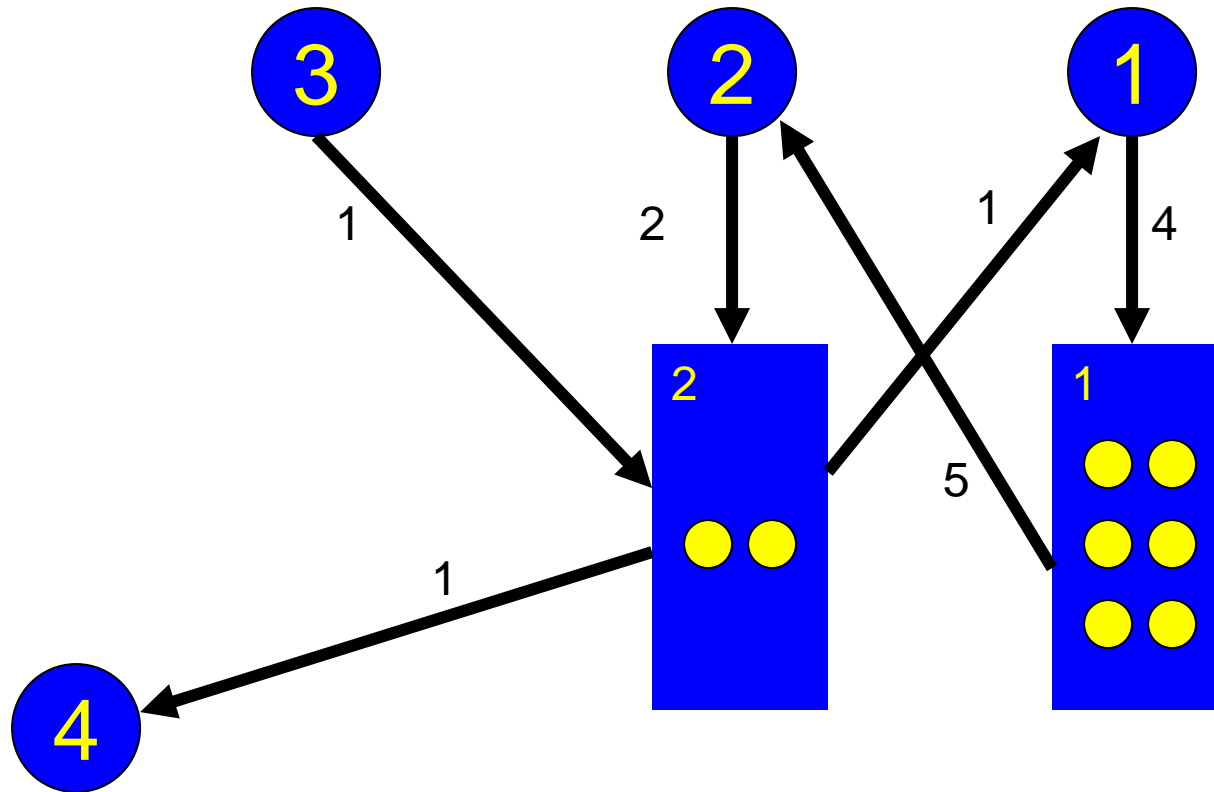


(P4, P3, P1, P2) or (P1, P4, P3, P2) or ...

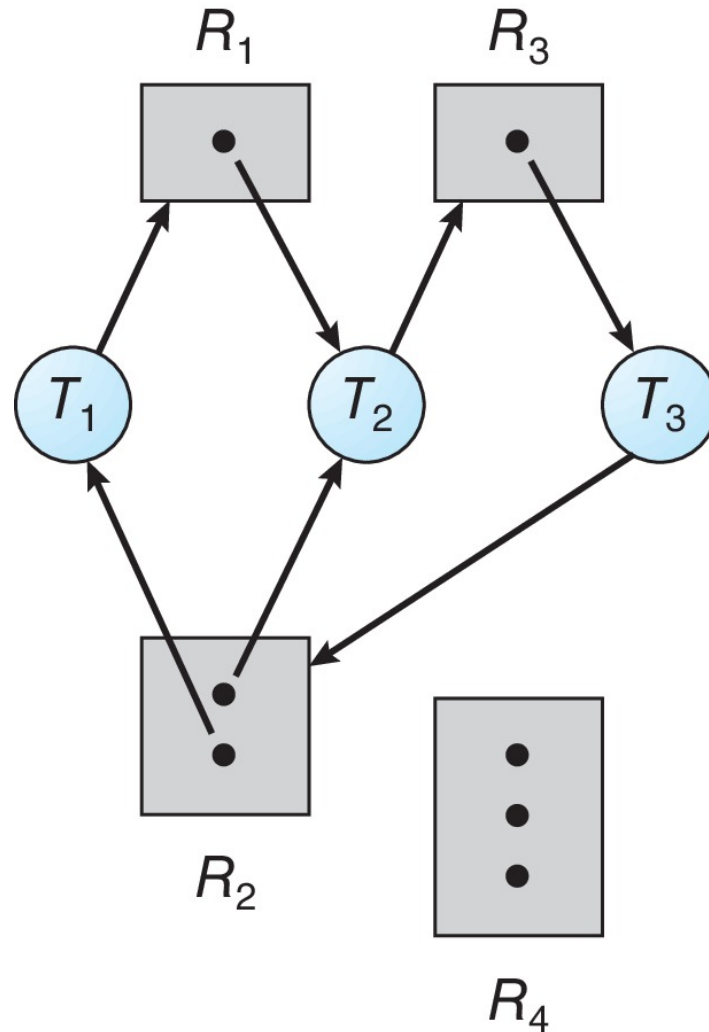
Another Graph With A Cycle But No Deadlock!



This graph is not reducible: The system is deadlocked



This graph is not reducible: The system is deadlocked



Comments

- If graph contains no cycles \Rightarrow no deadlock
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock
 - if several instances per resource type, possibility of deadlock
- It isn't common for systems to actually implement this kind of test
- However, we'll later use a version of the resource reduction graph as part of an algorithm called the "Banker's Algorithm"
- Idea is to schedule the granting of resources so as to avoid potentially deadlock states

Some questions you might ask

- Does the order in which we do the reduction matter?
 - Answer: No. The reason is that if a node is a candidate for reduction at step i , and we don't pick it, it remains a candidate for reduction at step $i+1$
 - Thus eventually, no matter what order we do it in, we'll reduce by every node where reduction is feasible

Some questions you might ask

- If a system is deadlocked, could this go away?
 - No, unless someone kills one of the threads or something causes a process to release a resource
 - Many real systems put time limits on “waiting” precisely for this reason.
 - When a process gets a timeout exception, it gives up waiting and this also can eliminate the deadlock
 - But that process may be forced to terminate itself because often, if a process can't get what it needs, there are no other options available!

Some questions you might ask

- Suppose a system isn't deadlocked at time T .
- Can we assume it will still be free of deadlock at time $T+1$?
 - No, because the very next thing it might do is to run some process that will request a resource...
 - ... establishing a cyclic wait
 - ... and causing deadlock

Methods for Handling Deadlocks

1. Reactive Approaches:

- Periodically check for evidence of deadlock
 - For example, using a graph reduction algorithm
- Then need a way to recover
 - Could blue screen and reboot the computer
 - Could pick a “victim” and terminate that thread
 - But this is only possible in certain kinds of applications
 - Basically, thread needs a way to clean up if it gets terminated and has to exit in a hurry
 - Often thread would then “retry” from scratch
- Despite drawbacks, database systems do this

Dealing with Deadlocks

2. Proactive Approaches:

– Deadlock Prevention

- OS ensures that one of the 4 necessary conditions can't hold
- This will prevent deadlock from occurring by constraining how requests for resources are made

– Deadlock Avoidance

- OS carefully allocates resources based on **future knowledge of resource requests** by processes
- So, deadlocks are avoided

Dealing with Deadlocks

3. Ignore the problem

- OS pretends deadlocks will never occur
- Cheapest solution
- Ostrich approach... but surprisingly common!
- Linux and Windows do the same!!
- Application programs then need to handle the deadlocks

Deadlock Prevention

Deadlock Prevention

- Can the OS prevent deadlocks?
- Prevention: Negate one of necessary conditions
 - Mutual exclusion:
 - Make resources sharable
 - Not always possible (spooling, locks?)
 - Hold and wait
 - Do not hold resources when waiting for another
 - ⇒ Request all resources before beginning the execution
 - ☞ Processes do not know what all they will need
 - ☞ Starvation (if waiting on many popular resources)
 - ☞ Low utilization (Need resource only for a bit)
 - Alternative: Release all resources before requesting anything new
 - Still has the last two problems

Deadlock Prevention

- Prevention: Negate one of necessary conditions
 - No preemption:
 - Make resources preemptable (2 approaches)
 - 1: Preempt requesting processes' current resources if all new resources are not available and let it re-acquire old and newly requesting resources
 - 2: Preempt resources of already waiting processes with the resources to satisfy request of current process
 - Good when easy to save and restore state of resource
 - CPU registers, memory space
 - Not good for printers, tape drives, locks
 - Bad if in middle of critical section and resource is a lock
 - Circular wait:
 - Impose partial ordering on resources, request them in order

Breaking Circular Wait

- Order resources (R_1, R_2, \dots) based on 1-to-1 $F: R \rightarrow N$
- Acquire resources in strictly increasing/decreasing order of enumeration
- Protocol for increasing scenario:
 - A process holding resource R_i can request R_j iff $F(R_j) > F(R_i)$
 - Or a process requesting R_j must have released all R_i s such that $F(R_i) \geq F(R_j)$
 - When requests to multiple instances of same resource, make the request a single operation

Breaking Circular Wait

- ☞ Ordering not always possible
 - ☞ Low resource utilization
 - ☞ It is up to application developers to write programs that follow the ordering!
 - **Witness** s/w tool in FreeBSD works as lock-order verifier and gives warnings when locks are acquired out-of-order
 - Example: T1 acquires lock1 and lock2 in that order. Then Witness records this order and gives warning if T2 tries to acquires lock2 first and then lock1 next
 - Note that imposing lock ordering does not guarantee deadlock prevention if locks can be acquired dynamically
 - Refer Figure 7.5 in Galvin textbook (9th ed) for an example program: Deadlock example with lock ordering for funds transfer between two accounts
- T1 → transaction(checking account, savings account, 10);
T2→ transaction(savings account, checking account, 20);

Two phase locking

- Acquire all resources before doing any work. If any is locked, release all, wait a little while, and retry

```
print_file {  
    do{  
        release all Resources;  
        tryLock(file);  
        tryAcquire(printer);  
        tryAcquire(disk);  
    } while(all not acquired);  
    read file from disk;  
    print file on printer;  
    release all Resources  
}
```

- Pro: dynamic, simple, flexible
- Con:
 - Could “spin” endlessly
 - How does cost grow with number of resources?
 - Hard to know what locks are needed a priori

Deadlock Avoidance

Deadlock Avoidance

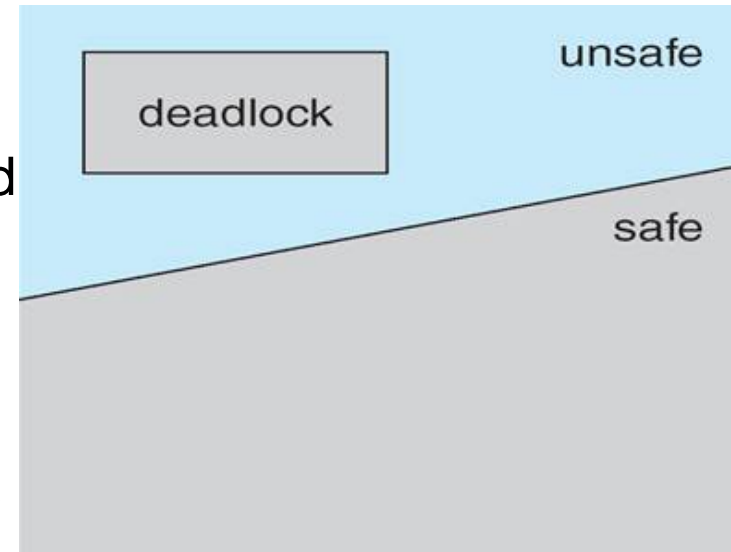
- The key here is that underlying 4 conditions (Mutex, hold & wait, no preemption, and circular wait) are still there, but you do something more clever to avoid getting into a deadlock
- It requires knowing future information
 - Max resource requirement of each process before they execute
- Can we guarantee that deadlocks will never occur?
- Avoidance Approach:
 - Before granting resource to a process, check if system state is **safe**
 - If the state is safe \Rightarrow no deadlock & grant the requested resource to the process

Safe State

- A state is said to be **safe**, if it has a sequence $\{P_1, P_2, \dots, P_n\}$ of ALL processes, such that for each P_i , the resources that P_i can still request can be satisfied by the currently available resources + the resources held by all P_j , where $j < i$
- That is:
 - If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
 - When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on
- State is safe because OS can definitely avoid deadlock
 - by blocking any new requests until safe order is executed
- This avoids circular wait condition
 - Process waits until safe state is guaranteed

Safe State

- If a system is in **safe state** → **no deadlocks**
 - For any possible sequence of resource requests, there exists at least one safe sequence of processing the requests that would succeed in granting all pending and future requests
- If a system is in **unsafe state** → **possibility of deadlock**
 - There is at least one sequence of future requests that leads to deadlock no matter what processing order is tried
- In a deadlocked state, the system has at least one deadlock.
- **Avoidance** → ensure that the system will never enter an **unsafe state**
 - **Conservative approach**



Safe State Example

- Suppose there are 12 tape drives in a system

	<u>max need</u>	<u>current usage</u>	<u>could ask for</u>
P0	10	5	5
P1	4	2	2
P2	9	2	7

3 drives available for granting

- Current state is safe because a safe sequence exists:

$\langle p1, p0, p2 \rangle$

p1 can complete with currently available resources (2 of 3)

p0 can complete with current (5) + available (5) after p1

p2 can complete with current (2) + available: 10 (after p1, p0)

- If p2 requests 1 drive
 - then it must wait to avoid unsafe state.

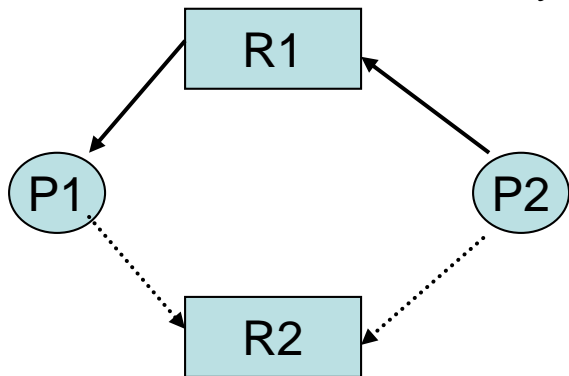
Deadlock Avoidance Algorithms

- Single instance of a resource type
 - Use a resource-allocation graph
- Multiple instances of a resource type
 - Use the Banker's Algorithm

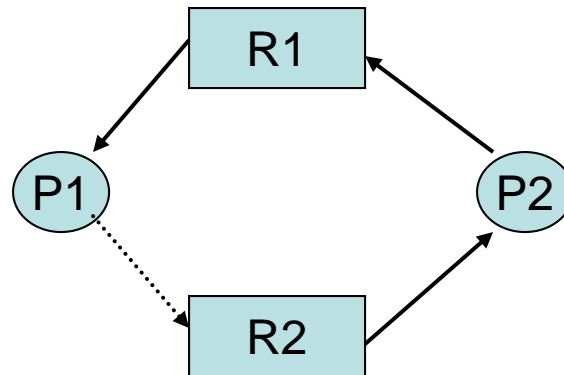
These are less conservative than Acquire-all/Release-all approach seen earlier in case of Two Phase Locking

Res. Alloc. Graph Algorithm

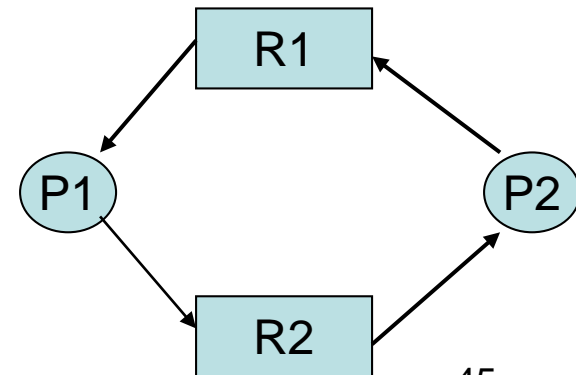
- Deadlock can be described using a *resource allocation graph, RAG*
- Works if only **one** instance of each resource type
- Algorithm:
 - Add a **claim edge**, $P_i \rightarrow R_j$ if P_i can request R_j in the future
 - Represented by a dashed line in graph
 - A request $P_i \rightarrow R_j$ can be granted only if:
 - Adding an **assignment edge** $R_j \rightarrow P_i$ does not introduce cycles
 - Since cycles imply unsafe state



Safe State



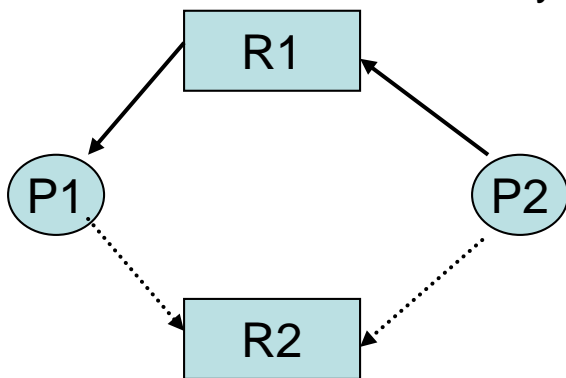
Unsafe State



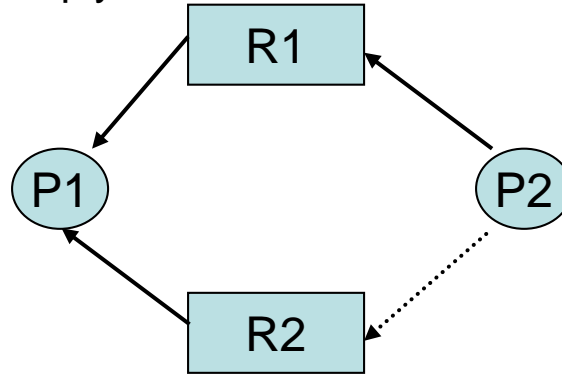
Deadlock State

Res. Alloc. Graph Algorithm

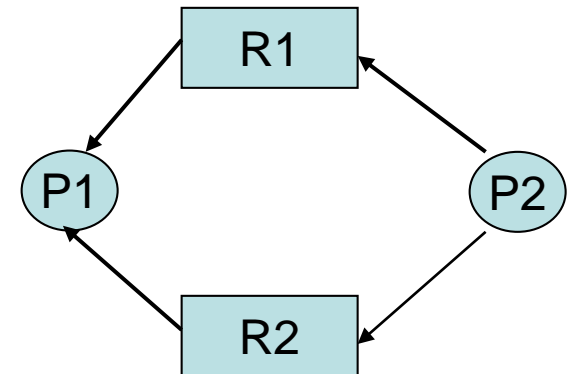
- Deadlock can be described using a *resource allocation graph, RAG*
- Works if only **one** instance of each resource type
- Algorithm:
 - Add a **claim edge**, $P_i \rightarrow R_j$ if P_i can request R_j in the future
 - Represented by a dashed line in graph
 - A request $P_i \rightarrow R_j$ can be granted only if:
 - Adding an **assignment edge** $R_j \rightarrow P_i$ does not introduce cycles
 - Since cycles imply unsafe state



Safe State



Safe State



Safe State

Banker's Algorithm

- Dijkstra developed it to improve on the performance of Acquire-all approach
- Suppose we know the “worst case” aka MAX resource needs of processes in advance
 - A bit like knowing the credit limit on your credit cards.
(This is why they call it the Banker's Algorithm)
 - Ensures that bank never allocates its cash reserves in such a way that it could no longer satisfy the needs of all customers.
- Observation: Suppose we just give some process ALL the resources it could need...
 - Then it will execute to completion.
 - After which it will give back the resources.

Banker's Algorithm

- So...
 - A process pre-declares its worst-case (MAX) resource needs
 - Then it asks for what it “really” needs, a little at a time and it may release some of the resources acquired
 - The algorithm decides when to grant requests by assigning resources
- It delays a request to avoid deadlocks unless:
 - It can find a sequence of processes...
 - such that it could grant their outstanding need...
 - ... so they would terminate...
 - ... letting it collect their resources...
 - ... and in this way it can execute everything to completion!

Banker's Algorithm

- How will it really do this?
 - The algorithm will just implement the graph reduction method for resource graphs
 - Graph reduction is “like” finding a sequence of processes that can be executed to completion
- So: given a request
 - Build a resource graph
 - See if it is reducible, only grant request if so
 - Else must delay the request until someone releases some resources, at which point can test again

Banker's Algorithm

- Toward right idea:
 - State maximum resource needs in advance, but each job may not require all of these Max no. of resources at a time to proceed
- Banker's algorithm (less conservative):
 - Allocate resources dynamically as and when a job needs
 - Evaluate each request and grant if some ordering of processes is still deadlock free afterward
 - **Technique: pretend each request is granted, then run safety algorithm,**
 - Check for $([Max_{node}] - [Alloc_{node}] \leq [Avail])$
 - **Grant request if result is deadlock free (conservative!)**
 - Keeps system in a "SAFE" state, i.e. there exists at least one safe sequence $\{P_1, P_2, \dots P_n\}$ with P_1 requesting all remaining resources, finishing & releasing its resources, then P_2 requesting all remaining resources, etc..
 - Algorithm allows the sum of maximum resource needs of all current processes to be greater than the total resources



Banker's Algorithm

- Decides whether to grant a resource request.
- Data structures:

n : integer	# of processes/threads
m : integer	# of resource types
$available[1..m]$	$available[i]$ is # of avail resources of type i
$max[1..n, 1..m]$	max demand of each P_i for each R_i
$allocation[1..n, 1..m]$	current allocation of resource R_j to P_i
$need[1..n, 1..m]$	max # resource R_j that P_i may still request
	$need_i = max_i - allocation_i$

let $request[i]$ be vector of # of resources of type j ($1 \leq j \leq m$)
process P_i wants

Basic Algorithm

1. If $\text{request}[i] > \text{need}[i]$ then
error (asked for too much)
2. If $\text{request}[i] > \text{available}[i]$ then
wait (can't supply it now)
3. Resources are available to satisfy the request
Let's assume that we satisfy the request. Then we would have:
$$\text{available} = \text{available} - \text{request}[i]$$
$$\text{allocation}[i] = \text{allocation}[i] + \text{request}[i]$$
$$\text{need}[i] = \text{need}[i] - \text{request}[i]$$

Now, check if this would leave system in a **safe state**.
if yes, grant the request for resources,
if no, then leave the state as is and cause process to wait.

Safety Check Algorithm

free[1..m] = available /* how many resources are available */
finish[1..n] = false (for all processes i) /* none finished yet */

Step 1: Find a process i such that finish[i]=false and need[i] <= free
/* find a process that can complete its request now */
if no such i exists, go to step 3 /* we're done */

Step 2: Found a process i:
finish [i] = true /* done with this process */
free = free + allocation [i]
/* assume this process is finished and add its allocation back
to the available list */
go to step 1

Step 3: If finish[i] = true for all i, the system is safe. Else Not

Example: Memory Page Allocation

- Consider a system having 8 memory pages
- Three processes
 - A requires MAX of 4 pages
 - B requires MAX of 5 pages
 - C requires MAX of 5 pages

Banker's Algorithm: Example

	<u>Allocation</u>				<u>Max</u>				<u>Available</u>		
	A	B	C		A	B	C		A	B	C
P0	0	1	0		7	5	3		3	3	2
P1	2	0	0		3	2	2				
P2	3	0	2		9	0	2				
P3	2	1	1		2	2	2				
P4	0	0	2		4	3	3				

this is a safe state: safe sequence <P1, P3, P4, P2, P0>

Suppose that P1 requests (1,0,2)

- add it to P1's allocation and subtract it from Available

Banker's Algorithm: Example

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	5	3	2	3	0
P1	3	0	2	3	2	2			
P2	3	0	2	9	0	2			
P3	2	1	1	2	2	2			
P4	0	0	2	4	3	3			

This is still safe: safe seq <P1, P3, P4, P0, P2>

In this new state,

P4 requests (3,3,0)

not enough available resources

P0 requests (0,2,0)

let's check resulting state

Banker's Algorithm: Example

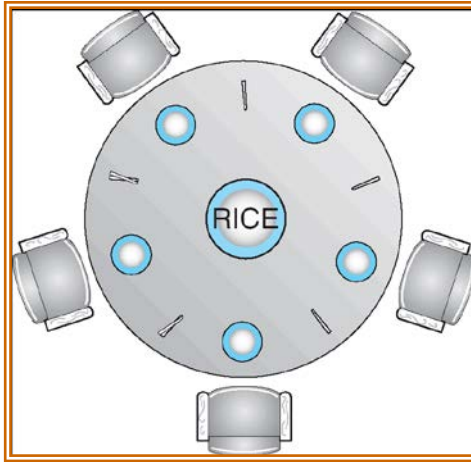
	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P0	0	3	0	7	5	3	2	1	0
P1	3	0	2	3	2	2			
P2	3	0	2	9	0	2			
P3	2	1	1	2	2	2			
P4	0	0	2	4	3	3			

This is unsafe state (why?)

So P0's request will be denied

Problems with Banker's Algorithm?

Banker's Algorithm Example



- Banker's algorithm with dining philosophers where a philosopher can pick any two of available chopsticks from the Center
 - “Safe” (won't cause deadlock) if when try to grab chopstick either:
 - Not last chopstick or
 - Is last chopstick but someone will have two afterwards
 - What if k-handed lawyers? Don't allow if:
 - It's the last one, and no one would have k
 - It's 2nd to last, and no one would have k-1
 - It's 3rd to last, and no one would have k-2
 - ...



Deadlocks: Part II

Detection and Recovery

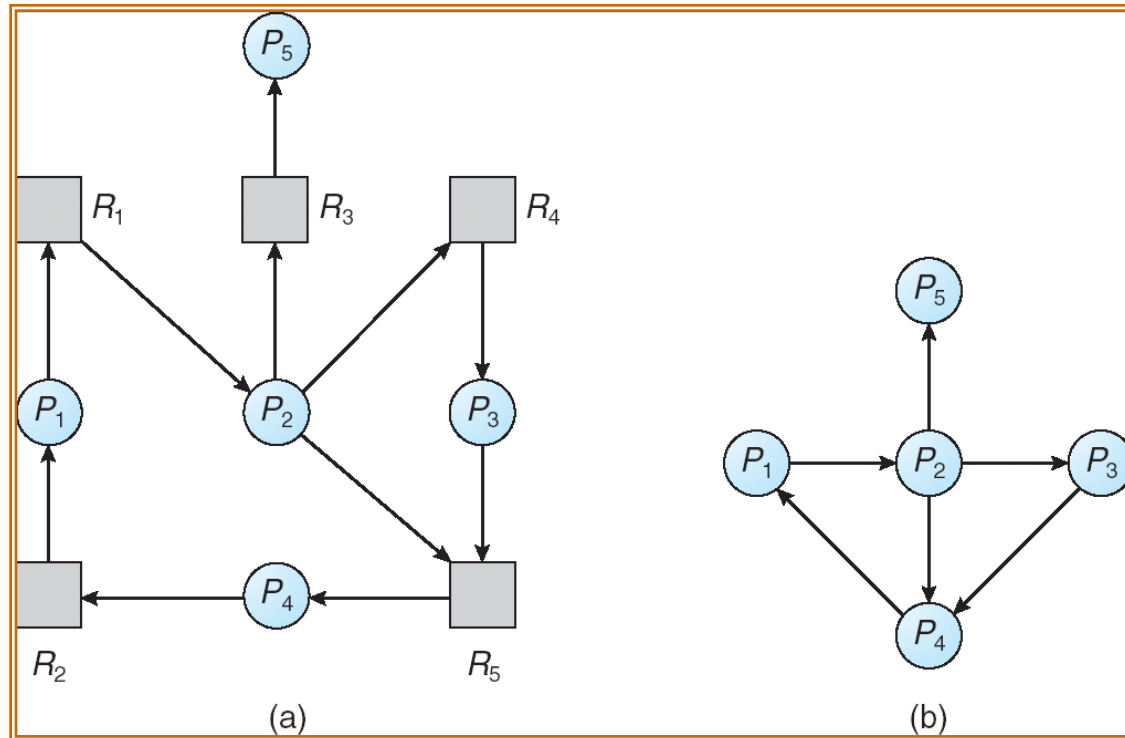
Deadlock Detection & Recovery

- If neither avoidance or prevention is implemented, deadlocks can (and will) occur.
- Coping with this requires:
 - Detection: finding out if deadlock has occurred
 - Keep track of resource allocation (who has what)
 - Keep track of pending requests (who is waiting for what)
 - Recovery: untangle the mess.
- Expensive to detect, as well as to recover

Using the RAG Algorithm to detect deadlocks

- Suppose there is only one instance of each resource
- Example 1: Is this a deadlock?
 - P1 has R2 and R3, and is requesting R1
 - P2 has R4 and is requesting R3
 - P3 has R1 and is requesting R4
- Example 2: Is this a deadlock?
 - P1 has R2, and is requesting R1 and R3
 - P2 has R4 and is requesting R3
 - P3 has R1 and is requesting R4
- Use a **wait-for graph**:
 - Collapse resources
 - An edge $P_i \rightarrow P_k$ exists only if RAG has $P_i \rightarrow R_j$ & $R_j \rightarrow P_k$
 - Cycle in wait-for graph \Rightarrow deadlock!

Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding wait-for graph

2nd Detection Algorithm

- What if there are multiple resource instances?
- Data structures:

n : integer # of processes

m : integer # of resources

$available[1..m]$ $available[i]$ is # of avail resources of type i

$request[1..n, 1..m]$ max current demand of each P_i for each R_i

$allocation[1..n, 1..m]$ current allocation of resource R_j to P_i

$finish[1..n]$ true if P_i 's request can be satisfied

let $request[i]$ be vector of # instances of each resource P_i wants

2nd Detection Algorithm

1. `work[] = available[]`
for all $i < n$, if `allocation[i] \neq 0` //not yet finished
then `finish[i] = false` else `finish[i] = true`
2. find an index i such that:
 `finish[i] = false;`
 `request[i] \leq work`
if no such i exists, go to Step 4.
3. `work = work + allocation[i]`
 `finish[i] = true`, go to Step 2
4. if `finish[i] = false` for some i ,
 then system is deadlocked with P_i in deadlock

Example

Finished = {F, F, F, F};

Work = Available = (0, 0, 1);

	R ₁	R ₂	R ₃
P ₁	1	1	1
P ₂	2	1	2
P ₃	1	1	0
P ₄	1	1	1

Allocation

	R ₁	R ₂	R ₃
P ₁	3	2	1
P ₂	2	2	1
P ₃	0	0	1
P ₄	1	1	1

Request

Example

Finished = {F, F, T, F};

Work = (1, 1, 1);

	R_1	R_2	R_3
P_1	1	1	1
P_2	2	1	2
P_3	1	1	0
P_4	1	1	1

Allocation

	R_1	R_2	R_3
P_1	3	2	1
P_2	2	2	1
P_3			
P_4	1	1	1

Request

Example

Finished = {F, F, T, T};

Work = (2, 2, 2);

	R ₁	R ₂	R ₃
P ₁	1	1	1
P ₂	2	1	2
P ₃	1	1	0
P ₄	1	1	1

Allocation

	R ₁	R ₂	R ₃
P ₁	3	2	1
P ₂	2	2	1
P ₃			
P ₄			

Request

Example

Finished = {F, T, T, T};

Work = (4, 3, 4);

	R ₁	R ₂	R ₃
P ₁	1	1	1
P ₂	2	1	2
P ₃	1	1	0
P ₄	1	1	1

Allocation

	R ₁	R ₂	R ₃
P ₁	3	2	1
P ₂			
P ₃			
P ₄			

Request

Example

Finished = {T, T, T, T};

Work = (4, 3, 4);

	R ₁	R ₂	R ₃
P ₁			
P ₂			
P ₃			
P ₄			

Allocation

	R ₁	R ₂	R ₃
P ₁			
P ₂			
P ₃			
P ₄			

Request

When to run Detection Algorithm?

- For every resource request?
- For every request that cannot be immediately satisfied?
- Once every hour?
- When CPU utilization drops below 40%?

Deadlock Recovery

- Killing one/all deadlocked processes
 - Crude, but effective
 - Keep killing processes, until deadlock broken
 - Repeat the entire computation
- Preempt resource/processes until deadlock broken
 - Selecting a victim (# resources held, how long executed)
 - Transactions: Rollback (partial or total) and Retry
 - Starvation (prevent a process from being executed)
- Proceed without a resource!

Java Concurrency Tools

- `java.lang.management.ThreadMXBean.findDeadlockedThreads()`
 - Part of the JVM management API
 - Can be used to detect deadlocks: returns you the threadIDs of threads currently blocking waiting to enter an object (and “ownable synchronizers”)
 - “It might be an expensive operation” - so when would you run it?
- `java.util.concurrent`
 - .Semaphore (and thus mutex)
 - .CountDownLatch & .CountDownLatch
 - .Exchanger (Nice - similar in flavor to co-routines from e.g. Scheme/Lisp)
- `java.util.concurrent.atomic`
 - A toolkit of classes that support lock-free programming (given H/W support):
AtomicInteger aint = new AtomicInteger(42);
aint.compareAndSet(whatIThinkItIs, whatIWantItToBe); //No lock needed..
- ..and more. Check out the following:
 - <http://docs.oracle.com/javase/1.5.0/docs/guide/concurrency/overview.html>
 - <http://docs.oracle.com/javase/tutorial/essential/concurrency/index.html>
 - <http://sourceforge.net/projects/javaconcurrenta/>
 - <https://www.ibm.com/developerworks/java/library/j-jtp10264/>

Summary

- Dining Philosophers Problem
 - Highlights need to multiplex resources
 - Context to discuss starvation, deadlock, livelock
- Four conditions for deadlocks
 - Mutual exclusion
 - Only one thread at a time can use a resource
 - Hold and wait
 - Thread holding at least one resource is waiting to acquire additional resources held by other threads
 - No preemption
 - Resources are released only voluntarily by the threads
 - Circular wait
 - \exists set $\{T_1, \dots, T_n\}$ of threads with a cyclic waiting pattern

Summary (2)

- Techniques for addressing Deadlock
 - Allow system to enter deadlock and then recover
 - Ensure that system will *never* enter a deadlock
 - Ignore the problem and pretend that deadlocks never occur in the system
- Deadlock prevention
 - Prevent one of four necessary conditions for deadlock
- Deadlock avoidance
 - Assesses, for each allocation, whether it has the potential to lead to deadlock
 - Banker's algorithm gives one way to assess this
- Deadlock detection and recover
 - Attempts to assess whether waiting graph can ever make progress
 - Recover if not

Reading Assignment

- Chapter 6 from OSPP by Anderson et al 2nd Edition
- Chapter 7 from OSC by Galvin et al 9th Edition
- Chapter 2 from MOS by Tanenbaum et al