Domain Specific Language	Reference to the paper
Tensor Comprehensions: Framework-Agnostic High-Performance Machine Learning Abstractions [A: Nicolas Vasilache, Oleksandr Zinenko, Theodoros Theodoridis, Priya Goyal, Zachary DeVito, William S. Moses, Sven Verdoolaege, Andrew Adams, Albert Cohen]	https://arxiv.org/abs/1802.04730
Tiramisu: a polyhedral compiler for expressing fast and portable code [A: Riyadh Baghdadi, Jessica Ray, Malek Ben Romdhane, Emanuele Del Sozzo, Abdurrahman Akkas, Yunming Zhang, Patricia Suriana, Shoaib Kamil, Saman Amarasinghe]	https://dl.acm.org/doi/10.5555/3314872.3314896
Diesel: DSL for linear algebra and neural net computations on GPUs [A: Venmugil Elango, Norm Rubin, Mahesh Ravishankar , Hariharan Sandanagobalane, Vinod Grover]	https://dl.acm.org/doi/10.1145/3211346.3211354
Caffe: Convolutional Architecture for Fast Feature Embedding [A: Yangqing Jia, Evan Shelhamer, Jeff Donahue, Sergey Karayev, Jonathan Long, Ross Girshick, Sergio Guadarrama, Trevor Darrell]	https://dl.acm.org/doi/10.1145/2647868.2654889
Halide: Decoupling algorithms from schedules for easy optimization of image processing pipelines [A: Jonathan Ragan-Kelley, Andrew Adams, Sylvain Paris, Marc Levoy, Saman Amarasinghe, Frédo Durand]	https://dl.acm.org/doi/10.1145/2185520.2185528
TVM: an automated end-to-end optimizing	https://dl.acm.org/doi/10.5555/3291168.3291211

compiler for deep learning [A: Tianqi Chen, Thierry Moreau, Ziheng Jiang, Lianmin Zheng, Eddie Yan, Meghan Cowan, Haichen Shen, Leyuan Wang, Yuwei Hu, Luis Ceze, Carlos Guestrin, Arvind Krishnamurthy]	
Latte: a language, compiler, and runtime for elegant and efficient deep neural networks [A: Leonard Truong, Rajkishore Barik, Ehsan Totoni, Hai Liu, Chick Markley, Armando Fox, Tatiana Shpeisman]	https://dl.acm.org/doi/10.1145/2908080.2908105
GraphIT and Universal Graph Framework	http://groups.csail.mit.edu/commit/papers/2020/ajay-sm-thesis.pdf
MLIR Affine dialect	https://mlir.llvm.org/docs/Dialects/Affine/
Iteration graphs in Tensor Algebra Compiler(TACO) It is a compiler that generates optimized code for sparse tensor multiplication. Its uniqueness lies in its ability to generate code that is as performant as Hand-written case-by-case kernels	http://tensor-compiler.org/kjolstad-oopsla17-tensor-compiler.pdf This is a nice talk on TACO: https://www.youtube.com/watch?v=yAtG64qV2nM