

# Network Programming Basics

Kotaro Kataoka

# Socket programming *with TCP*

## client must contact server

- server process must first be running
- server must have created socket (door) that welcomes client's contact

## client contacts server by:

- Creating TCP socket, specifying IP address, port number of server process
- *when client creates socket:* client TCP establishes connection to server TCP

- when contacted by client, *server TCP creates new socket* for server process to communicate with that particular client
  - allows server to talk with multiple clients
  - source port numbers used to distinguish clients (more in Chap 3)

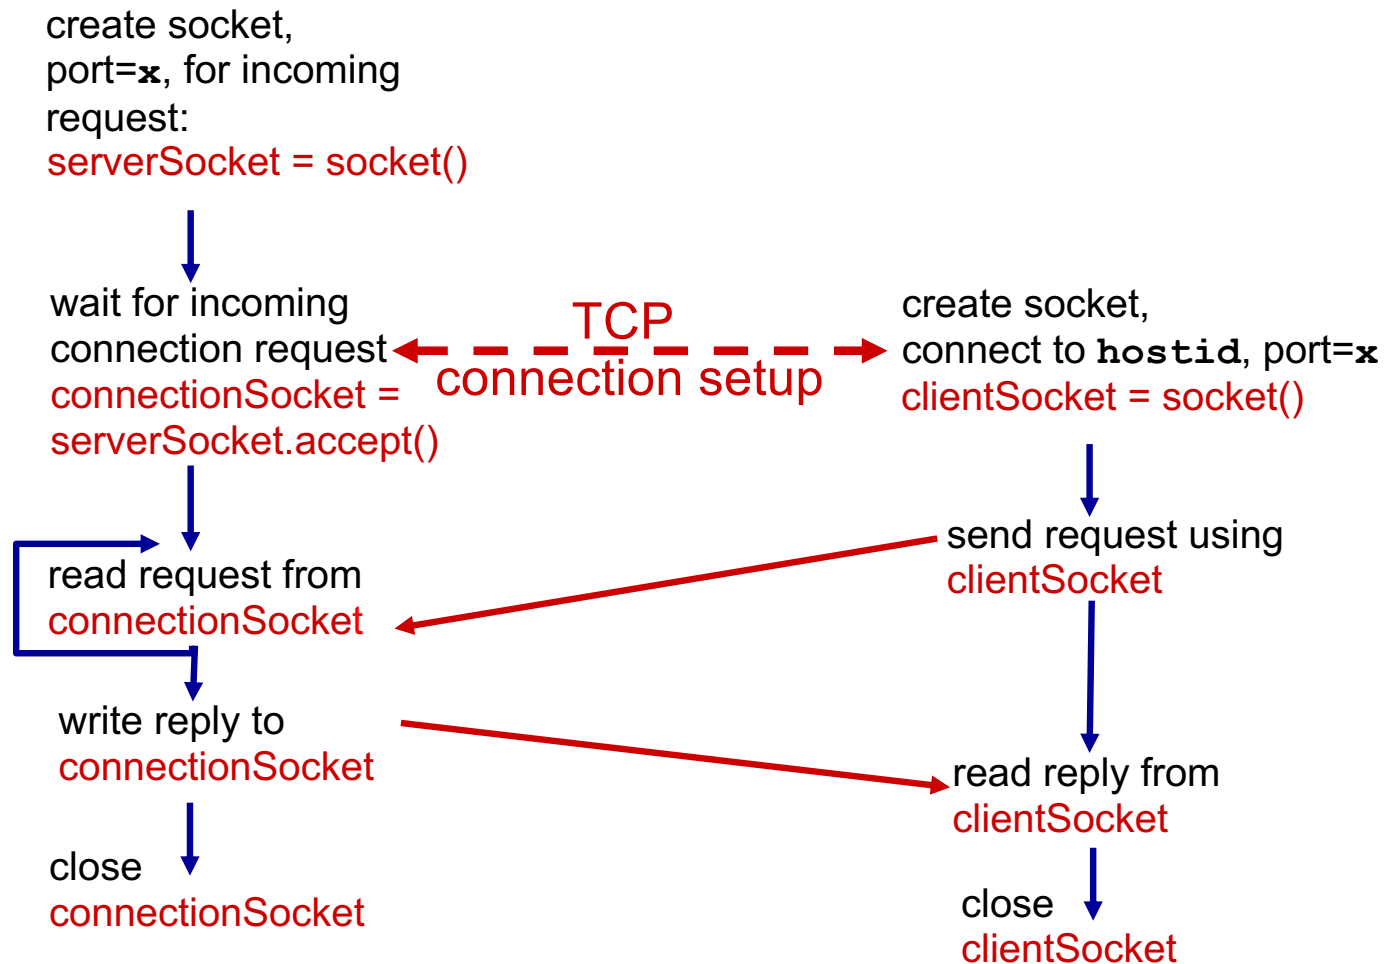
## application viewpoint:

TCP provides reliable, in-order byte-stream transfer (“pipe”) between client and server

# Client/server socket interaction:TCP

server (running on `hostid`)

client



# Echo Client with only one action

1. Program starts with Server **IP Address, Message**  
and **Server Port** `main()`
2. Create socket `socket()`
3. Set server parameters to socket
4. Connect to server `connect()`
5. Send message `send()`
6. Receive message `recv()`
7. Show message echoed by server `printf()`
8. Destroy socket and die `close()`

Let's get started.

# socket() System Call (1/2)

- Creates socket
- `int socket(int family, int type, int proto)`
  - Protocol Family
    - `AF_LOCAL / PF_LOCAL` Host-internal protocols
    - `AF_INET / PF_INET` Internet version 4 protocols
    - `AF_ROUTE / PF_ROUTE` Internal Routing protocol
    - `AF_INET6 / PF_INET6` Internet version 6 protocols
    - etc
  - Socket Type
    - `SOCK_STREAM`
    - `SOCK_DGRAM`
    - `SOCK_RAW`
  - Protocol
    - Normally “0” except the case of RAW

# socket() System Call (2/2)

- Return value
  - Success: socket descriptor
  - Failure: -1
- Example

```
int sd;
sd = socket(AF_INET, SOCK_STREAM, 0);
if (sd < 0) {
    something bad happened...
}
```

# struct sockaddr\_in

- Specification of a local or remote endpoint

```
#include <sys/socket.h>
```

```
#include <netinet/in.h>
```

```
#include <netinet/ip.h> /* superset of previous */
```

```
struct sockaddr_in {
```

```
    sa_family_t    sin_family; /* address family: AF_INET */
```

```
    in_port_t      sin_port;   /* port in network byte order */
```

```
    struct in_addr  sin_addr;   /* internet address */
```

```
};
```

```
/* Internet address. */
```

```
struct in_addr {
```


```
    uint32_t        s_addr;     /* address in network byte order */
```

```
};
```



# inet\_pton() Function

- Returns IP address in text format to that in binary format
  - Endianness
  - Network byte order (Big-endian)



# connect() System Call

- Initiate TCP connection on a socket
- Set the server address

```
struct sockaddr_in servAddr;  
memset(&servAddr, 0, sizeof(servAddr));  
servAddr.sin_family = AF_INET;  
int err = inet_pton(AF_INET, servIP, &servAddr.sin_addr.s_addr);  
if (err <= 0) {  
    perror("inet_pton() failed");  
    exit(-1);  
}  
servAddr.sin_port = htons(servPort);
```

- Connect to server

```
if (connect(sockfd, (struct sockaddr *) &servAddr, sizeof(servAddr)) < 0) {  
    perror("connect() failed");  
    exit(-1);  
}
```

# Echo Server

- |   |          |
|---|----------|
| 1. Program starts with <b>Server Port</b> | main()   |
| 2. Create socket                          | socket() |
| 3. Set server parameters to socket        | bind()   |
| 4. Wait for client                        | listen() |
| 5. Establish TCP connection               | accept() |
| 6. Receive message                        | recv()   |
| 7. Send back message (ECHO)               | send()   |
| 8. Repeat 5 to 7                          |          |

# bind() System Call

- Called on the server
- Bind a socket with a specific endpoint (me!!)
- `int bind(int sockfd, struct sockaddr *addr, int addrlen);`
- Example

```
struct sockaddr_in server;
```

```
memset((void *)&server, 0, sizeof(server));
```

```
server.sin_family = AF_INET;
```


```
server.sin_port = htons(12345);
```

```
server.sin_addr.s_addr = INADDR_ANY;
```

```
bind(sd, (struct sockaddr *)&server, sizeof(server));
```


# Creating Socket and Binding on Server

- `socket()` is called



# Creating Socket and Binding on Server

- Bind to listening port on server



```
[sfc-cpu:7:29] netstat -an
```

Active Internet connections (including servers)

Proto	Recv-Q	Send-Q	Local Address	Foreign Address	(state)
udp4	0	0	*.12345	*.*	

# listen() System Call

- int  
listen(int socket, int backlog);
- Listen on a socket and wait for a connection

# accept() System Call

- int  
accept(int socket, struct sockaddr \*restrict  
address, socklen\_t \*restrict address\_len);
- Accept a connection on a socket
- “address” contains the address of connecting host (i.e. client)



# send() and recv() System Call

- ssize\_t  
send(int socket, const void \*buffer, size\_t length, int flags);
- ssize\_t  
recv(int socket, void \*buffer, size\_t length, int flags);
- send() and recv() do not have an argument to store address information because they will be called after establishing connection (from and to are both known)

# Echo Client with Loop Action with BYE

- Extend the simple echo client
- You can repeat to type message
- You can stop your client by command “BYE”
- Program starts with Server **IP Address** and **Server Port**

# Echo Server with Multiple Clients

- Server serves for multiple clients
- Non blocking
  - Do not wait for the interrupt from a specific input
  - Use `select()`

# DNS?

GetAddrInfo.c

# addrinfo

- Used to store the result of DNS lookup from resolver

```
struct addrinfo {  
    int             ai_flags;  
    int             ai_family;  
    int             ai_socktype;  
    int             ai_protocol;  
    size_t          ai_addrlen;  
    struct sockaddr *ai_addr;  
    char            *ai_canonname;  
    struct addrinfo *ai_next;  
};
```

# getaddrinfo() (1/2)

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
```

```
struct addrinfo hints;           // Give initial hints to resolver
struct addrinfo *result, *rp;    // Result will be stored here
```

```
// Prepare the hints for resolver
memset(&hints, 0, sizeof(struct addrinfo));
hints.ai_family = AF_UNSPEC;      // IPv4 or v6 (You can specify either)
hints.ai_socktype = SOCK_STREAM; // Hardcoded TCP as dummy
hints.ai_protocol = IPPROTO_TCP; // Hardcoded TCP as dummy
hints.ai_flags = AI_CANONNAME;   // I want canonical name!!
```

# getaddrinfo() (2/2)

```
int s = getaddrinfo(hostName, NULL, &hints, &result);
```

```
for (rp = result; rp != NULL; rp = rp->ai_next) {  
    SHOW CORRESPONDING IPv4/v6 ADDRESS;  
    or  
    ATTEMPT connect() to IPv4/v6 ADDRESS;  
}
```

- getaddrinfo() takes hostname/address and/or port number as the 1<sup>st</sup> and 2<sup>nd</sup> arguments.

# Knowing the address family from rp

```
struct sockaddr_in *saln4;  
struct sockaddr_in6 *saln6;  
char addrString[INET6_ADDRSTRLEN];  
memset(addrString, 0, sizeof(addrString));  
  
switch (rp->ai_family) {  
case AF_INET:  
    saln4 = (struct sockaddr_in *) rp->ai_addr;  
    inet_ntop(rp->ai_family, &saln4->sin_addr.s_addr, addrString, sizeof(addrString));  
    break;  
  
case AF_INET6:  
    saln6 = (struct sockaddr_in6 *) rp->ai_addr;  
    inet_ntop(rp->ai_family, &saln6->sin6_addr.s6_addr, addrString, sizeof(addrString));  
    break;  
  
...
```



# Assignment

# Compulsory Assignment:

## Extending echo client/server functions

Q-1. Integrate `getaddrinfo()` as part of your client software so that the hostname of the server can be given as the command line option. This feature must be demonstrated by showing 1) the screenshot of successful execution of your client software with hostname, and 2) the screenshot of wireshark/tcpdump to prove that your client software sends a DNS query and receives a response.

Q-2. Add at least two features to Echo Client /Server, and demonstrate them. In the report, you must describe the new features with their benefit. Significance of the features will impact the marks given.

### Example Marks

- 1 Adding or recording a timestamp of a message exchange
- 5 Supporting IPv6 (with proper network configuration)
- 10 Supporting IPv6 and multicast (with extra clients)
- 10 Supporting SSL (with proper certificate configuration)

# Implementation Guidance

- Individual Assignment
- You may choose any programming language (C, JAVA, Python, Perl and etc.)
- The software must be based on Socket Programming.
- Wrappers of API must not be used (messaging etc). Use send/recv or read/write using TCP socket.
- Keep the record of **Reference. Which web site, documents, or textbook did you refer?**

# Submission Guidance

- Deadlines
  - Demo to TAs: December 12<sup>th</sup> and 13<sup>th</sup>
  - Submission of Report to Google Classroom: December 14<sup>th</sup>
- Requirement of Report
  - The core idea of your answer to each question. Better visibility like screenshot of application will be appreciated.
  - **Reference (web sites, books, etc.) which corresponds to each answer.** That's why, the reference is important. **Answers without appropriate reference may not get mark.**
  - Human readable source code as separate files
  - README as a separate file so that TAs and instructor can compile source code and execute the binary anytime
  - Submit all files as one tar ball or zip ball.

# Project: Networked Software

- Developing an application (routing protocol / server-client / P2P / Blockchain, etc.)
- Where did that idea come from?
  - Papers, blogs, your own finding as inspirations?
  - Does anybody else attempt to develop a similar application?
- Group Project
  - 4 to 8 members per group
  - Big or small or a group is not considered as evaluation criteria.
  - Individual contribution is not considered. Same marks will be offered to all members in the group.
- You MUST NOT re-use a project which has been or is going to be conducted as part of any other course.

# Leftover Contents

# IPv6 Address

```
#include <sys/socket.h>
```

```
#include <netinet/in.h>
```

```
struct sockaddr_in6 {  
    sa_family_t    sin6_family;    /* AF_INET6 */  
    in_port_t      sin6_port;      /* port number */  
    uint32_t        sin6_flowinfo; /* IPv6 flow information */  
    struct in6_addr sin6_addr;      /* IPv6 address */  
    uint32_t        sin6_scope_id; /* Scope ID (new in 2.4) */  
};
```

```
struct in6_addr {  
    unsigned char   s6_addr[16];    /* IPv6 address */  
};
```

# Compulsory Assignment #2: Making echo client/server “protocol independent”

Q-3. Revise echo client and server to be protocol independent (support both IPv4 and IPv6).

- Hint 1: `sockaddr` is too small for `sockaddr_in6`. `sockaddr_storage` has enough size to support both `sockaddr_in` and `sockaddr_in6`. (You will see this in server side program.)
- Hint 2: integrate `getaddrinfo` to avoid typing IPv6 address on your CLI
- Hint 3: you may use `hostname` (IPv4: “localhost”, IPv6: “ip6-localhost” address to develop / demonstrate the software on Ubuntu. They’re written in “/etc/hosts”.