

# SDN Hands-on

Kotaro Kataoka

# Outline: Let's Play with RYU & Mininet

- Completing installing of Mininet + Ryu
- Frequently used commands about OVS
- Running Learning Switch
- SDN App Development in a nutshell
  - Let's go deeper in the learning switch code
  - Let's implement Port Mirroring

# Installing Mininet and Upgrading Open vSwitch

- Mininet Installation
  - <https://github.com/mininet/mininet/wiki/Mininet-VM-Images>
- Launch Mininet VM (version 2.x.x) and check IP Address
  - mininet@mininet-vm:~\$ ifconfig
- Logging from VM host using SSH (-X for X11)
  - \$ ssh -X mininet@Mininet\_VM\_IP\_ADDRESS
- Upgrade Open vSwitch
  - <http://techandtrains.com/2013/10/24/another-way-to-upgrade-open-vswitch-in-mininet/>

# Installing SDN Controller to Mininet VM

- Installing Floodlight

`sudo apt-get install build-essential default-jdk ant python-dev eclipse`

Download floodlight @ <http://www.projectfloodlight.org/download/>

Please install Version 1.1 for developers (No VM Appliance) or Floodlight master(latest version) <https://github.com/floodlight/floodlight>

- Installing RYU using pip

`% pip install ryu`

- Installing from the source code

`% git clone git://github.com/osrg/ryu.git`

`% cd ryu; python ./setup.py install`

# Launching ryu-manager

```
$ ryu-manager --version
```

- Got Errors? Let's check required package is installed in Mininet-VM and re-install Ryu

- pkg\_resources.VersionConflict: (six 1.5.2 (/usr/lib/python2.7/dist-packages), Requirement.parse('six>=1.7.0'))

```
$ sudo pip install -U six
```

- pkg\_resources.DistributionNotFound: pbr>=0.6,!<0.7,<1.0

```
$ easy_install --upgrade pip
```

```
$ sudo python ./setup.py install
```

```
$ ryu-manager --version
```

# Launching Mininet with 3 hosts, 1 switch, 1 controller

- `sudo mn --topo single,3 --mac --switch ovsk --controller remote -x`
  - Topology of one switch and three hosts
  - Automatically sets the MAC address of the host
  - Open vSwitch as SDN switch
  - Uses external OpenFlow controller
  - xterm for generated nodes (hosts/switch/controller)
- Did 5 Windows come up?

# Showing OVS-Configuration

```
root@mininet-vm:~# ovs-vsctl show
1077578e-f495-46a1-a96b-441223e7cc22
Bridge "s1"
  Controller "tcp:127.0.0.1:6633"
  Controller "ptcp:6634"
  fail_mode: secure
  Port "s1-eth2"
    Interface "s1-eth2"
  Port "s1-eth1"
    Interface "s1-eth1"
  Port "s1"
    Interface "s1"
      type: internal
  Port "s1-eth3"
    Interface "s1-eth3"
  ovs_version: "2.0.2"
```

- Open vSwitch is abstracted as “Bridge”. “s1” is the name of switch

# Showing port and name on s1

```
root@mininet-vm:~# ovs-dpctl show
system@ovs-system:
  lookups: hit:127 missed:57 lost:0
  flows: 0
  port 0: ovs-system (internal)
  port 1: s1-eth3
  port 2: s1-eth1
  port 3: s1-eth2
  port 4: s1 (internal)
```



# Using OVS with OpenFlow 1.x

- Configuring OpenFlow 1.3 on s1

- Default version is 1.0

```
$ ovs-vsctl set bridge s1 protocols=OpenFlow1x
```

- And then check the empty flow table

```
$ ovs-ofctl -O OpenFlow13 dump-flows s1
```

- Datapath ID (DPID) of s1?

```
$ ovs-vsctl get bridge s1 datapath-id
```

# What happen if hosts attempt to communicate now?

- Ping fails because controller does not run any program

```
mininet> pingall
```

```
*** Ping: testing ping reachability
```

```
h1 -> X X
```

```
h2 -> X X
```

```
h3 -> X X
```

```
*** Results: 100% dropped (0/6 received)
```

# Lanching Learning Switch (OF1.x)

```
root@mininet-vm:~# ryu-manager --verbose  
ryu.app.simple_switch_1x
```

- Then let's do something again

```
mininet> pingall
```

# What did happen to Controller?

EVENT ofp\_event->SimpleSwitch13 EventOFPPacketIn

**packet in** 1 00:00:00:00:00:01 ff:ff:ff:ff:ff:ff 1

EVENT ofp\_event->SimpleSwitch13 EventOFPPacketIn

**packet in** 1 00:00:00:00:00:02 00:00:00:00:00:01 2

EVENT ofp\_event->SimpleSwitch13 EventOFPPacketIn

**packet in** 1 00:00:00:00:00:01 00:00:00:00:00:02 1

EVENT ofp\_event->SimpleSwitch13 EventOFPPacketIn

**packet in** 1 00:00:00:00:00:01 ff:ff:ff:ff:ff:ff 1

EVENT ofp\_event->SimpleSwitch13 EventOFPPacketIn

**packet in** 1 00:00:00:00:00:03 00:00:00:00:00:01 3

EVENT ofp\_event->SimpleSwitch13 EventOFPPacketIn

**packet in** 1 00:00:00:00:00:01 00:00:00:00:00:03 1

EVENT ofp\_event->SimpleSwitch13 EventOFPPacketIn

**packet in** 1 00:00:00:00:00:02 ff:ff:ff:ff:ff:ff 2

EVENT ofp\_event->SimpleSwitch13 EventOFPPacketIn

**packet in** 1 00:00:00:00:00:03 00:00:00:00:00:02 3

EVENT ofp\_event->SimpleSwitch13 EventOFPPacketIn

**packet in** 1 00:00:00:00:00:02 00:00:00:00:00:03 2

# Let's take a look of Flow Table

```
root@mininet-vm:~# ovs-ofctl -O OpenFlow13 dump-flows s1
```

```
OFPST_FLOW reply (OF1.3) (xid=0x2):
```

```
  cookie=0x0, duration=331.815s, table=0, n_packets=3, n_bytes=238,
```

```
priority=1,in_port=3,dl_dst=00:00:00:00:00:02 actions=output:2
```

```
  cookie=0x0, duration=331.834s, table=0, n_packets=373014, n_bytes=24618948,
```

```
priority=1,in_port=2,dl_dst=00:00:00:00:00:01 actions=output:1
```

```
  cookie=0x0, duration=331.825s, table=0, n_packets=3, n_bytes=238,
```

```
priority=1,in_port=3,dl_dst=00:00:00:00:00:01 actions=output:1
```

```
  cookie=0x0, duration=331.814s, table=0, n_packets=2, n_bytes=140,
```

```
priority=1,in_port=2,dl_dst=00:00:00:00:00:03 actions=output:3
```

```
  cookie=0x0, duration=331.824s, table=0, n_packets=2, n_bytes=140,
```

```
priority=1,in_port=1,dl_dst=00:00:00:00:00:03 actions=output:3
```

```
  cookie=0x0, duration=331.831s, table=0, n_packets=373036, n_bytes=16322899304,
```

```
priority=1,in_port=1,dl_dst=00:00:00:00:00:02 actions=output:2
```

```
  cookie=0x0, duration=395.824s, table=0, n_packets=9, n_bytes=546, priority=0
```

```
actions=CONTROLLER:65535
```

# Developing SDN Application using Ryu in a Nutshell

Thanks to

<http://sdnhub.org/tutorials/ryu/>

[https://github.com/osrg/ryu-book/blob/master/ja/source/switching\\_hub.rst](https://github.com/osrg/ryu-book/blob/master/ja/source/switching_hub.rst)

# Key Components of Ryu

- app/ – Contains set of applications that run on-top of the controller.
- base/ – Contains the base class for RYU applications. The RyuApp class in the app\_manager.py file is inherited when creating a new application.
- controller/ – Contains the required set of files to handle OpenFlow functions (e.g., packets from switches, generating flows, handling network events, gathering statistics etc).
- lib/ – Contains set of packet libraries to parse different protocol headers and a library for OFConfig. In addition, it includes parsers for Netflow and sFlow too.
- ofproto/ – Contains the OpenFlow protocol specific information and related parsers to support different versions of OF protocol (1.0, 1.2, 1.3, 1.4)
- topology/: Contains code that performs topology discovery related to OpenFlow switches and handles associated information (e.g., ports, links etc). Internally uses LLDP protocol.

# What to implement?

- Ability to **listen to asynchronous events** (e.g., **PACKET\_IN, FLOW\_REMOVED**) and to **observe events** using **ryu.controller.handler.set\_ev\_cls** decorator
- Ability to **parse incoming packets** (e.g., ARP, ICMP, TCP) and **fabricate packets to send out into the network**
- Ability to **create and send an OpenFlow/SDN message** (e.g., **PACKET\_OUT, FLOW\_MOD, STATS\_REQUEST**) to the programmable dataplane.




Let's take a look of  
simple\_switch13.py

# Importing Libraries

```
from ryu.base import app_manager
from ryu.controller import ofp_event
from ryu.controller.handler import CONFIG_DISPATCHER, MAIN_DISPATCHER
from ryu.controller.handler import set_ev_cls
from ryu.ofproto import ofproto_v1_3
from ryu.lib.packet import packet
from ryu.lib.packet import ethernet
```

Event Handlers

Packet Headers



Definition	Explanation
<code>ryu.controller.handler.HANDSHAKE_DISPATCHER</code>	Exchange of HELLO message
<code>ryu.controller.handler.CONFIG_DISPATCHER</code>	Waiting to receive SwitchFeatures message
<code>ryu.controller.handler.MAIN_DISPATCHER</code>	Normal status
<code>ryu.controller.handler.DEAD_DISPATCHER</code>	Disconnection of connection

# Initiating a new class

```
class SimpleSwitch13(app_manager.RyuApp):  
    OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]  
  
    def __init__(self, *args, **kwargs):  
        super(SimpleSwitch13, self).__init__(*args, **kwargs)  
        self.mac_to_port = {}
```

- Specify OpenFlow version to be 1.3
- Create Learning Table
  - Dictionary for “mac\_to\_port”

# Install table-miss flow entry during controller configures OVS

```
@set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
```

```
def switch_features_handler(self, ev):
```

```
    datapath = ev.msg.datapath
```

```
    ofproto = datapath.ofproto
```

```
    parser = datapath.ofproto_parser
```

During the configuration event of OVS...

```
    # install table-miss flow entry
```

```
    #
```

```
    # We specify NO BUFFER to max_len of the output action due to
```

```
    # OVS bug. At this moment, if we specify a lesser number, e.g.,
```

```
    # 128, OVS will send Packet-In with invalid buffer_id and
```

```
    # truncated packet data. In that case, we cannot output packets
```

```
    # correctly. The bug has been fixed in OVS v2.1.0.
```

```
    match = parser.OFPMatch()
```

```
    actions = [parser.OFPACTIONOutput(ofproto.OFPP_CONTROLLER,  
                                     ofproto.OFPCML_NO_BUFFER)]
```

```
    self.add_flow(datapath, 0, match, actions)
```

**Create a vacant match to let the switch send packet\_in to the controller if no flow entry matches in the flow table. Injected as a flow rule.**

# add\_flow ()

```
def add_flow(self, datapath, priority, match, actions, buffer_id=None):
    ofproto = datapath.ofproto
    parser = datapath.ofproto_parser

    inst = [parser.OFPInstructionActions(ofproto.OFPIT_APPLY_ACTIONS,
                                         actions)]

    if buffer_id:
        mod = parser.OFPFlowMod(datapath=datapath, buffer_id=buffer_id,
                                priority=priority, match=match,
                                instructions=inst)
    else:
        mod = parser.OFPFlowMod(datapath=datapath, priority=priority,
                                match=match, instructions=inst)

    datapath.send_msg(mod)
```

**“flow\_mod” message is sent from Controller to Switch**

# Unknown Packet Comes!

```
@set_ev_cls(ofp_event.EventOFPPacketIn, MAIN_DISPATCHER)
def _packet_in_handler(self, ev):
```

```
    # If you hit this you might want to increase
    # the "miss_send_length" of your switch
```

Every time packet\_in message comes,  
then this handler is called.

```
    if ev.msg.msg_len < ev.msg.total_len:
        self.logger.debug("packet truncated: only %s of %s bytes",
                           ev.msg.msg_len, ev.msg.total_len)
```

```
    msg = ev.msg
```

```
    datapath = msg.datapath
```

```
    ofproto = datapath.ofproto
```

From where did the packet come?

```
    parser = datapath.ofproto_parser
```

```
    in_port = msg.match['in_port']
```

```
    pkt = packet.Packet(msg.data)
```

```
    eth = pkt.get_protocols(ethernet.ethernet)[0]
```

```
    dst = eth.dst
```

We want the content of Ethernet Header

```
    src = eth.src
```

```
    dpid = datapath.id
```

```
    self.mac_to_port.setdefault(dpid, {})
```

Logging on Controller Contole

```
    self.logger.info("packet in %s %s %s %s", dpid, src, dst, in_port)
```

# Core Part as Learning Switch

```
# learn a mac address to avoid FLOOD next time.  
self.mac_to_port[dpid][src] = in_port
```

```
if dst in self.mac_to_port[dpid]:  
    out_port = self.mac_to_port[dpid][dst]  
else:  
    out_port = ofproto.OFPP_FLOOD
```


```
actions = [parser.OFPACTIONOutput(out_port)]
```

```
# install a flow to avoid packet_in next time  
if out_port != ofproto.OFPP_FLOOD:  
    match = parser.OFPMATCH(in_port=in_port, eth_dst=dst)  
    # verify if we have a valid buffer_id, if yes avoid to send both  
    # flow_mod & packet_out  
    if msg.buffer_id != ofproto.OFP_NO_BUFFER:  
        self.add_flow(datapath, 1, match, actions, msg.buffer_id)  
        return  
    else:  
        self.add_flow(datapath, 1, match, actions)
```

```
data = None  
if msg.buffer_id == ofproto.OFP_NO_BUFFER:  
    data = msg.data
```

```
out = parser.OFPPACKETOut(datapath=datapath, buffer_id=msg.buffer_id,  
                           in_port=in_port, actions=actions, data=data)  
datapath.send_msg(out)
```

On S1, H1 is reachable via Port X



Deciding outgoing port of the packet  
One port to DST or Flooding

Add a flow rule to the destination host  
at the OVS if it has been known.

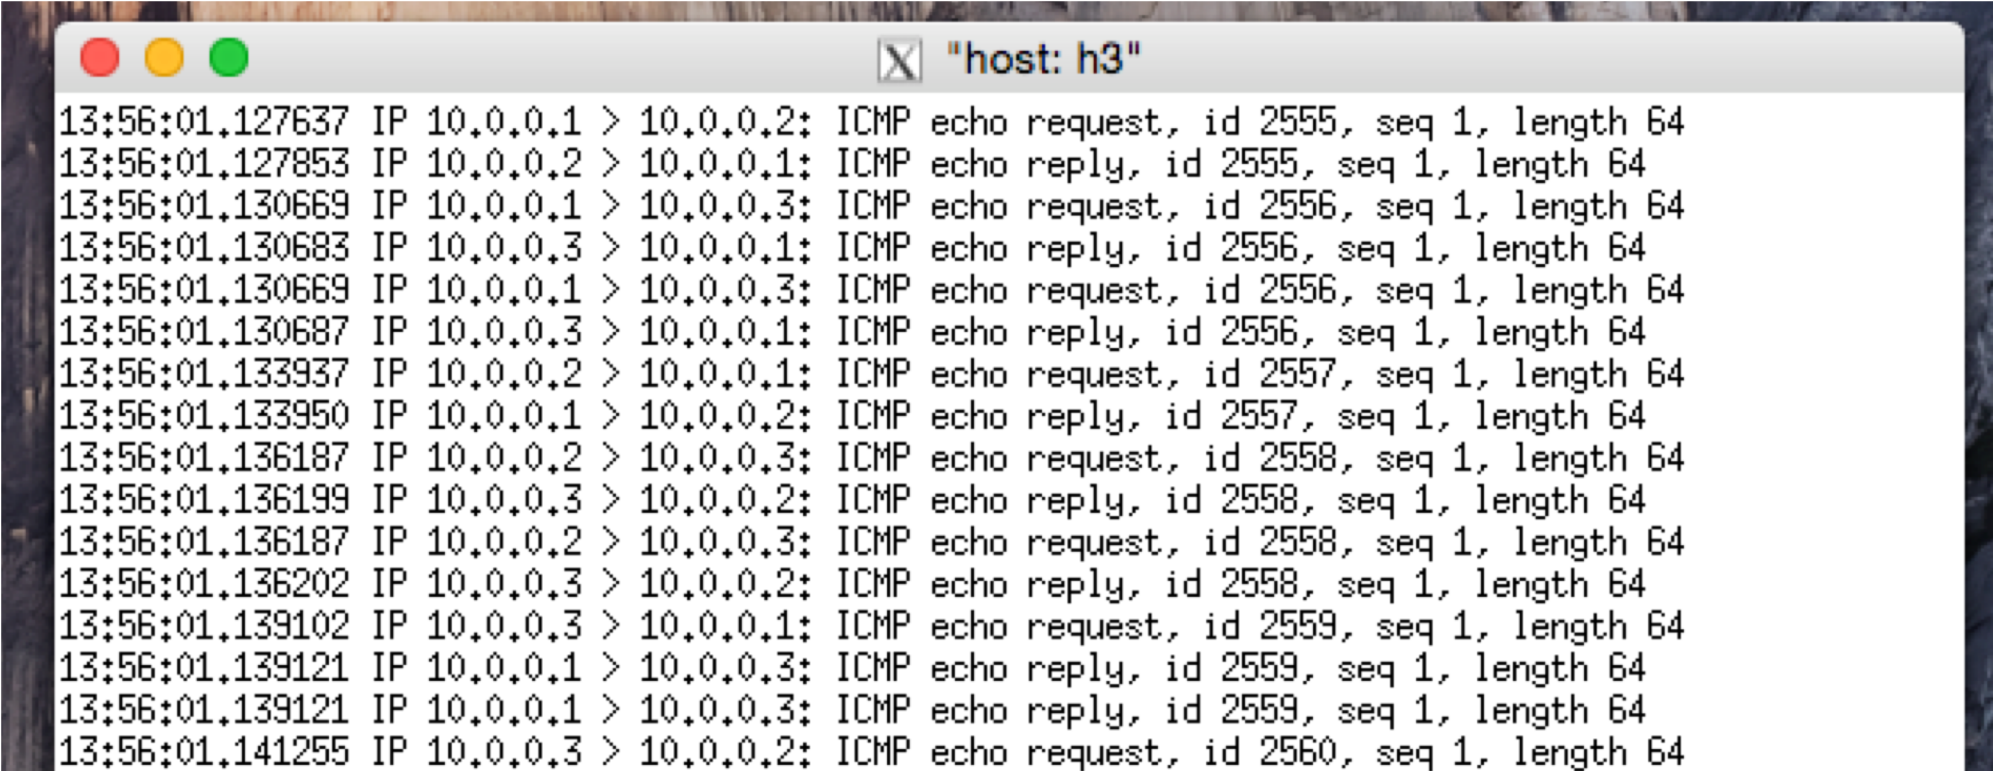
Packet Out!

# A Small Extension



# Enabling port mirroring on the learning switch

- One port on OVS copies all incoming packets to Mirroring Port (ex. Port 3 on S1)
- You may hard-code the mirroring port

A terminal window titled "host: h3" displays a series of network traffic logs. Each line shows a timestamp, MAC address, IP addresses, and a description of the packet type and its details. The logs show a sequence of ICMP echo requests and replies between three IP addresses: 10.0.0.1, 10.0.0.2, and 10.0.0.3. The timestamps are in the format HH:MM:SS.microseconds.

```
13:56:01.127637 IP 10.0.0.1 > 10.0.0.2: ICMP echo request, id 2555, seq 1, length 64
13:56:01.127853 IP 10.0.0.2 > 10.0.0.1: ICMP echo reply, id 2555, seq 1, length 64
13:56:01.130669 IP 10.0.0.1 > 10.0.0.3: ICMP echo request, id 2556, seq 1, length 64
13:56:01.130683 IP 10.0.0.3 > 10.0.0.1: ICMP echo reply, id 2556, seq 1, length 64
13:56:01.130669 IP 10.0.0.1 > 10.0.0.3: ICMP echo request, id 2556, seq 1, length 64
13:56:01.130687 IP 10.0.0.3 > 10.0.0.1: ICMP echo reply, id 2556, seq 1, length 64
13:56:01.133937 IP 10.0.0.2 > 10.0.0.1: ICMP echo request, id 2557, seq 1, length 64
13:56:01.133950 IP 10.0.0.1 > 10.0.0.2: ICMP echo reply, id 2557, seq 1, length 64
13:56:01.136187 IP 10.0.0.2 > 10.0.0.3: ICMP echo request, id 2558, seq 1, length 64
13:56:01.136199 IP 10.0.0.3 > 10.0.0.2: ICMP echo reply, id 2558, seq 1, length 64
13:56:01.136187 IP 10.0.0.2 > 10.0.0.3: ICMP echo request, id 2558, seq 1, length 64
13:56:01.136202 IP 10.0.0.3 > 10.0.0.2: ICMP echo reply, id 2558, seq 1, length 64
13:56:01.139102 IP 10.0.0.3 > 10.0.0.1: ICMP echo request, id 2559, seq 1, length 64
13:56:01.139121 IP 10.0.0.1 > 10.0.0.3: ICMP echo reply, id 2559, seq 1, length 64
13:56:01.139121 IP 10.0.0.1 > 10.0.0.3: ICMP echo reply, id 2559, seq 1, length 64
13:56:01.141255 IP 10.0.0.3 > 10.0.0.2: ICMP echo request, id 2560, seq 1, length 64
```

# Question

- How to modify the sample code?
- Think by your self first!

# Answer

```
if dst in self.mac_to_port[dpid]:
    out_port = self.mac_to_port[dpid][dst]
else:
    out_port = ofproto.OFPP_FLOOD

actions = [parser.OFPActionOutput(out_port)]
```

Add 3 lines here!!

```
if in_port != 3:
    if out_port != ofproto.OFPP_FLOOD:
        actions = [parser.OFPActionOutput(out_port), parser.OFPActionOutput(3)]
```

```
# install a flow to avoid packet_in next time
if out_port != ofproto.OFPP_FLOOD:
    match = parser.OFPMatch(in_port=in_port, eth_dst=dst)
    # verify if we have a valid buffer_id, if yes avoid to send both
    # flow_mod & packet_out
    if msg.buffer_id != ofproto.OFP_NO_BUFFER:
        self.add_flow(datapath, 1, match, actions, msg.buffer_id)
        return
    else:
        self.add_flow(datapath, 1, match, actions)
```

# Running Your Code

- Stop the running controller (Ctrl-C on Controller)

- Flush the flow table of S1

```
# ovs-ofctl -O OpenFlow13 del-flows s1
```

- Run your code

```
# ryu-manager simple_switch_13-Mirror.py
```

# Oops... Doesn't look very good?

- 2 out\_ports for in\_port 1 and 2 is correct, but...

```
root@mininet-vm:~# ovs-ofctl -O OpenFlow13 del-flows s1
root@mininet-vm:~# ovs-ofctl -O OpenFlow13 dump-flows s1
OFPST_FLOW reply (OF1.3) (xid=0x2):
  cookie=0x0, duration=2.736s, table=0, n_packets=0, n_bytes=0, priority=0 ac
tions=CONTROLLER:65535
root@mininet-vm:~# ovs-ofctl -O OpenFlow13 dump-flows s1
OFPST_FLOW reply (OF1.3) (xid=0x2):
  cookie=0x0, duration=8.335s, table=0, n_packets=4, n_bytes=224, priority=1,
in_port=3,d1_dst=00:00:00:00:00:02 actions=output:2
  cookie=0x0, duration=8.355s, table=0, n_packets=3, n_bytes=182, priority=1,
in_port=2,d1_dst=00:00:00:00:00:01 actions=output:1,output:3
  cookie=0x0, duration=8.349s, table=0, n_packets=4, n_bytes=224, priority=1,
in_port=3,d1_dst=00:00:00:00:00:01 actions=output:1
  cookie=0x0, duration=8.331s, table=0, n_packets=3, n_bytes=182, priority=1,
in_port=1,d1_dst=00:00:00:00:00:03 actions=output:3,output:3
  cookie=0x0, duration=8.337s, table=0, n_packets=3, n_bytes=182, priority=1,
in_port=2,d1_dst=00:00:00:00:00:03 actions=output:3,output:3
  cookie=0x0, duration=8.342s, table=0, n_packets=2, n_bytes=84, priority=1,i
n_port=1,d1_dst=00:00:00:00:00:02 actions=output:2,output:3
  cookie=0x0, duration=20.198s, table=0, n_packets=9, n_bytes=882, priority=0
actions=CONTROLLER:65535
root@mininet-vm:~#
```

# Additional Hands-on?

- SDN-based Router
  - Develop network topology with multiple subnets
  - Giving IP address from different subnet to each host
  - Give static IP address to each interface on OVS -> Default gateway
  - Handling ARP messages from hosts
  - Handling IP packets from hosts
  - You may forget updating TTL, returning ICMP error messages

Done