

# Towards Function Approximation Methods

Easwar Subramanian

TCS Innovation Labs, Hyderabad

Email : [easwar.subramanian@tcs.com](mailto:easwar.subramanian@tcs.com) / [cs5500.2020@iith.ac.in](mailto:cs5500.2020@iith.ac.in)

October 07, 2021

## 1 Review

## 2 Function Approximation Methods

# Review

- Given a MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  and a policy  $\pi$ , the **prediction problem** involves computing

★ State-value function :

$$V^{\pi}(s) = \mathbb{E}_{\pi} \left( \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s \right)$$

★ Action value function :

$$Q^{\pi}(s, a) = \mathbb{E}_{\pi} \left( \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s, a_t = a \right)$$

- Given a MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ , the **control problem** involves finding optimal value functions  $V_*$  and  $Q_*$  or equivalently finding optimal policy  $\pi_*$ .

Model based implies, we know  $\mathcal{P}$  and  $\mathcal{R}$

- Iterative Policy Evaluation (for prediction problem)

$$V_{k+1}(s) \leftarrow \sum_a \pi(s, a) \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V_k(s')]$$

- Value Iteration (for control problem)

$$V_*(s) \leftarrow \max_a \left[ \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a (\mathcal{R}_{ss'}^a + \gamma V_*(s')) \right]$$

## Drawbacks of DP Algorithms

- Requires full prior knowledge of the dynamics of the environment
- Can be implemented only on small, discrete state spaces

- Key idea is to estimate the following expectations using samples

$$\begin{aligned} V^\pi(s) &= \mathbb{E}_\pi \left( \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s \right) \\ &= \mathbb{E}_\pi [r_{t+1} + \gamma V^\pi(s_{t+1}) | s_t = s] \end{aligned}$$

- Monte-Carlo and TD Methods
- **One Step TD for  $V$**  : If the *transition*  $(s, r, s')$  is observed at time  $t$  under policy  $\pi$ , then

$$V(s) \leftarrow V(s) + \alpha_t [r + \gamma V(s') - V(s)]$$

- **One step TD for  $Q$**  : Given the transition  $(s, a, r, s')$  from  $\pi$ , sample  $a' \sim \pi(s')$  and update

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$$

- Policy is always  $\epsilon$ -greedy with  $\epsilon$  decay
- **SARSA Update** : Given a trajectory segment  $(s, a, r, s', a')$  generated by the  $\epsilon$ -greedy policy, the policy evaluation steps involves the following update

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$$

and policy improvement is done using  $\epsilon$ -greedy with respect to current  $Q$

- **Q-learning update (Watkins)** : Given a trajectory segment  $(s, a, r, s')$  generated by the  $\epsilon$ -greedy policy, update

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

## Drawback


- Not scalable to large state and action spaces

# Function Approximation Methods



# On the need for Function Approximators


- To solve large scale RL problems
  - ★ Game of Backgammon :  $10^{20}$  states
  - ★ Game of Go :  $10^{170}$  states
  - ★ Even Atari games have large state space



$|S|$  is very large : Curse of Dimensionality

- ▶ Value function have been basically lookup tables.
- ▶ Solution for large MDP's is to use function approximators
  - ★ Generalize from seen to unseen states
- ▶ Function approximators could be
  - ★ Linear function approximator
  - ★ **Neural networks**
  - ★ Decision tree
  - ★ ...

# Neural Network Approximators




# Policy Evaluation Using Neural Networks

The value of a policy  $\pi$  is given by

$$\begin{aligned} V^{\pi}(s) &= \mathbb{E}_{\pi} \left( \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s \right) \\ &= \mathbb{E}_{\pi} [r_{t+1} + \gamma V^{\pi}(s_{t+1}) | s_t = s] \end{aligned}$$

**Question :** How do we compute the above expectations using neural networks ?



- Roll-out  $m$  trajectories from state  $s$  and observe rewards

# Value Function Fitting using Monte Carlo

- Consider a MDP with a finite horizon  $H$

$$V^\pi(s) \approx \frac{1}{m} \left[ \sum_{j=1}^m \left[ \sum_{k=0}^H \left( \gamma^k r_{t+k+1}^i | s_t = s \right) \right] \right]$$

- Need to reset the simulator back to state  $s$  (Not always possible)
- Alternative : Roll-out single sample estimate (high variance, but OK)
- Collect training data for as many states as possible and regress thereafter

$$\left( s_i, \underbrace{\left[ \sum_{k=0}^H \left( \gamma^k r_{t+k+1}^i | s_t = s \right) \right]}_{=y_i} \right)$$

---

## Algorithm Monte Carlo Based Value Function Fitting

---

Initialize number of iterations  $N$

**for**  $i = 1$  to  $N$  **do**

    Perform a roll-out from an initial state  $s_i$  (could be any state from  $\mathcal{S}$ )

    Calculate targets  $y_i$  using Monte-Carlo roll outs

$$y_i = \left[ \sum_{k=0}^H \left( \gamma^k r_{t+k+1}^i | s_t = s_i \right) \right]$$

    Form input-output pairs  $(s_i, y_i)$  ( $N$  datapoints in total)

**end for**

Perform supervised regression with loss function

$$L(\phi) = \frac{1}{2} \sum_{i=1}^N [V_{\phi}^{\pi}(s_i) - y_i]^2$$

- Needs complete sequences, suitable only for episodic tasks

# Fitted V Iteration

We observe transition  $(s, a, r, s')$  at time  $t$ ; Using one step look-ahead,

$$\begin{aligned} V^\pi(s) &= \mathbb{E}_\pi [r + \gamma V^\pi(s') | s_t = s] \\ &\approx r + \gamma V^\pi(s') \text{ (Bootstrap } V^\pi) \end{aligned}$$

Using function approximators, we get,

$$V_\phi^\pi(s) \approx r + \gamma V_\phi^\pi(s')$$

- Directly use the previous fitted value function  $V_\phi^\pi$
- Collect training data,

$$\left( s_i, \underbrace{r + V_\phi^\pi(s'_i)}_{=y_i} \right)$$

- Perform supervised regression

$$L(\phi) = \frac{1}{2} \sum_{i=1}^N [V_\phi^\pi(s_i) - y_i]^2$$



---

## Algorithm Fitted V Iteration

---

- 1: Initialize number of iterations  $N$
- 2: **for**  $j = 1$  to  $N$  **do**
- 3:   Sample  $K$  transitions  $(s, a, r, s')$  using policy  $\pi$
- 4:   **for**  $i = 1$  to  $K$  **do**
- 5:     Calculate targets  $y_i$  using one step TD approximation

$$y_i = \left[ r + V_{\phi_j}^{\pi}(s'_i) \right]$$

- 6:     Form input-output pairs  $(s_i, y_i)$  ( $K$  datapoints in total)
- 7:   **end for**
- 8:   Perform supervised regression (Optimizer : RProp) using loss function

$$L(\phi_j) = \frac{1}{2} \sum_{i=1}^K \left[ V_{\phi_j}^{\pi}(s_i) - y_i \right]^2$$

and get a new function approximator with new weights  $\phi_{j+1}$

- 9: **end for**

# Optimal Value Function : Control

Bellman optimality equation for  $V_*$  is given by,

$$V_*(s) \leftarrow \max_a \left[ \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a (\mathcal{R}_{ss'}^a + \gamma V_*(s')) \right] \approx \max_a E [r_{t+1} + \gamma V_*(s_{t+1}) | s_t = s]$$

**Question :** How do we get a sample estimate for transition  $(s, a, r, s')$  for  $V_*$  ?

$$V(s) \approx \max_a [r + \gamma V(s')]$$



- To compute max over  $a$ , we need to know the outcome of all actions starting from  $s$ . Mostly not possible and costly as well.
- For model free control, we use approximators for  $Q$  and not  $V$

Bellman optimality equation for  $Q_*$

$$Q_*(s, a) = \left[ \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \left( \mathcal{R}_{ss'}^a + \gamma \max_{a'} Q_*(s', a') \right) \right] \approx \mathbb{E} \left[ r_{t+1} + \gamma \max_{a'} Q_*(s_{t+1}, a') \mid s_t = s, a_t = a \right]$$

- ▶ Max is inside the expectation; that's ok
- ▶ For transitions  $(s, a, r, s')$  we can compute  $r + \gamma \max_{a'} Q(s', a')$
- ▶ Does not require simulating over actions
- ▶ Use the previous fitted optimal Q function  $Q_\phi^*$  like in fitted V iteration
- ▶ Collect training data,

$$\left( s_i, \underbrace{r + \gamma \max_{a'} Q_\phi(s'_i, a'_i)}_{=y_i} \right)$$

- ▶ Perform supervised regression

$$L(\phi) = \frac{1}{2} \sum_{i=1}^N [Q_\phi(s_i, a_i) - y_i]^2$$

---

## Algorithm Fitted Q Iteration

---

- 1: Initialize number of iterations  $N$
- 2: **for**  $j = 1$  to  $N$  **do**
- 3:   Sample  $K$  transitions  $(s, a, r, s')$  using any behaviour policy  $\mu$
- 4:   **for**  $i = 1$  to  $K$  **do**
- 5:     Calculate targets  $y_i$  using one step TD approximation

$$y_i = \left[ r + \gamma \max_{a'} Q_{\phi_j}(s'_i, a') \right]$$

- 6:     Form input-output pairs  $(s_i, y_i)$  ( $K$  Datapoints in total)
- 7:   **end for**
- 8:   Perform supervised regression (Optimizer : RProp) using loss function

$$L(\phi_j) = \frac{1}{2} \sum_{i=1}^K \left[ Q_{\phi_j}(s_i, a_i) - y_i \right]^2$$

and get a new function approximator with new weights  $\phi_{j+1}$

- 9: **end for**