

Advanced Policy Gradients - II

Easwar Subramanian

TCS Innovation Labs, Hyderabad

Email : easwar.subramanian@tcs.com / cs5500.2020@iith.ac.in

Novemer 20, 2021

Overview of this Lecture

- 1 Review
- 2 Approximations to Trust Region Formulation
- 3 Natural Policy Gradient
- 4 Fisher Information Matrix and KL Divergence
- 5 Relationship of Natural Gradient to Policy Gradient
- 6 Other Algorithms
- 7 Proximal Policy Optimization

Review

Policy Optimization Problem

The performance of a policy π_θ is given by

$$J(\theta) = V^{\pi_\theta}(s_0) = \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t r_{t+1} | s_0 = s \right]$$

where $\gamma < 1$ is the discount factor of the MDP

General form for gradient of the performance measure is given by

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^{\infty} \nabla_\theta \log \pi_\theta(a_t | s_t) \Psi_t \right]$$

Disadvantages are :

- ▶ Sample Inefficiency : **on-policy expectation**
- ▶ Distance in parameter space \neq policy space

We recast the optimization problem using a **surrogate loss function**

$$\arg \max_{\pi'} J(\pi') = \arg \max_{\pi'} [J(\pi') - J(\pi_0)] \approx \mathcal{L}_{\pi_0}(\pi')$$

where

$$\mathcal{L}_{\pi_0}(\pi') = \mathbb{E}_{\tau \sim \pi_0} \left[\sum_{t=0}^{\infty} \gamma^t \frac{\pi'(a_t | s_t)}{\pi_0(a_t | s_t)} A^{\pi_0}(s_t, a_t) \right]$$

The approximation is valid if policies π' and π_0 are '**close**' in terms of their KL divergence

Relative Policy Performance Bound

We can have a **relative policy performance bound** using KL divergence to measure the goodness of the approximation obtained

$$\left[J(\pi') - (J(\pi_0) + \mathcal{L}_{\pi_0}(\pi')) \right] \leq C \sqrt{\mathbb{E}_{s \sim d^{\pi_0}} [D_{KL}(\pi' || \pi_0)[s]]}$$

This gives rise to an optimization routine with the following iterative procedure with π_{k+1} and π_k are related by

$$\pi_{k+1} = \arg \max_{\pi'} \left[\mathcal{L}_{\pi_k}(\pi') - C \sqrt{\mathbb{E}_{s \sim d^{\pi_k}} [D_{KL}(\pi' || \pi_k)[s]]} \right]$$

Performance guarantee

$$[J(\pi_{k+1}) - J(\pi_k)] \geq 0$$

- C is quite high when γ is close to 1 and hence choosing step size becomes an issue

-
- 1: Initialize π_0
 - 2: **for** $k = 0, 1, 2, \dots$ until convergence **do**
 - 3: Sample a trajectory τ from policy π_k
 - 4: Compute advantage function $A^{\pi_{\theta_k}}(a_t, s_t)$ for all (s_t, a_t) pairs in the trajectory τ
 - 5: Solve the optimization problem

$$\pi_{k+1} = \arg \max_{\pi'} L_{\pi_k}(\pi') - C \sqrt{\mathbb{E}_{s \sim d^{\pi_k}} \left[D_{KL}(\pi' || \pi_k)[s] \right]}$$


6: **end for**

Issues are :

- ▶ C is quite high when γ is close to 1 $\left(C = \frac{4\varepsilon\gamma}{1-\gamma^2} \alpha^2 \right)$
- ▶ Consequently, step size becomes too small

Majorize Maximize Framework

Majorize-Maximize framework is used to solve the optimization step



- ▶ Instead of KL penalty, use KL constraint
- ▶ Can control worst case error through constraint upper limit

$$\pi_{k+1} = \arg \max_{\pi'} [L_{\pi_k}(\pi')]$$
$$\text{such that } \mathbb{E}_{s \sim d^{\pi_k}} D_{KL}(\pi' || \pi_k)[s] \leq \delta$$

- ▶ From the constraint, **steps respect a notion of distance in policy space**
- ▶ Above constrained optimization is basis of many algorithms, Natural Policy Gradient (NPG), truncated NPG, TRPO and PPO
- ▶ The objective and the constraint can be estimated from the roll-out of old policies – **sample efficient**
- ▶ Update is **invariant** to parametrization

Approximations to Trust Region Formulation

Trust Region Formulation

We have the following optimization problem

$$\begin{aligned}\pi_{k+1} &= \arg \max_{\pi'} [\mathcal{L}_{\pi_k}(\pi')] \\ \text{such that } &\bar{D}_{KL}(\pi' || \pi_k) \leq \delta\end{aligned}$$

The constraint on the optimization problem is the trust region with size δ and some guarantees on performance improvement are there within the trust region

For parametrized policies the optimization can be written as

$$\begin{aligned}\pi_{\theta_{k+1}} &= \arg \max_{\pi_{\theta}} [\mathcal{L}_{\pi_{\theta_k}}(\pi_{\theta})] \\ \text{such that } &\bar{D}_{KL}(\pi_{\theta} || \pi_{\theta_k}) \leq \delta\end{aligned}$$

How do we solve it ?

- ▶ Linear approximation for the objective
- ▶ Quadratic approximation for the constraint

Approximation of Objective Function

Taylor series expansion for function $f(x)$ around point a is given by

$$f(x) = f(a) + f'(a)(x - a) + \frac{1}{2}f''(a)(x - a)^2 + \dots$$

- Using Taylor series expansion on objective function $\mathcal{L}_{\pi_{\theta_k}}(\pi_{\theta})$ around θ_k (upto first order term) gives us

$$\mathcal{L}_{\pi_{\theta_k}}(\pi_{\theta}) \approx \mathcal{L}_{\pi_{\theta_k}}(\pi_{\theta_k}) + g^T(\theta - \theta_k) \quad \text{where } g \doteq \nabla_{\theta} \mathcal{L}_{\pi_{\theta_k}}(\pi_{\theta})|_{\theta=\theta_k}$$

- Recall that g is exactly the policy gradient (from previous lecture !)

$$\nabla_{\theta} \mathcal{L}_{\pi_{\theta_k}}(\pi_{\theta})|_{\theta=\theta_k} = \mathbb{E}_{\tau \sim \pi_{\theta_k}} \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log(\pi_{\theta_k}(a_t|s_t))|_{\theta=\theta_k} \gamma^t A^{\pi_{\theta_k}}(s_t, a_t) \right]$$

- Objective function is simplified to

$$\theta_{k+1} = \arg \max_{\theta} g^T(\theta - \theta_k)$$

Using Taylor series expansion on the constraint (around θ_k ; upto second order) gives us

$$\bar{D}_{KL}(\pi_\theta || \pi_{\theta_k}) \approx \cancel{\bar{D}_{KL}(\pi_{\theta_k} || \pi_{\theta_k})}^0 + \cancel{\nabla_\theta \bar{D}_{KL}(\pi_\theta || \pi_{\theta_k})|_{\theta=\theta_k}}^0 + \nabla_\theta^2 \bar{D}_{KL}(\pi_\theta || \pi_{\theta_k}) |_{\theta=\theta_k}$$

The first order term $\nabla_\theta \bar{D}_{KL}(\pi_\theta || \pi_{\theta_k})$ evaluates to zero since the expectation of the score function is zero

$$\nabla_\theta \bar{D}_{KL}(\pi_\theta || \pi_{\theta_k}) = \nabla_\theta \mathbb{E}_{\pi_\theta} [\log \pi_\theta] - \nabla_\theta \mathbb{E}_{\pi_{\theta_k}} [\log \pi_{\theta_k}] = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta] = 0$$

Therefore, we are left only with the second order term

$$\bar{D}_{KL}(\pi_\theta || \pi_{\theta_k}) \approx \frac{1}{2}(\theta - \theta_k)^T H (\theta - \theta_k) \quad \text{where } H \doteq \nabla_\theta^2 \bar{D}_{KL}(\pi_\theta || \pi_{\theta_k}) |_{\theta=\theta_k}$$

Natural Policy Gradient

The optimization problem is now simplified as

$$\begin{aligned}\theta_{k+1} &= \arg \max_{\theta} g^T (\theta - \theta_k) \\ \text{such that } &\frac{1}{2} (\theta - \theta_k)^T H (\theta - \theta_k) \leq \delta\end{aligned}$$

Linear objective with quadratic constraint

Solution to the approximate problem obtained using Lagrange multiplier method

$$\theta_{k+1} = \theta_k + \sqrt{\frac{2\delta}{g^T H^{-1} g}} H^{-1} g$$

The term $H^{-1}g$ is called the Natural gradient

Algorithm Natural Policy Gradient

- 1: Initialize π_0
- 2: **for** $k = 0, 1, 2, \dots$ **do**
- 3: Collect trajectories D_k on policy $\pi_k = \pi_{\theta_k}$
- 4: Estimate all advantages $A^{\pi_{\theta_k}}(s_t, a_t)$
- 5: Form sample estimates for policy gradients \hat{g}_k (using advantage estimates)
- 6: Form sample estimates for the Hessian of KL divergence
- 7: Compute the Natural Policy Gradient update

$$\theta_{k+1} = \theta_k + \sqrt{\frac{2\delta}{g_k^T H_k^{-1} g_k}} H_k^{-1} g_k$$

8: **end for**

Fisher Information Matrix and KL Divergence

- ▶ Let $p(x|\theta)$ be a probability distribution parameterized by θ .
- ▶ Score function of a parameterized probability distribution is given by

$$s(\theta) = \nabla_{\theta} \log p(x|\theta) ,$$

- ▶ For a parameter vector θ , Fisher Information Matrix is given by,

$$F = \mathbb{E}_{p(x|\theta)} \left[\nabla_{\theta} \log p(x|\theta) \nabla_{\theta} \log p(x|\theta)^T \right] .$$

- ▶ The sample estimate of the above expectation is given by,

$$F = \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \log p(x_i|\theta) \nabla_{\theta} \log p(x_i|\theta)^T . \quad (1)$$

- ▶ **Claim** : Fisher Information Matrix F is the Hessian of KL-divergence between two probability distributions $p(x|\theta')$ and $p(x|\theta)$ evaluated at $\theta' = \theta$

$$\text{KL}[p(x|\theta') \parallel p(x|\theta)] = \mathbb{E}_{p(x|\theta)} [H_{\log p(x|\theta)}] = F$$

- ▶ Natural policy gradient algorithm gives an update-rule in which updates are pre-multiplied by H^{-1}
- ▶ The Hessian of the KL-divergence is the Fischer Information Matrix given by

$$F = \mathbb{E}_{\pi_{\theta}} \left[\nabla \log \pi_{\theta}(\cdot|s) \nabla \log \pi_{\theta}(\cdot|s)^T \right]$$

- ▶ The NPG direction $H^{-1}g$ is **co-variant**; i.e. it points in same direction irrespective of the parametrization that is used to compute it

Relationship of Natural Gradient to Policy Gradient

Consider the following optimization problem

$$\pi_{\theta_{k+1}} = \arg \max_{\pi_{\theta}} \left[\mathcal{L}_{\pi_{\theta_k}}(\pi_{\theta}) \right]$$

such that $\|\theta - \theta_k\|^2 \leq \delta$

After **linearising** the objective, the optimization problem is now,

$$\theta_{k+1} = \arg \max_{\theta} g^T(\theta - \theta_k) \text{ such that } (\theta - \theta_k)^2 \leq \delta$$

This is the original policy gradient problem !!

We move a small distance in parameter space in the direction of the gradient

Natural policy gradient problem is given by,

$$\pi_{\theta_{k+1}} = \arg \max_{\pi_{\theta}} \left[\mathcal{L}_{\pi_{\theta_k}}(\pi_{\theta}) \right]$$

such that $\bar{D}_{KL}(\pi_{\theta} || \pi_{\theta_k}) \leq \delta$

After **linearising** the objective and **quadratising** the constraint, the optimization problem is then given by,

$$\theta_{k+1} = \arg \max_{\theta} g^T(\theta - \theta_k) \text{ such that } \frac{1}{2}(\theta - \theta_k)^T F (\theta - \theta_k) \leq \delta$$

- ▶ Vanilla policy gradient has the right objective but "*incorrect*" constraint (Euclidean penalty instead of KL penalty)
- ▶ Recall that, policy iteration (from MDP lectures) obtain policy improvement with no constraint

Other Algorithms

- ▶ **Problem** : For neural networks, the dimensionality of parameter θ are high. High computational cost in inverting the matrix H
- ▶ **Solution** : Use the **conjugate gradient algorithm** to compute $H^{-1}g$ without inverting H
- ▶ Resultant algorithm : Truncated Natural Policy Gradient
- ▶ ACTKR algorithm uses KFAC technique to solve the inverse Hessian computation problem

- ▶ Another problem with NPG update is that - might not be robust to trust region size δ
 - ★ δ may be too large in some iterations and can degrade the performance
- ▶ Because of quadratic approximation, the KL-divergence constraint may be violated
- ▶ Monotonic improvement may not occur in all iterations

- ▶ Enforce improvement in surrogate (i.e. $\mathcal{L}_{\pi_{\theta_k}}(\pi_{\theta}) \geq 0$)
- ▶ Enforce KL constraint
- ▶ How ? Backtracking line search with exponential decay

Algorithm Line Search for TRPO

- 1: Compute the proposed policy step $\Delta_k = \sqrt{\frac{2\delta}{g_k^T H_k^{-1} g_k}} H_k^{-1} g_k$
 - 2: **for** $j = 0, 1, 2, \dots, N$ **do**
 - 3: Compute proposed update $\theta = \theta_k + \alpha_j \Delta_k$
 - 4: **If** $\mathcal{L}_{\pi_{\theta_k}}(\pi_{\theta}) \geq 0$ and $\bar{D}_{KL}(\theta || \theta_k) \leq \delta$
 - 5: Accept the update $\theta = \theta_k + \alpha_j \Delta_k$
 - 6: **Else**
 - 7: Find another α_j (Reduce α_j)
 - 8: **end for**
-

Algorithm Trust Region Policy Optimization

- 1: Initialize π_0
- 2: **for** $k = 0, 1, 2, \dots$ **do**
- 3: Collect trajectories D_k on policy $\pi_k = \pi_{\theta_k}$
- 4: Estimate all advantages $A^{\pi_{\theta_k}}(s_t, a_t)$
- 5: Form sample estimates for policy gradients \hat{g}_k (using advantage estimates)
- 6: Form sample estimates for the Hessian of KL divergence / FIM
- 7: Use conjugate gradient to obtain FIM estimate H^{-1}
- 8: Estimate step size α using backtracking line search to enforce KL constraint and monotonic improvement
- 9: Compute the Natural Policy Gradient update

$$\theta_{k+1} = \theta_k + \alpha \Delta_k$$

10: **end for**

Proximal Policy Optimization

Proximal Policy Optimization is a family of methods that approximately enforce without actually computing the natural gradient

► Adaptive KL Penalty

$$\pi_{\theta_{k+1}} = \arg \max_{\pi_{\theta}} \left[\mathcal{L}_{\pi_{\theta_k}}(\pi_{\theta}) - \beta \bar{D}_{KL}(\pi_{\theta} || \pi_{\theta_k}) \right]$$

Penalty co-efficient β is changed between iterations to approximately enforce KL constraint

► Clipped Objective (Simpler to implement, no need to check KL constraint; works well)

$$\mathcal{L}_{\pi_{\theta_k}}^{CLIP}(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta_k}} \left[\sum_{t=0}^T \min(r_t(\theta) A_t^{\pi_{\theta_k}}, \text{clip}(r_t(\theta), 1 - \varepsilon, 1 + \varepsilon) A_t^{\pi_{\theta_k}}) \right]$$

where $r_t(\theta)$ is the importance sampling ratio between target policy π_{θ} and behaviour policy π_{θ_k} and policy update takes place as

$$\pi_{\theta_{k+1}} = \arg \max_{\pi_{\theta}} \mathcal{L}_{\pi_{\theta_k}}^{CLIP}(\pi_{\theta})$$