

# Deep Q Networks - Variants

Easwar Subramanian

TCS Innovation Labs, Hyderabad


Email : [easwar.subramanian@tcs.com](mailto:easwar.subramanian@tcs.com) / [cs5500.2020@iith.ac.in](mailto:cs5500.2020@iith.ac.in)

October 14, 2021

- 1 Review
- 2 Efficacy of DQN Algorithm
- 3 Double DQN
- 4 Prioritized Experience Replay
- 5 Practical Tips

# Review

- ▶ Online algorithm like Q-learning in tabular case
- ▶ No sequential correlation in data samples
- ▶ Some stability with respect to gradient updates



- ▶ In an online setting, use  $\epsilon$ -greedy policy to periodically feed the buffer with newer experiences
- ▶ Use FIFO like mechanism to maintain size
- ▶ Sample a random minibatch of transitions ( $B$  transitions) to perform gradient descent (random sampling ensure samples for SGD are no longer correlated)
- ▶ Variance of the gradient estimate is also low compared to gradient computed using one sample

- ▶ Use an older set of weights to compute the targets
- ▶ Called **Target Network**
- ▶ Loss term is given by

$$L_i(\phi_i) = \left[ \mathbb{E}_{(s,a,r,s') \in D} \left( Q_{\phi_i}(s,a) - \underbrace{r + \max_{a'} Q_{\phi'_i}(s',a')}_{\text{target}} \right)^2 \right]$$


- ▶ Target network is kept constant for a while (every  $C$  steps) before being changed
  - ★ Every  $C$  steps the weights of the original network is copied to target network

---

## Algorithm DQN Algorithm

---

- 1: Initialize replay memory  $D$  to capacity  $N$
  - 2: Initialize action value function  $Q$  with parameters  $\phi$
  - 3: Initialize target action value function  $\hat{Q}$  with parameters  $\phi' = \phi$
  - 4: **for** episodes = 1 to  $M$  **do**
  - 5:     Initialize start state  $s_1$
  - 6:     **for** steps  $t = 1$  to  $T$  **do**
  - 7:         Select action  $a_t$  using  $\epsilon$ -greedy policy
  - 8:         Execute action  $a_t$  and store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $D$
  - 9:         Sample random minibatch (size  $B$ ) of transitions from  $D$
  - 10:        **for**  $b = 1$  to  $B$  **do**
  - 11:          Calculate targets for each transitions (Bellman backup or reward)
  - 12:        **end for**
  - 13:        Perform a gradient descent step on  $(y_i - Q_\phi(s_t, a_t))^2$  w.r.t  $\phi$
  - 14:        Every  $C$  steps set  $\hat{Q} = Q$
  - 15:     **end for**
  - 16: **end for**
-



## Polyak Averaging

- Replace target network update step (Step 14) by

$$\phi' : \phi' \leftarrow \tau \phi' + (1 - \tau) \phi$$

- Typical value for  $\tau = 0.99$




# Efficacy of DQN Algorithm

- ▶ Mnih et al. introduced Deep Q-Network (DQN) algorithm, applied it to ATARI games
- ▶ Used deep learning / ConvNets, published in early stages of deep learning craze (one year after AlexNet)
- ▶ Popularized ATARI (Bellemare et al., 2013) as RL benchmark
- ▶ Outperformed baseline methods, which used hand-crafted features

---

<sup>2</sup>Slide content from Schulman


# DQN on Atari <sup>2</sup>




Pong



Enduro




Beamrider



Q\*bert


- ▶ 49 ATARI 2600 games
- ▶ From pixels to actions
- ▶ The change in score is the reward
- ▶ Same algorithm
- ▶ Same function approximator
- ▶ Same hyperparameters
- ▶ Roughly human-level performance on 29 out of 49 games

<sup>2</sup>Slide content from Minh



- Convolutional neural network architecture
- History of 4 frames as input
- One output per action ( $Q(s, a)$ ) – expected reward for action  $a$

# Profile of Q Function Approximator



# Demonstration - Ping Pong

Random Policy


After 5.2 Millon Epochs

# Demonstration - Ping Pong

After 8 Million Epochs


After 9.5 Millon Epochs

# Are the Q-Values Meaningful ?





# On Tracking the Training Process



# Double DQN

- ▶ Consider single-state MDP with 2 actions.
- ▶ Both actions have zero mean rewards (the agent does not know this information)
- ▶ Let  $\hat{Q}(\cdot, a_1)$  and  $\hat{Q}(\cdot, a_2)$  be (unbiased) finite sample estimates of  $Q$  for action  $a_1$  and  $a_2$  respectively
- ▶ The agent will prefer the action which has maximum  $\hat{Q}$  based on sample estimates, although both actions have same expected mean reward

- ▶ Consider single-state MDP with 2 actions.
- ▶ One action has  $-\epsilon$  ( $\epsilon$  positive and small ) mean reward and the other action has zero mean reward.
- ▶ Let  $\hat{Q}(\cdot, a_1)$  and  $\hat{Q}(\cdot, a_2)$  be (unbiased) finite sample estimates of  $Q$  for action  $a_1$  and  $a_2$  respectively
- ▶ In this case, the agent might choose action  $a_1$  , depending on how the maximum  $\hat{Q}$  values that is based on sample estimates show up, although action  $a_2$  is clearly better in expectation

Extending the analogy, the (tabular or deep) Q-learning algorithm may actually pick the suboptimal action for target computation and this causes over estimation of Q-values.


- ▶ The problem with vanilla tabular Q-Learning is that the same samples are being used to decide which action is the best (highest expected reward), and the same samples are also being used to estimate that action-value
- ▶ Break up the Q-learning update rule as follows

$$Q(s_t, a) \leftarrow Q(s_t, a) + \alpha (r_{t+1} + \gamma Q(s_{t+1}, \arg \max Q(s_{t+1}, a)) - Q(s_t, a))$$


- ▶ If action value is overestimated, then it is chosen as the best action, and its overestimated value is used as the target

# Persistence of the Problem in DQN

- ▶ For transition  $(s_t, a_t, r_{t+1}, s_{t+1})$  the TD target of the  $Q$ -value update is  $r_{t+1} + \gamma \arg \max_{a'} Q_{\phi}(s_{t+1}, a')$
- ▶  $Q_{\phi}(\cdot, \cdot)$  is a noisy estimate during training phase
- ▶ Therefore,  $\max Q_{\phi}(\cdot, \cdot)$  would typically be overestimated during training



# Evidence from Atari Games



- ▶ There are two identical fair coins (we don't know they are fair)
- ▶ If a coin lands on head, we get one dollar; otherwise we lose a dollar
- ▶ Interested in answering the following questions
  - ★ Which coin will yield more money in future flips ?
  - ★ How much can we expect to win or lose per flip using the coin from previous question ?
- ▶ Two ways to answer
  - ★ Flip each coin  $n$  few times and answer both questions
  - ★ Flip each coin  $n_1$  times, answer the first question; collect fresh  $n_2$  samples to answer second question based on the answer to the first question

The idea behind the second method is that we use separate samples to choose the best action and separate samples to use Q-values



- ▶ Have two different set of samples to decide the action and to evaluate the target
- ▶ The idea is to use two  $Q$  functions
- ▶ In the tabular  $Q$ -learning setting, for each transition quadruple  $(s_t, a_t, r_{t+1}, s_{t+1})$  we flip a fair coin to decide any of the two update steps given below,

$$Q_1(s_t, a_t) \leftarrow Q_1(s_t, a_t) + \alpha (r_{t+1} + \gamma Q_2(s_{t+1}, \arg \max Q_1(s_{t+1}, a)) - Q_1(s_t, a_t))$$

$$Q_2(s_t, a_t) \leftarrow Q_2(s_t, a_t) + \alpha (r_{t+1} + \gamma Q_1(s_{t+1}, \arg \max Q_2(s_{t+1}, a)) - Q_2(s_t, a_t))$$

- ▶ In the DQN setting, we already have two  $Q$  networks (to address the moving target problem). So, we can take advantage of that by using the following update rule.
- ▶ In Double DQN, targets for the transition  $(s_t, a_t, r_{t+1}, s_{t+1})$  is computed as follows,

$$Q^{original}(s_t, a_t) \leftarrow r + \gamma Q^{target}(s, \arg \max Q^{original}(s, a))$$

- ▶ In the above equation, the  $\leftarrow$  actually means assigning targets which can then be picked up while sampling the replay buffer
- ▶ The fundamental idea is to use two  $Q_\phi$ 's (both are noisy estimates of true  $Q$ ) in different ways so that the overestimation problem mellows down

# Prioritized Experience Replay

- ▶ Replaying all transitions with equal probability is suboptimal
- ▶ Replay transitions in proportion to absolute Bellman error

$$\left| r + \gamma \max_{a'} Q_{\phi'}(s', a') - Q_{\phi}(s, a) \right|$$

- ▶ Leads to much faster learning

- ▶ TD error for vanilla DQN is

$$\delta_i = r_t + \gamma \max_{a \in \mathcal{A}} Q_{\phi'}(s_{t+1}, a) - Q_{\phi}(s_t, a_t)$$

- ▶ TD error for DDQN is

$$\delta_i = r_t + \gamma Q_{\phi'}(s_{t+1}, \operatorname{argmax}_{a \in \mathcal{A}} Q_{\phi}(s_{t+1}, a)) - Q_{\phi}(s_t, a_t)$$

- ▶ Priority for each entry in replay buffer  $D$  is given by  $p_i = |\delta_i| + \epsilon$



- ▶ Sample from replay buffer according to probability distribution

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}$$

with  $\alpha$  determining the level of prioritization

- ▶ In order to compute the expectation

$$\min_{\phi} \mathbb{E}_{(s_t, a_t, r_t, s_{t+1}) \sim D} \left[ \left( r_t + \gamma \max_{a \in \mathcal{A}} Q_{\phi'}(s_{t+1}, a) - Q_{\phi}(s_t, a_t) \right)^2 \right],$$

it is essential to use the importance sampling weights in each mini-batch of the gradient update

$$w_i = \left( \frac{1}{N} \cdot \frac{1}{P(i)} \right)^{\beta}$$

- ▶ The parameter  $\beta$  is an annealing term that is low (0.4 to 0.8) in the beginning of the training and tends to 1 towards the end of the training.
- ▶ The question on optimal choices of  $\alpha$  and  $\beta$  during various phases of training depends on task at hand

# Practical Tips

- ▶ DQN is more reliable on some tasks than others. Test your implementation on reliable tasks like Pong and Breakout: if it doesn't achieve good scores, something is wrong
- ▶ Large replay buffers improve robustness of DQN, and memory efficiency is important
- ▶ DQN converges slowly - for ATARI it is often necessary to wait for 10-40 million frames (couple of hours to a day of training on GPU) to see results significantly better than random policy. **Be Patient**
- ▶ Always run at least two different seeds when experimenting
- ▶ Learning rate scheduling is beneficial. Try high learning rates in initial exploration period

---

<sup>2</sup>Slide content from Schulman



## Practical Tips for DQN <sup>2</sup>

- ▶ Try non-standard exploration schedules
- ▶ Do use Double DQN with prioritized experience replay – significant improvement
- ▶ Use Huber loss on Bellman error

$$L_{\delta}(a) = \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta, \\ \delta(|a| - \frac{1}{2}\delta), & \text{otherwise.} \end{cases}$$

