

Policy Gradients Methods

Easwar Subramanian

TCS Innovation Labs, Hyderabad

Email : easwar.subramanian@tcs.com / cs5500.2020@iith.ac.in

October 21, 2021

- 1 Review
- 2 Policy Gradient Formulation
- 3 Variance Reduction Techniques

Review

- Last week, we parametrized value functions using parameter ϕ

$$V_{\phi}^{\pi}(s) = V^{\pi}(s)$$

$$Q_{\phi}^{\pi}(s, a) = Q^{\pi}(s, a)$$

- Policy was directly generated from value functions (greedy or ϵ greedy)

$$\pi_*(a|s) = \begin{cases} 1 & \text{if } a = \arg \max_{a \in \mathcal{A}} Q_*(s, a) \\ 0 & \text{Otherwise} \end{cases}$$


- In the next couple of lectures, we will directly parametrize the policy

$$\pi_{\theta}(a|s) = P(a|s, \theta)$$


- We will consider model free control with parametrized policies

Why Policy Optimization ?

- Often policies (π) are simpler than value functions (V or Q)



- Computing optimal V is bit of problem (we did not see any control algorithms for V)
- With state-value functions Q , computing $\arg \max$ over actions gets tricky when action space is large or continuous
- Better convergence properties
- Can learn stochastic policies



- If action space is discrete
 - ★ Network could output a vector of probabilities (softmax)
- If action space is continuous
 - ★ Network could output the parameters of a distribution (For e.g., mean and variance of a Gaussian)

A policy $\pi(\cdot)$ is parametrized by parameter θ and denoted by π_θ

Performance of a policy π_θ is given by

$$J(\theta) = V^{\pi_\theta}(s) = \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t r_{t+1} | s_0 = s \right]$$

Goal of RL is to find a policy

$$\pi_\theta^* = \arg \max_{\pi_\theta} V^{\pi_\theta}(s) = \arg \max_{\pi_\theta} \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t r_{t+1} | s_0 = s \right]$$


We will look for π_θ^* in class of stochastic policies by finding θ that maximizes $J(\theta)$

- ▶ Let $J(\theta)$ be the policy objective function
- ▶ Policy gradient algorithms search for a local maximum in $J(\theta)$ by ascending the gradient of the policy, w.r.t. parameters θ

$$\Delta\theta = \alpha \nabla_{\theta} J(\theta)$$

- ▶ $\nabla_{\theta} J(\theta)$ is the policy gradient and
- ▶ α is the step size parameter

Policy Gradient Formulation



- ▶ Let policy π be parametrized by θ and denoted by π_θ
- ▶ Let $\tau \sim \pi_\theta$ denote the state-action sequence given by $s_0, a_0, s_1, a_1, \dots, s_t, a_t, \dots$
- ▶ Then, $P(\tau; \theta)$ be the probability of finding a trajectory τ with policy π_θ

$$P(\tau; \theta) = P(s_0) \prod_{t=0}^{\infty} \pi_\theta(a_t | s_t) P(s_{t+1} | s_t, a_t)$$

We can define $G(\tau)$ discounted cumulative reward obtained by following trajectory τ

$$G(\tau) = \sum_{t=0}^{\infty} \gamma^t r_{t+1}$$

Objective function $J(\theta)$ for policy gradient approach is written as,

$$J(\theta) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{t=0}^{\infty} \gamma^t r_{t+1} | \pi_{\theta} \right] = \sum_{\tau \sim \pi_{\theta}} [P(\tau; \theta) G(\tau)]$$

Goal is to find θ^* such that

$$\theta^* = \arg \max_{\theta} J(\theta)$$

$$J(\theta) = \sum_{\tau} [P(\tau; \theta)G(\tau)]$$

Taking gradient with respect to θ gives


$$\begin{aligned}\nabla_{\theta} J(\theta) &= \nabla_{\theta} \left(\sum_{\tau} [P(\tau; \theta)G(\tau)] \right) \\ &= \sum_{\tau} \nabla_{\theta} [P(\tau; \theta)G(\tau)] \\ &= \sum_{\tau} \frac{P(\tau; \theta)}{P(\tau; \theta)} [\nabla_{\theta} P(\tau; \theta)] G(\tau) \\ &= \sum_{\tau} \frac{\nabla_{\theta} P(\tau; \theta)}{P(\tau; \theta)} P(\tau; \theta) G(\tau) \\ &= \sum_{\tau} \nabla_{\theta} \log P(\tau; \theta) P(\tau; \theta) G(\tau) \quad \left(\because \nabla_{\theta} \log f(x) = \frac{\nabla_{\theta} f(x)}{f(x)} \right)\end{aligned}$$

$$\nabla_{\theta} J(\theta) = \sum_{\tau \sim \pi_{\theta}} \nabla_{\theta} \log P(\tau; \theta) P(\tau; \theta) G(\tau) = \mathbb{E}_{\tau \sim \pi_{\theta}} [\nabla_{\theta} \log P(\tau; \theta) G(\tau)]$$

Sample based estimate is given by

$$\nabla_{\theta} J(\theta) \approx \frac{1}{K} \sum_{i=1}^K \nabla_{\theta} \log P(\tau^{(i)}; \theta) G(\tau^{(i)})$$

Policy Gradient : Intuition



- ▶ Increase the probability of paths with positive $G(\tau)$
- ▶ Decrease the probability of paths with negative $G(\tau)$
- ▶ Formalize the notion of 'trial and error'

$$\nabla_{\theta} J(\theta) \approx \frac{1}{K} \sum_{i=1}^K \nabla_{\theta} \log P(\tau^{(i)}; \theta) G(\tau^{(i)})$$

Is the above formula good enough for implementation ?

$$P(\tau; \theta) = P(s_0) \prod_{t=0}^{\infty} \pi_{\theta}(a_t | s_t) P(s_t | s_{t-1}, a_t)$$

$$\begin{aligned} \nabla_{\theta} \log P(\tau^{(i)}; \theta) &= \nabla_{\theta} \log \left[\prod_{t=0}^{\infty} \underbrace{P(s_{t+1}^{(i)} | s_t^{(i)}, a_t^{(i)})}_{\text{dynamics model}} \cdot \underbrace{\pi(a_t^{(i)} | s_t^{(i)})}_{\text{policy}} \right] \\ &= \nabla_{\theta} \left[\sum_{t=0}^{\infty} \log P(s_{t+1}^{(i)} | s_t^{(i)}, a_t^{(i)}) + \sum_{t=0}^{\infty} \log \pi(a_t^{(i)} | s_t^{(i)}) \right] \\ &= \nabla_{\theta} \sum_{t=0}^{\infty} \log \pi(a_t^{(i)} | s_t^{(i)}) = \underbrace{\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t^{(i)} | s_t^{(i)})}_{\text{policy gradient}} \end{aligned}$$

The following formulation provides an unbiased estimate of the policy gradient and we can calculate it without using the dynamics model

$$\begin{aligned}\nabla_{\theta} J(\theta) &\approx \frac{1}{K} \sum_{i=1}^K \left[\nabla_{\theta} \log P(\tau^{(i)}; \theta) \right] G(\tau^{(i)}) \\ \nabla_{\theta} J(\theta) &\approx \frac{1}{K} \sum_{i=1}^K \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t^{(i)} | s_t^{(i)}) \right] \left[\sum_{t=0}^{\infty} \gamma^t r_{t+1}^{(i)} \right]\end{aligned}$$

Algorithm REINFORCE : MC based Policy Gradient

- 1: Initialize policy network π with parameters θ_1 and learning rate α
- 2: **for** $n = 1$ to N **do**
- 3: Sample K trajectories from π_{θ_n}
- 4: Calculate gradient estimate

$$\nabla_{\theta_n} J(\theta_n) \approx \frac{1}{K} \sum_{i=1}^K \left[\sum_{t=0}^{\infty} \nabla_{\theta_n} \log \pi(a_t^{(i)} | s_t^{(i)}) \right] \left[\sum_{t=0}^{\infty} \gamma^t r_{t+1}^{(i)} \right]$$

- 5: Perform gradient update

$$\theta_{n+1} = \theta_n + \alpha \nabla_{\theta_n} J(\theta_n)$$

- 6: **end for**
-

Policy Gradient

$$\nabla_{\theta} J(\theta) \approx \frac{1}{K} \sum_{i=1}^K \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t^{(i)} | s_t^{(i)}) \right] \left[\sum_{t=0}^{\infty} \gamma^t r_{t+1}^{(i)} \right]$$

Maximum Likelihood

$$\nabla_{\theta} J_{ML}(\theta) \approx \frac{1}{K} \sum_{i=1}^K \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t^{(i)} | s_t^{(i)}) \right]$$

(Supervised Learning : Given s_t find a_t)

- ▶ The gradient estimate, thus calculated, is unbiased but has high variance (reason : we are sampling stochastic paths)
- ▶ Hence the gradient descent is slow to converge
- ▶ Some variance reduction techniques are required in practice

Variance Reduction Techniques

Discount Factor and Variance Reduction

Gradient estimate is given by,

$$\nabla_{\theta} J(\theta) \approx \frac{1}{K} \sum_{i=1}^K \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t^{(i)} | s_t^{(i)}) \right] \left[\sum_{t=0}^{\infty} \gamma^t r_{t+1}^{(i)} \right]$$

One can rewrite the above equation as

$$\nabla_{\theta} J(\theta) \approx \frac{1}{K} \sum_{i=1}^K \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t^{(i)} | s_t^{(i)}) \cdot \left[\sum_{k=0}^{\infty} \gamma^k r_{k+1}^{(i)} \right] \right]$$

- ▶ For infinite horizon MDPs having $\gamma < 1$ not only helps in proving convergence of algorithms but also helps reduce variance of the policy gradient estimate
- ▶ Ignoring reward terms 'far' into the future gives us a reasonable approximation to policy gradient but with lower variance

Score function in policy gradient is the term

$$\nabla_{\theta} \log \pi(a_t | s_t)$$

Expectation of the score function is zero


$$\begin{aligned} \mathbb{E}_{a_t | s_t} [\nabla_{\theta} \log \pi(a_t | s_t)] &= \int_{a_t} \pi(a_t | s_t) \nabla_{\theta} \log \pi(a_t | s_t) da_t \\ &= \int_{a_t} \nabla_{\theta} \pi(a_t | s_t) da_t \\ &= \nabla_{\theta} \int_{a_t} \pi(a_t | s_t) da_t \\ &= \nabla_{\theta} 1 = 0 \end{aligned}$$


Causality : Policy at time t' cannot affect reward at time t when $t < t'$.

- When we take an action at timestep t , it can only affect the rewards from timesteps t and onwards.

Recall that,

$$\nabla_{\theta} J(\theta) \approx \frac{1}{K} \sum_{i=1}^K \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t^{(i)} | s_t^{(i)}) \cdot \left[\sum_{k=0}^{\infty} \gamma^k r_{k+1}^{(i)} \right] \right]$$





$$G(\tau) = \sum_{t=0}^{\infty} \gamma^t r_{t+1}$$

Let $\tau_{a:b}$ denote the states and actions visited from time a to b and

$$G_{a:b}(\tau) = \sum_{t=a}^b \gamma^t r_{t+1}$$

Therefore for any time t , we have,

$$G(\tau) = G_{0:t-1}(\tau) + G_{t-1:\infty}(\tau)$$

Figure Source:
Jie-Han-Chen:SlideShare

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t | s_t) \cdot \left[\sum_{k=0}^{\infty} \gamma^k r_{k+1} \right] \right] \\&= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t | s_t) G(\tau) \right] \\&= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t | s_t) G_{0:t-1}(\tau) + \sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t | s_t) G_{t:\infty}(\tau) \right] \\&= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t | s_t) G_{0:t-1}(\tau) \right] \\&\quad + \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t | s_t) G_{t:\infty}(\tau) \right]\end{aligned}$$

Consider evaluating the expectation of the first term

$$\begin{aligned}\mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t | s_t) G_{0:t-1}(\tau) \right] &= \left[\sum_{t=0}^{\infty} G_{0:t-1}(\tau) \mathbb{E}_{\pi_{\theta}} \nabla_{\theta} \log \pi(a_t | s_t) \right] \\ &= \sum_{t=0}^{\infty} G_{0:t-1} \cdot 0 = 0\end{aligned}$$


Therefore, the policy gradient estimate with temporal structure is given by

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t | s_t) G_{t:\infty}(\tau) \right]$$


The sample estimate of the gradient expression is given by

$$\nabla_{\theta} J(\theta) \approx \frac{1}{K} \sum_{i=1}^K \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi(a_t^{(i)} | s_t^{(i)}) \cdot \left[\sum_{t'=t}^{\infty} \gamma^{t'-t} r_{t'+1}^{(i)} \right] \right]$$

- The above policy gradient estimate with temporal structure is also an unbiased estimate of the true policy gradient but has **lower variance** since it has '*thrown out*' a few terms



Need for a Baseline



What if all paths have positive reward sum ?

Can we subtract a baseline without biasing the gradient ?


Let $b(s_t)$ be a baseline that is conditioned on s_t . Then,

$$\mathbb{E}_{a_t|s_t} [b(s_t) \nabla_{\theta} \log \pi(a_t|s_t)] = b(s_t) \mathbb{E}_{a_t|s_t} [\nabla_{\theta} \log \pi(a_t|s_t)] = 0$$

Therefore,

$$\begin{aligned} \nabla_{\theta} J(\theta) &= \mathbb{E}_{\tau \sim \pi_{\theta}} [\nabla_{\theta} \log \pi(a_t|s_t) \cdot G_{t:\infty}(\tau)] \\ &= \mathbb{E}_{\tau \sim \pi_{\theta}} [\nabla_{\theta} \log \pi(a_t|s_t) \cdot G_{t:\infty}(\tau)] - \mathbb{E}_{\tau \sim \pi_{\theta}} [b(s_t) \nabla_{\theta} \log \pi(a_t|s_t)] \\ &= \mathbb{E}_{\tau \sim \pi_{\theta}} [\nabla_{\theta} \log \pi(a_t|s_t) \cdot [G_{t:\infty}(\tau) - b(s_t)]] \end{aligned}$$

Need for a Baseline



A good choice for baseline :

$$b = \mathbb{E}(G(\tau)) \approx \frac{1}{K} \sum_{i=1}^K G(\tau^{(i)})$$

- Constant Baseline

$$b = \mathbb{E}(G(\tau)) \approx \frac{1}{K} \sum_{i=1}^K G(\tau^{(i)})$$

- Time Dependent Baseline

$$b_t = \frac{1}{K} \sum_{i=1}^K G_{t:\infty}(\tau^{(i)})$$

- Optimal Baseline


$$b = \frac{\mathbb{E}_{\tau}(\nabla_{\theta} \log \pi(a_t | s_t)^2 G_{t:\infty}(\tau))}{\mathbb{E}_{\tau}(\nabla_{\theta} \log \pi(a_t | s_t)^2)}$$

- State dependent expected return

$$b(s) = \mathbb{E}_{\pi_{\theta}}[r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | s_t = s] = V^{\pi}(s)$$

State dependent expected return

$$b(s) = \mathbb{E}_{\pi_{\theta}}[r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \cdots | s_t = s] = V^{\pi}(s)$$



Algorithm Vanilla Policy Gradient Algorithm

- 1: Initialize policy network π with parameters θ_1 learning rate α and baseline b
- 2: **for** $n = 1$ to N **do**
- 3: Sample K trajectories by executing the policy π_{θ_n}
- 4: At each time step of each trajectory compute $G_t = \sum_{t'=t}^{\infty} \gamma^{t'-t} r_{t+1}$ and advantage estimate $A_t = G_t - b(s_t)$
- 5: Calculate gradient estimate

$$\nabla_{\theta_n} J(\theta_n) \approx \frac{1}{K} \sum_{i=1}^K \left[\sum_{t=0}^{\infty} \nabla_{\theta_n} \log \pi(a_t^{(i)} | s_t^{(i)}) A_t \right]$$

- 6: Perform gradient update

$$\theta_{n+1} = \theta_n + \alpha \nabla_{\theta_n} J(\theta_n)$$

- 7: **end for**
-

- ▶ The REINFORCE and Vanilla policy gradient as described above is on-policy
 - ★ There is an off-policy way to do policy gradient algorithms
- ▶ We do learning by Monte-Carlo roll-outs
 - ★ Will be addressed by Actor-Critic method