

28/09/2021

CS 6160 Cryptology Lecture 7: Practical Constructions of Block Ciphers

Maria Francis

September 28, 2021

Block Ciphers

- A block cipher must behave like a random permutation.
- On ℓ -bit strings, there $2^\ell!$ permutations possible and to represent a permutation we need $\ell \cdot 2^\ell$ bits.
- For $\ell > 20$ it is impractical and infeasible for $\ell > 50$. For modern block ciphers block length ℓ is ≥ 128 .
- We need a set of permutations with a concise description (a short key) but still behaves like a random permutation.
- Changing one bit in the input of E_k should yield an independent result, should affect every bit of output.
- Does not mean every output bit should be changed! That is not random. Every bit is changed with probability roughly half.

Confusion-Diffusion paradigm

- Shannon again! He gave a way for constructing concise random-looking permutations.
- We need to construct a random looking permutation E with a large block length from many random looking (or random) permutations $\{f_i\}$ with small block length.
- For e.g: We need to build E with block length 128 bits.
- The key k for E will specify 16 permutations f_1, \dots, f_{16} that have 8-bit block length.
- Given $x \in \{0, 1\}^{128}$, we parse it as 16-bytes $x_1 \dots x_{16}$ and then define,

$$E_K(x) = f_1(x_1) \circ \dots \circ f_{16}(x_{16}).$$

$\{f_i\}$ s introduce *confusion* into F .

Confusion-Diffusion paradigm

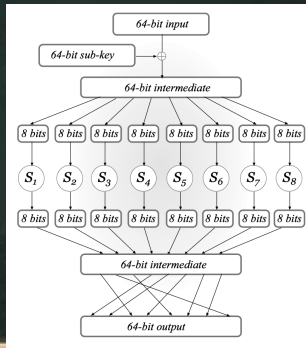
- Note that permutation on 8 bits can be represented using $\log(2^8!) \approx 1600$ bits much smaller than $128 \cdot 2^{128}$ bits.
- But E is not pseudorandom. For eg: if x and x' differ only in their first bit then $E_k(x)$ and $E_k(x')$ will differ only in their first byte regardless of the key.
- This is why we need diffusion, output bits are permuted using a mixing permutation to spread a local change everywhere in the block.
- The confusion diffusion steps constitute a round and such rounds are repeated several times.

Substitution-Permutation Networks (SPN)

- SPN implements the confusion-diffusion paradigm.
- For round functions we give it a particular form rather than choosing all possible permutations on say 8-bit strings.
- We fix a **substitution function** S called **S-box** and using k we define $S(k \oplus x)$.
- For example: consider a SPN with a 64-bit block length based on a collection of 8-bit S -boxes S_1, \dots, S_8 .
- Each round we do the following operations on 64-bit input m :
 - ▶ **Key mixing**: Set $x = m \oplus k$, where k is the current-round sub-key.
 - ▶ **Substitution**: Set $x = S_1(x_1) \circ \dots \circ S_8(x_8)$ where x_i is the i th byte of x .
 - ▶ **Permutation**: Permute all the bits of x to obtain the output of the round which is the input of next round.

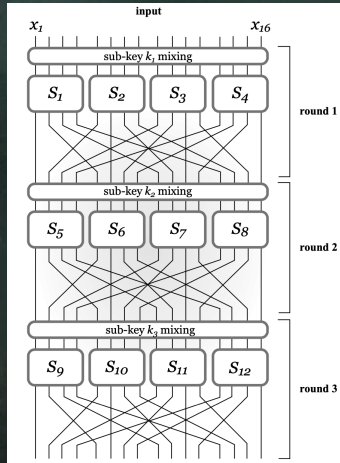
SPN rounds

- After the final round, there is also a **key-mixing step** and then we get the final output of the cipher.
- Kerckhoff's principle: *S*-boxes and mixing permutations are public.
- The key-mixing step is key, *the substitution-permutation steps add no security as no key is involved*.



SPN multiple rounds

High level structure of an SPN with 16-bit block length and 4-bit S-boxes in each round.



SPN

- The actual key of the block cipher is called master key.
- For each round a round key is generated. They derived from the master key using a **key schedule**.
- It can be a simple algorithm : take different subsets of the bits of the master key. Or could be a complex relation.
- An r -round SPN has **r full rounds of key-mixing, S-box substitution, application of a mixing permutation and teh final key-mixing step.**
- We need then $r + 1$ sub-keys.

SPN

- SPN is invertible given the key, i.e. given the output of a SPN and the key we can retrieve the input.
- Note that if a single round is inverted then the entire SPN can be inverted.
- How to invert a single round?
 - ▶ invert the mixing permutation, which is just rearrangement of bits
 - ▶ S -boxes are permutations too and therefore one-one, which means it can be inverted too.
 - ▶ The result XOR with sub-key gives original input of the round.

Thus we have:

Let E be a keyed function defined by an SPN in which S -boxes are permutations. Then regardless of the key schedule and the number of rounds E_k is a permutation for any k .

Security of SPN – Avalanche Effect

- The number of rounds, design of S -boxes, mixing permutations, key scheduling algo are important in determining the security.
- **Avalanche effect for block ciphers** : a small change in the input must “affect” every bit of the output.
- How to implement it in SPN ?
 - ▶ S -boxes are designed so that changing a single bit of its input changes at least **two bits** of the S -box.
 - ▶ Mixing permutations are such that **the output bits of any given S -box are used as input to multiple S -boxes in the next round.**

How does it yield the avalanche effect?

- Assume the S -boxes have 8 bits input and output and block length of the cipher is 128 bits.
- Consider what happens when you are giving two inputs that differ in a single bit.
- After the first round, **the intermediate values differ in exactly two bit positions**. How?
 - ▶ XORing with sub-key maintains 1-bit difference, so the input to S -box is identical **except for one**.
 - ▶ For that S -box since there is 1-bit difference, the output will generate a 2-bit difference.
 - ▶ The mixing permutation applied to the results rearranges positions but maintains the 2-bit difference.

How does it yield the avalanche effect?

- After the first round, the mixing permutation spreads the two bit-positions to **two different S-boxes in the second round**.
- This is the case even after XORing with the sub-key for the next round.
- Now in the second round **two S-boxes receive inputs differing in single input**, thus resulting in an intermediate value that **differ in 4 bits**.
- Same argument means 8 bits after 3 rounds, 16 bits after 4th round and all 128 bits after 7th round.
- The arguments are not that precise and there could be fewer differences than expected at the end of some round, so we need to use more than 7 rounds.
- But 7 gives us a lower bound: less than that and we can distinguish this cipher from a random permutation.

How to design S -boxes?

- Choose them at random? Not the best way!
 - ▶ Let a S -box S be s.t. it takes 4-bit inputs and let x and x' be two distinct values with $y = S(x)$
 - ▶ Consider choosing uniform $y' (\neq y)$ as $y' = S(x')$.
 - ▶ There are 4 strings that differ from y in only 1 bit and so $\frac{4}{15}$ is the probability that we will choose y' that does not differ from y in more than 1 bit.
 - ▶ When we consider all pairs that differ in exactly one bit, the probability is compounded.
- So S -boxes have to be carefully designed!
- For a block cipher to be strongly pseudorandom, the avalanche effect should apply to its inverse as well.

Changing every bit of the output should affect every bit of the input. Answer: More rounds!

Feistel Networks

- Another approach for constructing block ciphers.
- Major advantage/difference over SPN: the underlying functions (analogous to S -boxes) need not be invertible.
- It gives a way to construct an invertible function from non-invertible components.
- This helps give the block cipher an “unstructured” behavior as requiring all the components to be invertible introduces structure.
- Also invertibility makes designing S -boxes hard.

Feistel Networks


- It operates with a series of rounds.
- In each round, a **keyed round function** is applied.
- They need not be invertible.
- Typically build from S -boxes and mixing permutations but a Feistel network can deal with **any round function**.
- Balanced Feistel networks –
 - ▶ i th round - \hat{f}_i takes sub-key k_i and $\ell/2$ **bit string and outputs an $\ell/2$ -bit string**.
 - ▶ Master key k is used to derive sub-keys k_i for each round i .

$$f_i : \{0, 1\}^{\ell/2} \rightarrow \{0, 1\}^{\ell/2}$$

$$f_i(x) = \hat{f}_i(k_i, x)$$

i th round of Feistel Networks

- The input ℓ -bit is divided into two $\ell/2$ bit halves: left: L_{i-1} and right: R_{i-1}
- Output (L_i, R_i) is $L_i = R_{i-1}$ and $R_i = L_{i-1} \oplus f_i(R_{i-1})$.
- For r -round Feistel Network, the ℓ -bit input is parsed as (L_0, R_0) and the output is ℓ -bit (L_r, R_r) .



Inverting a Feistel Network

- It is invertible **regardless of the $\{f_i\}$ (and thus \hat{f}_i)**
- To show that a Feistel network is invertible it is enough to show that each round of the network is invertible given $\{f_i\}$.
- Given the output (L_i, R_i) of the i th round, we show how to compute (L_{i-1}, R_{i-1}) :

$$R_{i-1} := L_i$$

$$L_{i-1} := R_i \oplus f_i(R_{i-1})$$

Note that we took f_i , not its inverse.

Let E be a keyed function defined by a Feistel network. Then regardless of the round functions $\{\hat{f}_i\}$ and the number of rounds, E_k is an efficiently invertible permutation for all k .

DES : Data Encryption Standard

- The most common example for a block cipher, but an old one (1970s) and is considered a legacy encryption scheme.
- Evolved from IBM's submission, **Lucifer**, to the NIST's request for encryption algorithms that could be standardized.
- A well-engineered algorithm with an immense influence in the field.
- Short $k = 56$ bits and $n = 64$ bits and has 16 rounds of a **Feistel network**.
- Triple DES (3DES) introduced in 1998 was used in finance and payment industries until July 2017 when NIST proposed retiring it. In Nov 2017, NIST allowed for **restricted usage**, no longer be used in TLS, IPsec or for large files.

DES

- General consensus: except for short key everything else is well-designed.
- Best attack **in practice** is exhaustive search over 2^{56} keys.
- There are theoretical attacks that we will discuss that assumes less computation but these attacks assume certain conditions that are difficult to realize in practice.
- We give a high-level overview of the design to understand the basic ideas but it is not a full specification!

Design of DES

- DES cipher has a 16 round Feistel cipher, block length $\ell = 64$ bits and key length $n = 56$ bits.
- The same round function \hat{f} is used in each round.
- Each round function takes a 48-bit sub-key and a 32-bit input (half a block) since it is a **balanced** Feistel network.
- Key scheduling algorithm, *KeySchedule* takes the 56-bit key and gives as output 48 bit length sub-keys, k_1, \dots, k_{16} .
- *KeySchedule* algorithm : gives k_i a permuted subset of 48 bits of the master key.
 - ▶ Divide 56 bit master key to left and right halves of 28 bits.
 - ▶ For the leftmost 24 bits of k_i , we take some subset of the 28 bits of the left half of master key.
 - ▶ And the same for the rightmost 24 bits of k_i .
 - ▶ How it is done is all public!

Key Scheduling of DES

The permutations applied on the master key:

<i>PC-1</i>								<i>PC-2</i>							
57	49	41	33	25	17	9		14	17	11	24	1	5		
1	58	50	42	34	26	18		3	28	15	6	21	10		
10	2	59	51	43	35	27		23	19	12	4	26	8		
19	11	3	60	52	44	36		16	7	27	20	13	2		
63	55	47	39	31	23	15		41	52	31	37	47	55		
7	62	54	46	38	30	22		30	40	51	45	33	48		
14	6	61	53	45	37	29		44	49	39	56	34	53		
21	13	5	28	20	12	4		46	42	50	36	29	32		

The *KeySchedule* algorithm:

```

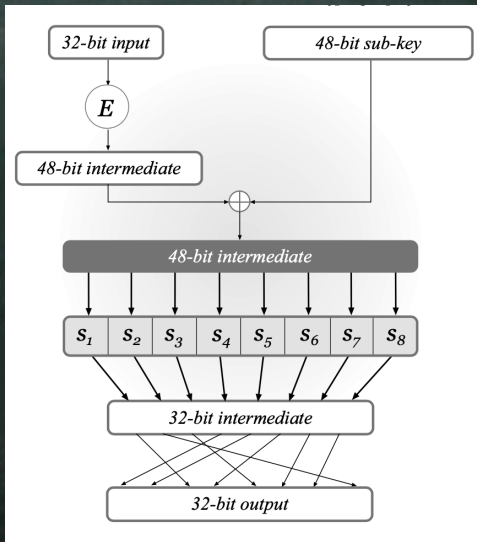
Algorithm KeySchedule( $K$ )  //  $|K| = 56$ 
   $K \leftarrow PC-1(K)$ 
  Parse  $K$  as  $C_0 \| D_0$ 
  for  $r = 1, \dots, 16$  do
    if  $r \in \{1, 2, 9, 16\}$  then  $j \leftarrow 1$  else  $j \leftarrow 2$  fi
     $C_r \leftarrow \text{leftshift}_j(C_{r-1})$  ;  $D_r \leftarrow \text{leftshift}_j(D_{r-1})$ 
     $K_r \leftarrow PC-2(C_r \| D_r)$ 
  return( $K_1, \dots, K_{16}$ )

```

DES Round Function

- Also called **DES mangler function** : a substitution-permutation network.
- Each round the input to \hat{f} is $k_i \in \{0, 1\}^{48}$ and $x \in \{0, 1\}^{32}$.
- We first expand x to a 48-bit value $x' (= \text{Exp}(x))$, by simply duplicating half the bits of x .
- We then do computation just like before in SPN:
 - ▶ x' is XORed with k_i which is 48 bits long.
 - ▶ Divide the result into 8 blocks each of 6 bits long.
 - ▶ Each 6-bit block is passed into a S -box and yields a 4-bit output.
 - ▶ Concatenate the result of 8 S -boxes and we have a 32-bit result.
 - ▶ Then a mixing permutation.
- Main difference: S -boxes are not invertible, inputs are longer than outputs.

DES Round Function



S-boxes in DES

- The eight *S*-boxes are the **non-linear part of DES** and they were very carefully designed. Even a small change makes DES vulnerable to attacks.
- Each *S*-box takes a 6-bit input and gives a 4-bit output.
- View it as a table with 4 rows and 16 columns with each entry a 4-bit entry.
- The $4 \cdot 16 = 64 = 2^6$ entries of the table can be indexed by the 6-bit input: 1 and 6 bits for row and the other bits for column.

S-boxes in DES

S ₁	:		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	0	0	14	4	13	1	2	15	11	8	3	10	6	12	5	9	0	7
	1	0	1	0	15	7	4	14	2	13	1	10	6	12	11	9	5	3
	2	1	0	4	1	14	8	13	6	2	11	15	12	9	7	3	10	5
	3	1	1	15	12	8	2	4	9	1	7	5	11	3	14	10	0	6
S ₂	:		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	0	0	15	1	8	14	6	11	3	4	9	7	2	13	12	0	5	10
	1	0	1	3	13	4	7	15	2	8	14	12	0	1	10	6	9	11
	2	1	0	0	14	7	11	10	4	13	1	5	8	12	6	9	3	2
	3	1	1	13	8	10	1	3	15	4	2	11	6	7	12	0	5	14
S ₃	:		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	0	0	10	0	9	14	6	3	15	5	1	13	12	7	11	4	2	8
	1	0	1	13	7	0	9	3	4	6	10	2	8	5	14	12	11	15
	2	1	0	13	6	4	9	8	15	3	0	11	1	2	12	5	10	14
	3	1	1	1	10	13	0	6	9	8	7	4	15	14	3	11	5	2
S ₄	:		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	0	0	7	13	14	3	0	6	9	10	1	2	8	5	11	12	4	15
	1	0	1	13	8	11	5	6	15	0	3	4	7	2	12	1	10	14
	2	1	0	10	6	9	0	12	11	7	13	15	1	3	14	5	2	8
	3	1	1	3	15	0	6	10	1	13	8	9	4	5	11	12	7	2
S ₅	:		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	0	0	2	12	4	1	7	10	11	6	8	5	3	15	13	0	14	9
	1	0	1	14	11	2	12	4	7	13	1	5	0	15	10	3	9	8
	2	1	0	4	2	1	11	10	13	7	8	15	9	12	5	6	3	0
	3	1	1	11	8	12	7	1	14	2	13	6	15	0	9	10	4	5
S ₆	:		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	0	0	12	1	10	15	9	2	6	8	0	13	3	4	14	7	5	11
	1	0	1	10	15	4	2	7	12	9	5	6	1	13	14	0	11	3
	2	1	0	9	14	15	5	2	8	12	3	7	0	4	10	1	13	11
	3	1	1	4	3	2	12	9	5	15	10	11	14	1	7	6	0	8
S ₇	:		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	0	0	4	11	2	14	15	0	8	13	3	12	9	7	5	10	6	1
	1	0	1	13	0	11	7	4	9	1	10	14	3	5	12	2	15	8
	2	1	0	1	4	11	13	12	3	7	14	10	15	6	8	0	5	9
	3	1	1	6	11	13	8	1	4	10	7	9	5	0	15	14	2	3
S ₈	:		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	0	0	13	2	8	4	6	15	11	1	10	9	3	14	5	0	12	7
	1	0	1	1	15	13	8	10	3	7	4	12	5	6	11	0	14	9
	2	1	0	7	11	4	1	9	12	14	2	0	6	10	13	15	3	5
	3	1	1	2	1	14	7	4	10	8	13	15	12	9	0	3	5	6

DES avalanche effect

- Changing one input bit to a S -box changes at least two output bits.
- The mixing permutation ensures that the **four output bits of any S -box will affect the input of 6 S -boxes of next round!**
- Consider two inputs (L_0, R_0) and (L'_0, R_0) which differ by only a single bit in the left half.
- After first round, the intermediate values only differ by **one bit** but now the difference is in the **right half**.
- If you assume that the bit is not duplicated in the expansion step, the input to S -box differs only in one bit again.
- This generates **a two bit difference after S -box computation**.
- This means after two rounds there is a total of 3 bits difference.

DES avalanche effect

- The mixing permutation spreads the two bit difference of right halves to **two different S-boxes**, resulting in a 4-bit difference in the right halves. And of course you have a 2-bit difference in left halves.
- As with a SPN, after 7 rounds all 32 bits in the right half and after 8 rounds all 32 bits in the left half are affected.
- With 16 rounds we are good!
- How to attack reduced rounds of both SPN and DES?
- Note that DES is a Feistel construction but has only **dependent weak pseudorandom round functions**.
- **Luby-Rackoff Theorem**: *Three rounds of Feistel construction with **independent PRFs as round functions** makes the block cipher a weak PRP and **four rounds makes it a strong PRP**.*

How to improve DES?

- Ultimately every attack for DES is exhaustive search of key space!
- Increase key size to 128 and then use a different key schedule but still choosing a 48-bit sub-key!
- Or change the *S*-boxes and use a larger sub-key for each round.
- Any small **seemingly insignificant** change and we are worried about its security!
- Preferred approach: use DES as a blackbox and try and build something that invokes this DES.

Double Encryption

- Let E be a block cipher with an n -bit key and ℓ -bit block length.
- We define a new cipher, E' with key length $2n$:

$$E'_{k_1, k_2}(x) := E_{k_2}(E_{k_1}(x)),$$

where k_1, k_2 are independent keys.

- For DES, we get a 2DES which takes a 112 bit key.
- Looks good since time 2^{112} is out of reach.
- **Meet-in-the-middle Attack** : an attack that needs time roughly 2^n .

Meet-in-the-middle Attack

Assume the adversary is given a single input/output pair (x, y) ,
 $y = E'_{k_1^*, k_2^*}(x) = E_{k_2^*}(E_{k_1^*}(x))$, for unknown k_1^*, k_2^* .

- For each $k_1 \in \{0, 1\}^n$, compute $z := E_{k_1}(x)$ and store (z, k_1) in a list L .
- For each $k_2 \in \{0, 1\}^n$, compute $z := E_{k_2}^{-1}(y)$ and store (z, k_2) in a list L' .
- We say $(z_1, k_1) \in L$ and $(z_2, k_2) \in L'$ are a match if $z_1 = z_2$. For each match add (k_1, k_2) to a set S .
- This attacks takes time $\mathcal{O}(n \cdot 2^n)$ and requires space $\mathcal{O}((n + \ell) \cdot 2^n)$.
- Also (k_1^*, k_2^*) is in S .
- With more input/output pairs we can identify the key with high probability.

Triple Encryption

What if I apply the block cipher three times?

- **Variant 1:** Choose three independent keys k_1, k_2, k_3 ,

$$E''_{k_1, k_2, k_3} = E_{k_3}(E_{k_2}^{-1}(E_{k_1}(x))).$$

- **Variant 2:** Choose three independent keys k_1, k_2, k_3 ,

$$E''_{k_1, k_2} = E_{k_1}(E_{k_2}^{-1}(E_{k_1}(x))).$$

Note that the second E is inverse permutation. This is not an issue since if E is a strong pseudorandom function then E^{-1} is also one.

Security of the Variants

- Variant 1: Key length is $3n$, but security is only 2^{2n} because we can do meet-in-the-middle attack.
- Not as secure as $3n$ keys but still secure for practical purposes.
- Variant 2: Key length is $2n$ and it gives security against attacks running in time 2^{2n} .
- Why use inverse? For backward compatibility. Set $k_1 = k_2 = k_3$ then we get a single invocation of E with k_1 .
- Triple DES (3DES) : Standardized in 1999 and widely used even now.
- It can be any of the two variants.
- It has relatively small block length and it is relatively slow.

AES – Advanced Encryption Standard

- In Jan 1997, NIST announced a competition for a new block cipher – AES.
- In Oct 2000, Rijndael algorithm designed by two Belgian cryptographers Rijmen and Daemen was the final choice and which later after some modifications became AES.
- Winner was selected based on efficiency, performance in hardware, flexibility, etc.
- AES block cipher has a 128-bit block length and can be used with 128, 192, 256-bit keys.
- The length of key affects the key schedule, the number of rounds but not the high-level.

AES – Advanced Encryption Standard

- AES uses SPN not Feistel network.
- During computation a 4-by-4 array of bytes called state is modified.
- The first state is defined by the input which is 128 bits or 16 bytes.
 1. **AddRoundKey**: Every round a 128-bit subkey is derived from the master key and is interpreted as a 4-by-4 array. XOR it with the sub-key to get new state array.
 2. **SubBytes**: Each byte of the state array is replaced by another byte according to a *S*-box, a bijection over $\{0,1\}^8$.
 3. **ShiftRows**: Bytes in each row of the state array are shifted to the left.
 4. **MixColumns**: An invertible transformation (linear transformation over a finite field, $GF(2^8)$) is applied to the four bytes in each column. Each input byte affects all four output bytes. Together with ShiftRows, it provides diffusion.

AES – Advanced Encryption Standard

- MixColumns: If two inputs differ in b bytes then the transformation yields outputs differing in at least $5 - b$ bytes.
- In the final round, MixColumns is replaced with AddRoundKey: to avoid the adversary from inverting last three stages which are not dependent on the key.
- Essentially AES is a SPN :
 - ▶ the input is XORed with the round key
 - ▶ a small, **invertible function** is applied to chunks of the resulting value
 - ▶ And then mix the bits of result to obtain diffusion.

Security of AES

- Number of rounds : 10 for 128 bit key, 12 for 192 bit key, 14 rounds for 256 bit key.
- To date, there are no practical cryptanalytic attacks that are significantly better than an exhaustive search for the key.
- But no simple hardness assumption for the security for AES.
- AES-256 is widely believed to be secure against quantum computers too!
 - ▶ In a quantum computer, the time taken for a search problem scales slower than for a classical one: search for an n -bit key is equiv to a $n/2$ -bit key. So AES-256 effectively has a 128-bit key against a quantum attacker, so still ok!
- AES-128, AES-192, AES-256 : in IPSec, TLS, iOS Keychain
- 3DES: some Mac OSes for their Keychain

AES and DES as PRPs

- DES:

$$K \times X \longrightarrow X, X = \{0, 1\}^{64}, K = \{0, 1\}^{56}$$

- AES-128/192/256:

$$K \times X \longrightarrow X, K = \{0, 1\}^{128/192/256}, X = \{0, 1\}^{128}$$

- These PRFs are not proven to be secure PRFs.
- They have been subjected to intense scrutiny by cryptographers so we assume they are secure PRFs.