

09/10/2020

# CS 6160 Cryptology Lecture 11: CCA-Security & Message Authentication Codes (MACs)

Maria Francis

October 09, 2020

# Defining CCA-security

- Chosen-ciphertext attack is even more powerful than eavesdropping and chosen-plaintext attacks.
- $\mathcal{A}$  has the ability to not only obtain encryptions of messages of its choice like CPA but also obtain decryptions of ciphertexts of its choice.
- Formally, that means  $\mathcal{A}$  has access to a decryption oracle as well as an encryption oracle.

# CCA indistinguishability experiment

$\text{PrivK}_{\mathcal{A}, \Pi}^{\text{cca}}(1^n)$ : We have the SKE  $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ ,  $\mathcal{A}$  and  $1^n$  security parameter.

1.  $k$  is generated by running  $\text{Gen}(1^n)$ .
2.  $\mathcal{A}$  has oracle access to  $\text{Enc}_k(\cdot)$  and  $\text{Dec}_k(\cdot)$ .
3. It outputs a pair of messages  $m_0, m_1$  of the same length.
4. A uniform bit  $b \in \{0, 1\}$  is chosen and then a **challenge ciphertext**,  $c \leftarrow \text{Enc}_k(m_b)$  is computed and given to  $\mathcal{A}$ .
5.  $\mathcal{A}$  continues to have oracle access to  $\text{Enc}_k(\cdot)$  and  $\text{Dec}_k(\cdot)$ . **But it is not allowed to query on the challenge ciphertext.**
6. Eventually,  $\mathcal{A}$  outputs a bit  $b'$ .
7. Output is 1 if  $b' = b$  and 0 otherwise. If output is 1 we say that  $\mathcal{A}$  succeeds.

## CCA-secure

Now that we have an indistinguishability experiment, we can have the security definition.

### Definition

A SKE  $\Pi$  has **indistinguishable encryptions under a chosen-ciphertext attack** or is **CCA-secure** if for PPT adversaries  $\mathcal{A}$  there is a negligible function  $\text{negl}$  such that:

$$\Pr[\text{PrivK}_{\mathcal{A}, \Pi}^{\text{cca}}(1^n) = 1] \leq \frac{1}{2} + \text{negl}(n).$$

Also, if a scheme has indistinguishable encryptions under a chosen-ciphertext attack then it has indistinguishable **multiple** encryptions under a chosen-ciphertext attack.

**Adversary has unlimited access to the decryption oracle except a request for the decryption of the challenge ciphertext itself.**

# CCA in the real-world

- The adversary may not get honest parties to decrypt arbitrary ciphertexts but it may be able to influence what gets decrypted.
- An adversary sends encrypted messages to the bank on the behalf of the user and see what the result is.
- For e.g what if the ciphertext corresponding to an ill-formed plaintext is sent and then the adversary may be able to deduce that from the bank's reaction.
- If encryption is part of the authentication protocol where one party sends a ciphertext to the other (using a PK), decrypts it (using SK) and returns the result to claim it is indeed him/her.  
The honest party being authenticated here is a decryption oracle!

# CCA secure schemes so far?

- No, none of them were CCA-secure.
- Consider the CPA-secure encryption scheme we discussed in last lecture:

$$Enc_k(m) = \langle r, F_k(r) \oplus m \rangle.$$

- Let  $\mathcal{A}$  run a CCA-indistinguishability exp with  $m_0 = 0^n$  and  $m_1 = 1^n$ .
- On seeing  $c = \langle r, s \rangle$ ,  $\mathcal{A}$  flips the first bit of  $s$  and asks for decryption of  $c'$ .
- $c \neq c'$  and so this query is allowed.
- Decryption oracle returns either  $10^{n-1}$  (in which case  $b = 0$ ) or  $01^{n-1}$  (in which case  $b = 1$ ).

## CCA secure schemes so far?

- The problem was we were able to manipulate the ciphertexts in a way that you still got a meaningful ciphertext with a relation to the original plaintext. We need to avoid this to avoid these attacks.
- I.e. **CCA-security implies non-malleability**.
- A non-malleable encryption scheme is s.t. if the adversary tries to modify a given ciphertext, the result is **either an invalid ciphertext or one that decrypts to a plaintext having no relation to the original one**.
- So far the CCA attack we discussed seems contrived, but you will see in public-key crypto they are more easy to visualize.
- For more practical CCA attacks in SKE, read Section 3.7.2 in Katz and Lindell textbook.
- We will also see how to build CCA-secure schemes in the coming lectures.



# Chosen-ciphertext attacks (CCA)

- Some literature classifies them into two: **Nonadaptive (CCA1 or Lunchtime attacks)** and **adaptive (CCA2)**.
- CCA1 or lunchtime attack is when  $\mathcal{A}$  has access to the decryption oracle only for a limited amount of time like for e.g. when the user of a computer is out for lunch.
- It is modelled as  $\mathcal{A}$  chooses the ciphertexts to be decrypted without seeing any of the resulting plaintexts.
- An adaptive CCA or CCA2 is when  $\mathcal{A}$  has unlimited access to the decryption oracle, before and after the challenge ciphertext, with the only condition that the challenge one cannot be queried.
- This is the one we looked at formally.



# How to achieve CCA security?

- So at the end of the day, we have an active adversary Mallory that can send ciphertexts to Bob and get them decrypted.
- A layman view of CCA.
- What can Bob do?
- Use a CPA-secure encryption scheme but have Bob accept and decrypt only ciphertexts produced by Alice.
- Mallory can not create new ciphertexts that will be accepted by Bob.
- Building a CCA-secure  $\Pi$  now reduces to a problem of building a CPA-secure  $\Pi$  with message authentication.

# Message Authentication

- Secure communication till now has been how to do **secret communication over an open channel**.
- Basically, we showed how encryption (private key) can be used to prevent an Eve or Mallory from learning anything about the messages over an unprotected channel.
- **What if our security concerns are not related to secrecy?**
- One other concern: **how to guarantee integrity of our message? i.e. message authentication?**
- Each honest party should be able to identify when a message it receives is indeed sent by the party that is claiming it has sent it and **not modified in between**.

# Message Authentication

- Consider the case when a user communicated with a bank and requests for a transfer of some money:
  - ▶ How does the bank know if it has indeed come from the user?
  - ▶ Even if it is from a legitimate user, how can we be sure that the message has not been tampered? As in the account number has not been changed for example?
- Standard error-correcting techniques won't work since they are for random errors and not for malicious ones which know what exactly to be changed.

# Message Authentication

Another scenario where message integrity takes precedence: Web cookies.

- When you go to buy something from a shopping site, any state generated by a session for e.g: contents of your shopping cart is placed as a **cookie** with you, the client.
- It is sent to the server as part of each message you sent.
- If there are some items for which you get some special discount (user-specific pricing) then the server needs to make sure you have not modified the cookie to alter the prices of the product!
- Note: none of the details are secret.

# Encryption Vs Message Authentication

- Encryption does not solve the problem of message authentication.
- Encryption using stream ciphers:
  - ▶  $c := G(k) \oplus m$  where  $G$  is a PRG.
  - ▶ Flipping any bit in  $c$  results in the same bit being flipped in the message.
  - ▶ Thus given a  $c$ , an encryption for  $m$  we can obtain a  $c'$  s.t. its decryption is the same as  $m$  with one bit flipped.
  - ▶ We can do the same attack for OTPs, **so not even perfect secrecy guarantees integrity.**

# What about block ciphers?

- Same attack we described above works for OFB and CTR mode since they both generate a pseudorandom stream using block ciphers.
- This implies CPA-secure encryption schemes is not enough to prevent message tampering.
- What about CBC or ECB?
  - ▶ Note that ECB does not even provide CPA-security!
  - ▶ Both modes need inverting a PRF  $F$  and  $F_k^{-1}(x)$  and  $F_k^{-1}(x')$  will be unrelated even if  $x$  and  $x'$  differ in only one bit.
  - ▶ But still there can be predictable changes in the plaintext.
  - ▶ In ECB, flipping a bit in the  $i$ th block of ciphertext changes only the  $i$ th block of plaintext.

# What about block ciphers?

- In CBC mode, flipping  $j$ th bit of  $IV$  changes only the  $j$ th bit of  $m_1$ .
- All other plaintext blocks remain unchanged as  $m_i = F_k^{-1}(c_i) \oplus c_{i-1}$  and  $c_i$  and  $c_{i-1}$  are not modified.
- The first block of CBC (possibly header info) can be changed arbitrarily!



# Defining MACs

- The communicating parties need to share a secret – in private key setting Alice and Bob share a secret key  $k$  generated by the *Gen* algorithm.
- A *MAC tag* or just *tag*  $t$  based on  $m$  and  $k$  is computed. It is computed by the *MAC* algorithm.,  $t \leftarrow \text{MAC}_k(m)$ .
- It could be *randomized* and hence the  $\leftarrow$ .
- On receiver end, he runs the *deterministic Verify* algorithm on  $(m, t)$  to ensure that the given tag is valid, if valid output is a bit  $b = 1$ , else 0.
- Correctness: For all  $k$  generated by *Gen* and all messages  $m$ ,  $\text{Verify}_k(m, \text{MAC}_k(m)) = 1$ .

# Defining MACs



Alice

$$t \leftarrow \text{MAC}_k(m)$$

$(m, t)$



Bob

$$\text{Verify}_k(m, t) = 1$$

# Defining MACs



Alice

 $(m, t)$ 

Bob

 $(m', t)$  or  $(m', t')$ 

$Verify_k(m', t) = 0$   
Fails!



Mallory

# Security of MACs

- No efficient adversary should be able to generate a valid tag on a different message **that was not previously sent**.
- We need to formalize this idea, so we need to define a break of the scheme.
- Note that : an eavesdropping adversary can see all the messages sent by these parties along with their corresponding MAC tags.
- The adversary may also be able to influence the content of these messages, whether directly or indirectly. For e.g: the user changes the contents of the cookie stored on his computer.
- To formalize this we allow for  $\mathcal{A}$  to request tags for **any messages of its choice**, i.e. access to an **MAC oracle**  $MAC_k(\cdot)$ .

# Breaking MACs

- A break is when  $\mathcal{A}$  produces  $(m, t)$ 
  1.  $t$  is a valid tag for  $m$ .
  2.  $\mathcal{A}$  did not request a MAC tag on  $m$  from the oracle.
- First case covers when honest parties are fooled into thinking that  $m$  came from a honest party.
- Second is a **replay attack** which is a serious attack *but not considered a break of a MAC*. i.e.,  $\mathcal{A}$  copies  $(m, t)$  sent previously by one of the legitimate parties.
- Before defining security we as usual give an experiment.

# Message authentication experiment

$MAC - \text{forge}_{\mathcal{A}, \Pi}(1^n)$ :

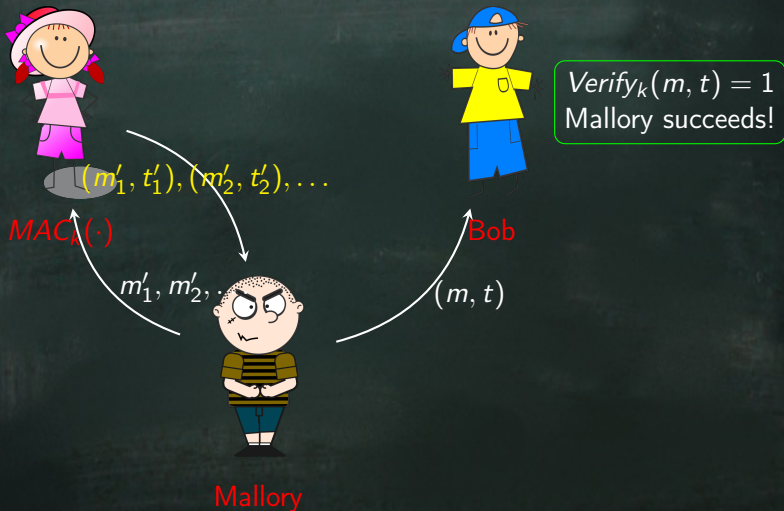
1. A key  $k$  is generated by  $\text{Gen}(1^n)$ .
2.  $\mathcal{A}$  is given input  $1^n$  and access to  $\text{MAC}_k(\cdot)$ .
3. After polynomial such accesses  $\mathcal{A}$  outputs  $(m, t)$ . Let  $\mathcal{Q}$  denote the set of all queries that  $\mathcal{A}$  makes to the oracle.
4.  $\mathcal{A}$  succeeds with output 1 iff (1)  $\text{Verify}_k(m, t) = 1$  and (2)  $m \notin \mathcal{Q}$ .

## Definition

A message authentication code  $\Pi = (\text{Gen}, \text{MAC}, \text{Verify})$  is **existentially unforgeable under an adaptive chosen-message attack**, or just **secure**, if for PPT  $\mathcal{A}$ , there is a negligible function  $\text{negl}$  such that:

$$\Pr[\text{MAC} - \text{forge}_{\mathcal{A}, \Pi}(1^n) = 1] \leq \text{negl}(n).$$

# Message authentication experiment





# Too strong a definition?

- $\mathcal{A}$  is allowed to request MAC tags for **any messages of its choice**.
- $\mathcal{A}$  succeeds if it can output a valid tag on **any previously unauthenticated message**.
- In real-life it would be **meaningful messages** that it should request MAC tags for and also succeed only for such a subset.
- What is meaningful? Too Application specific.
- Replay attacks are serious but MACs cannot work against them. Use sequence numbers or time-stamps.

# Strong MACs

- The specification has been that  $\mathcal{A}$  cannot generate a valid tag on a new message that was never previously authenticated.
- But we can still have this scenario: an attacker might be able to generate a new tag on a previously authenticated message.
- More precisely, a *MAC* guarantees that if an attacker learns tags  $t_1, \dots$ , on  $m_1, \dots$ , then it will not be able to forge a valid tag  $t$  on any message  $m \notin \{m_1, \dots\}$ .
- $\mathcal{A}$  may still forge a different valid tag  $t_1'$  on  $m_1$ .
- *MAC* – *sforge* takes care of that by considering  $(m, t) \in \mathcal{Q}$ , i.e. not messages but pairs of oracle queries and responses.
- A **strong MAC** is one in which the probability of  $\mathcal{A}$  succeeding *MAC* – *sforge* is negligible.

# Canonical Verification and Strong MAC

- If  $MAC$  is deterministic, canonical verification is simply re-computing the tag and check for equality.
- $Verify_k(m, t)$  will first do  $\bar{t} := MAC_k(m)$  and output 1 iff  $t = \bar{t}$ .
- If  $\Pi$  uses canonical verification then  $MAC$  is deterministic.
- A deterministic  $MAC$  can use canonical verification, **but it doesn't have to**.
- The following is easy to see (Assignment question!)

## Theorem

Let  $\Pi = (Gen, MAC, Verify)$  be a secure MAC that uses canonical verification ( $\Rightarrow$  **deterministic**). Then  $\Pi$  is a strong MAC.